

# SDNRacer

Concurrency Analysis for SDNs



***Ahmed El-Hassany***

Jeremie Miserez

Pavol Bielik

Laurent Vanbever

Martin Vechev

***ETH*** zürich

<http://sdnracer.ethz.ch>

# SDNRacer

Finds violations in SDN controllers:

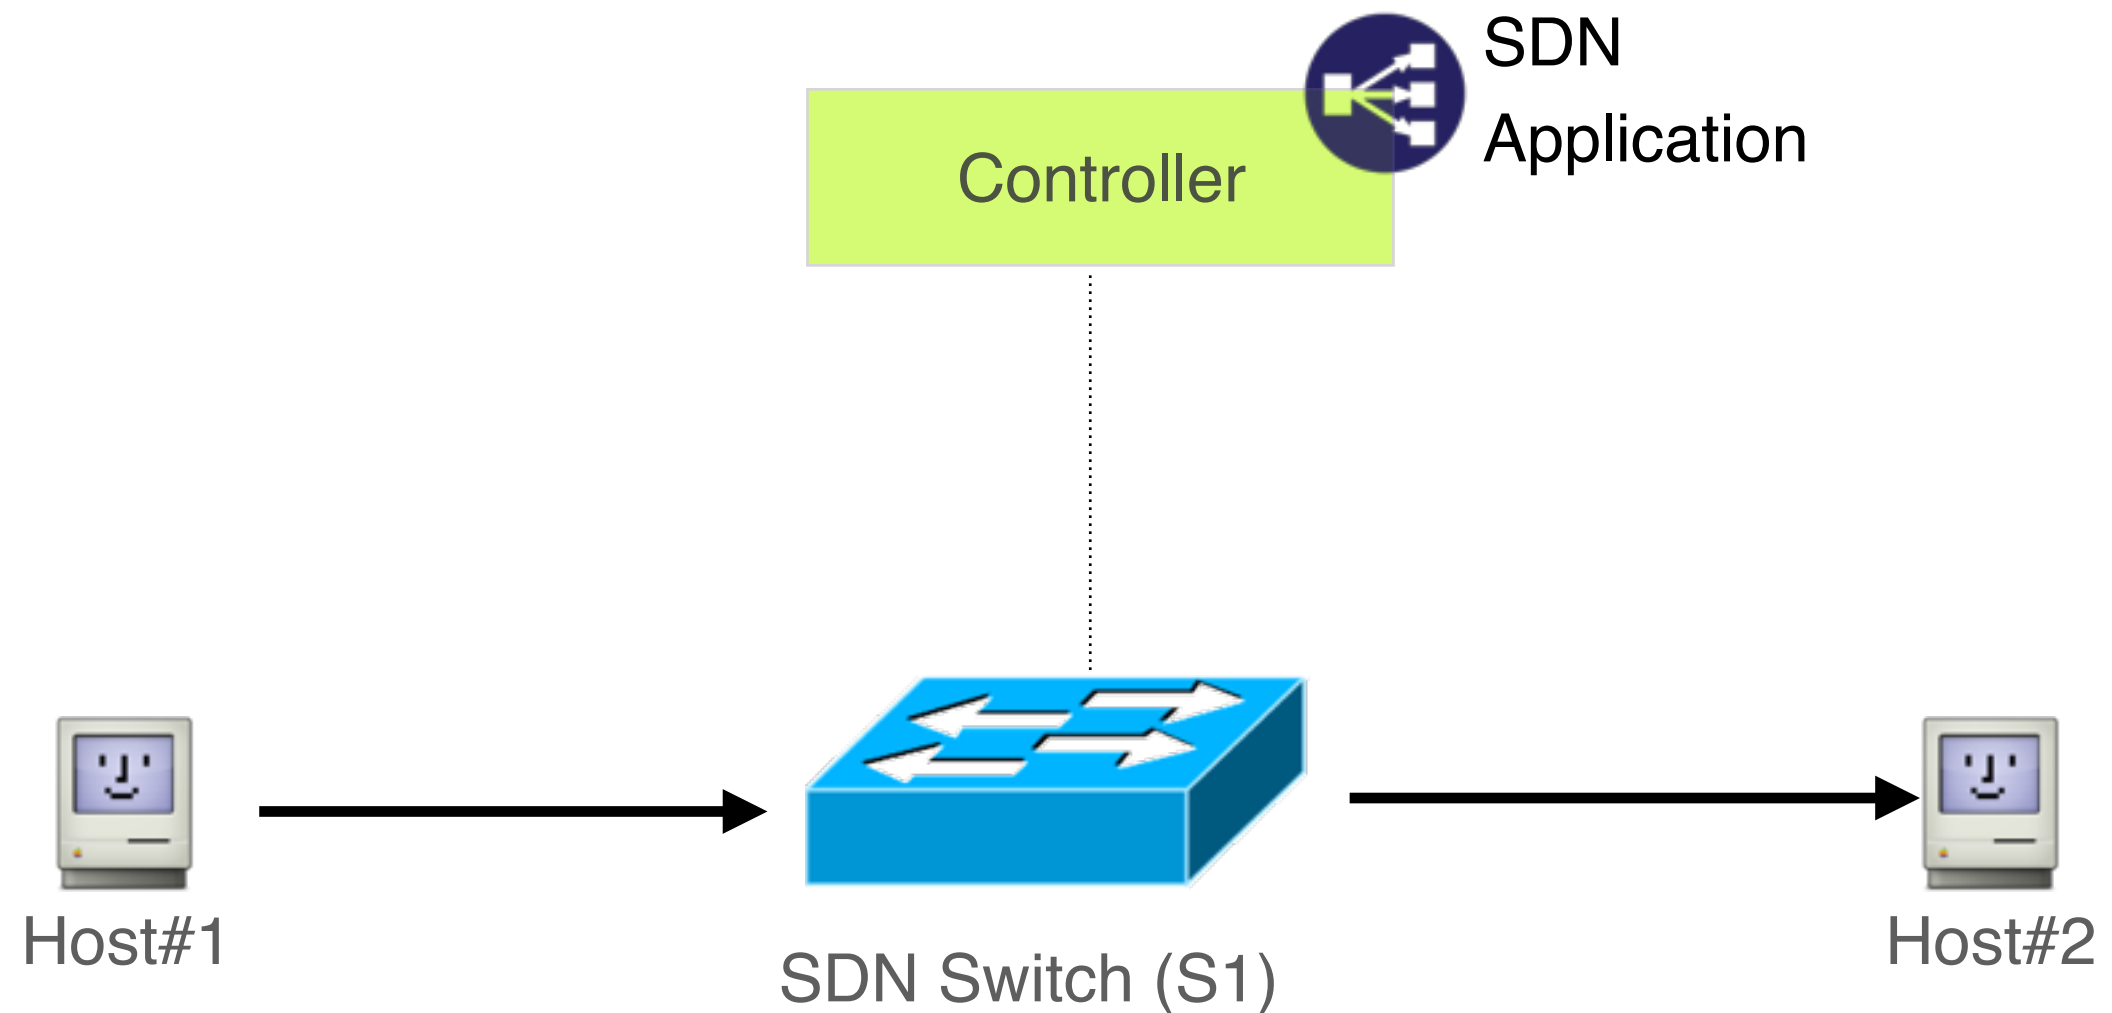
**Race Freedom**


**Update Isolation**

**Packet Coherence**

Violations of these properties can cause serious bugs in the network.

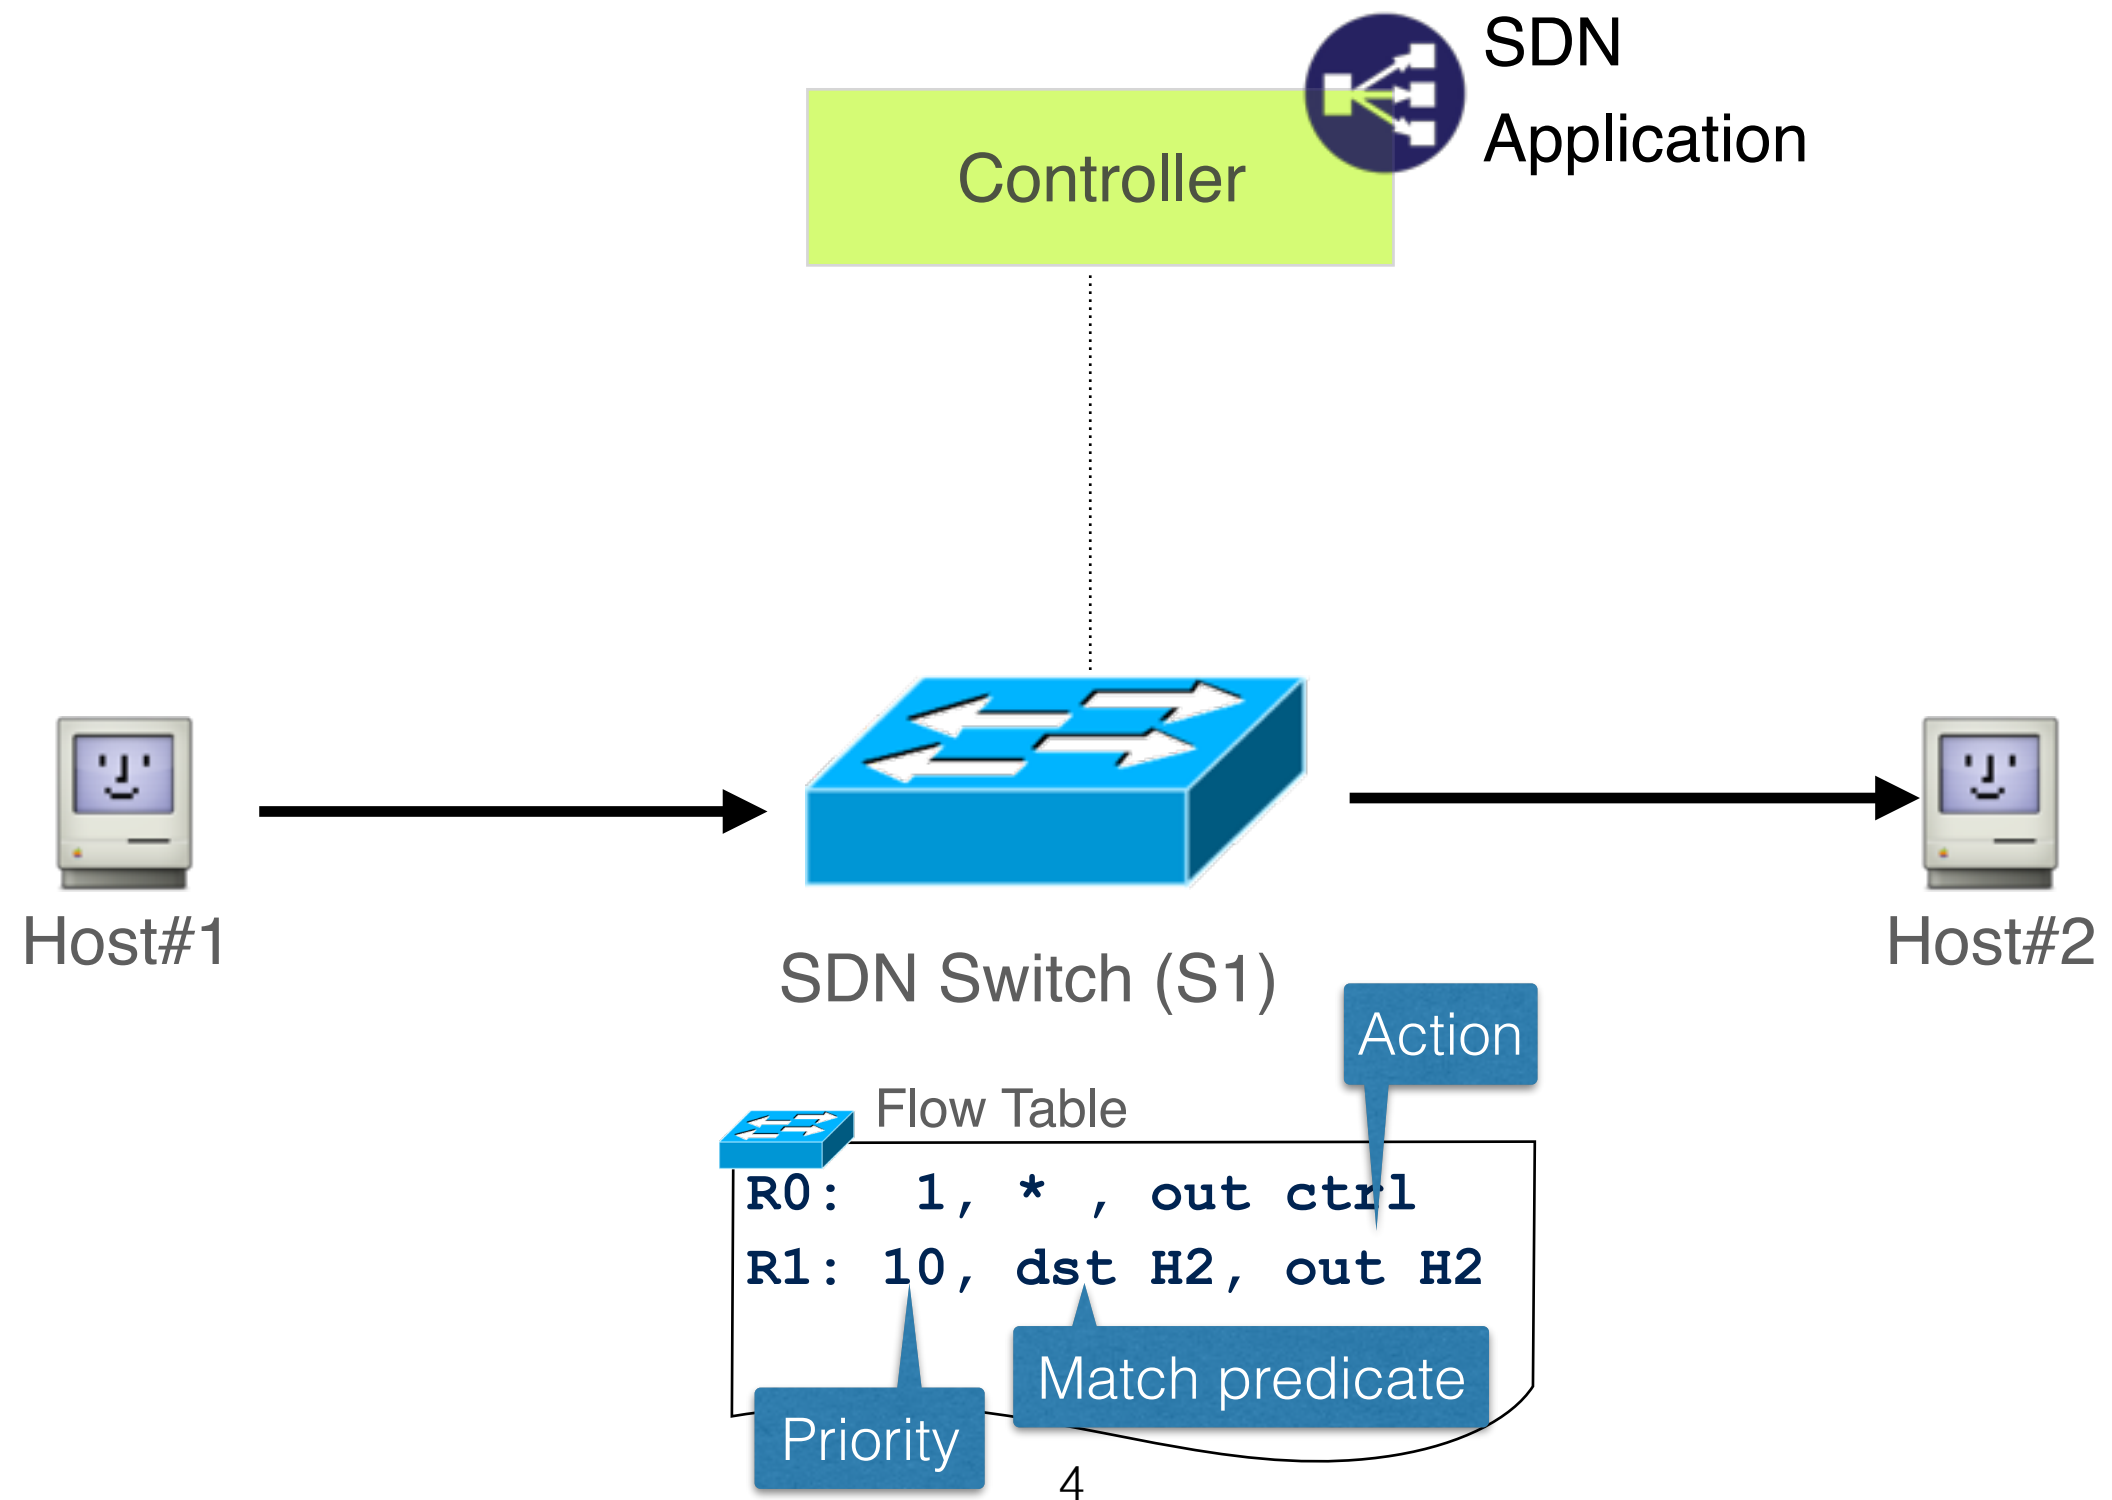
# SDN Overview



 Flow Table


R0:	1, *	out ctrl
R1:	10, dst H2	out H2

# SDN Overview



# Example SDN App: Load-Balancer

# Load Balancer Application



# Load Balancer Application



Controller

```
if dst == server:
    rep = rep[idx] idx = (idx+1)%2
    install_path(src, rep)
    install_path(rep, src)
    packet_out(pkt, in sw)
```

# Load Balancer Application



Controller

Round-Robin  
Server Selection

```
if dst == server:  
    rep = rep[idx] idx = (idx+1)%2  
    install_path(src, rep)  
    install_path(rep, src)  
    packet_out(pkt, in sw)
```



# Load Balancer Application



Controller

```
if dst == server:  
    rep = rep[idx]  idx = (idx+1)%2  
install_path(src, rep)  
install_path(rep, src)  
packet_out(pkt, in sw)
```

1. Find the shortest Path.
2. Write a flow entry on each switch on the path.

# Load Balancer Application

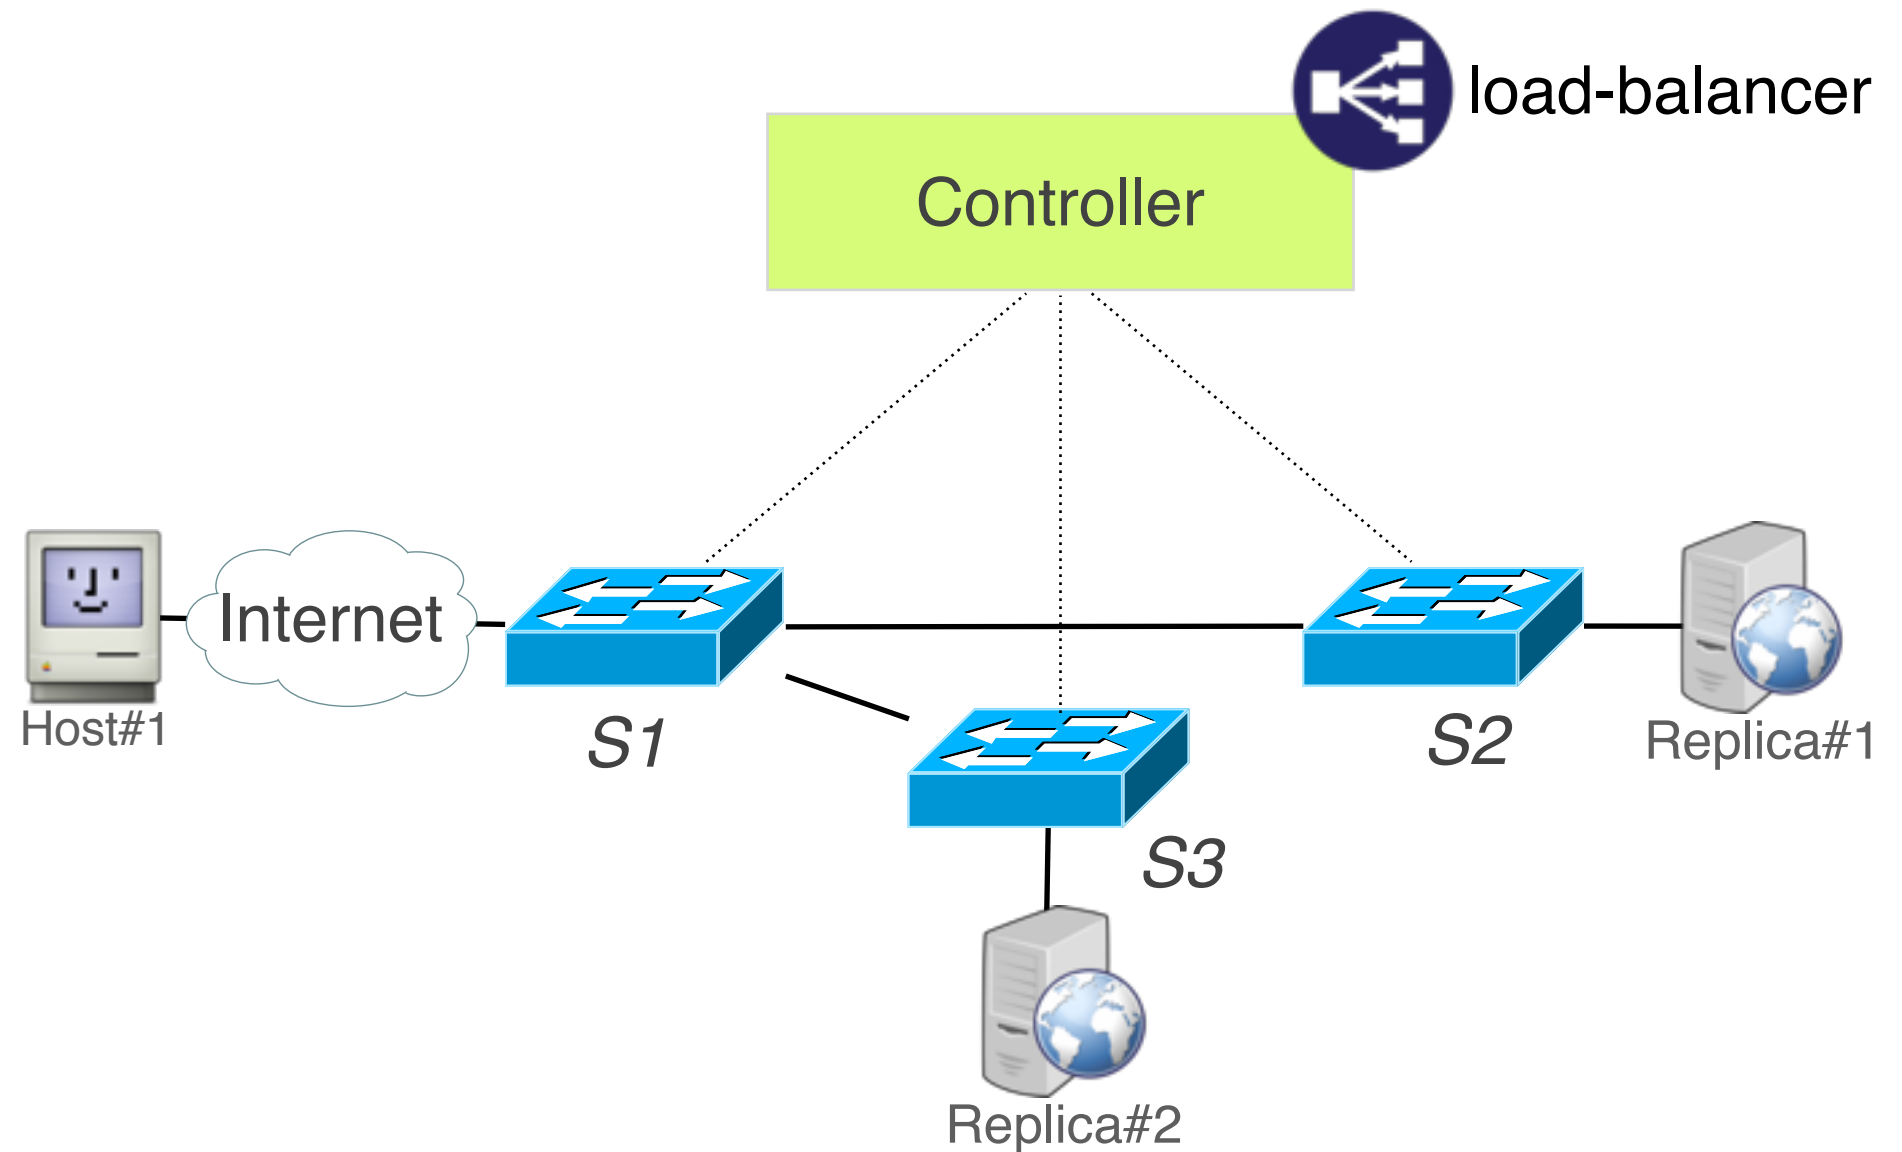


Controller

```
if dst == server:  
    rep = rep[idx]  idx = (idx+1)%2  
    install_path(src, rep)  
    install_path(rep, src)  
packet_out(pkt, in sw)
```

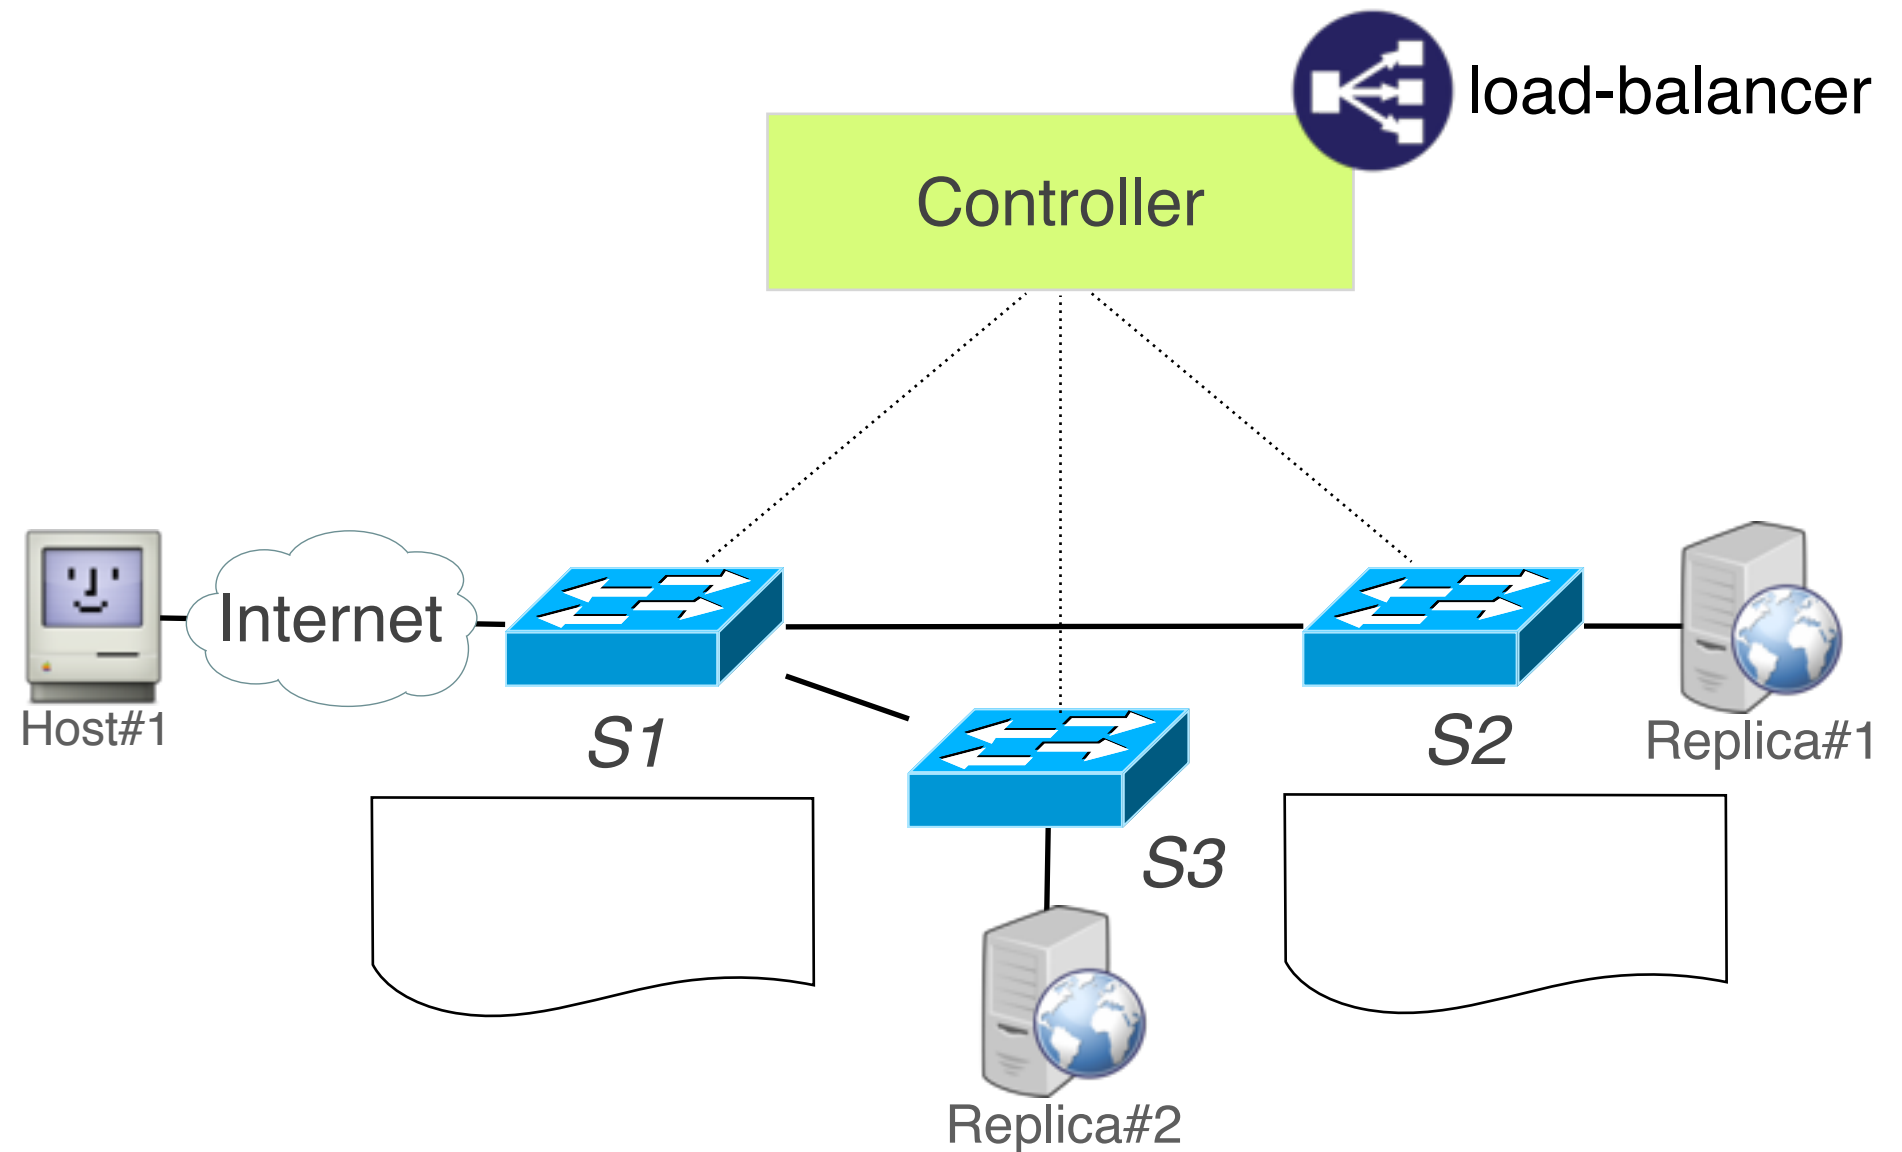
Send the packet back to the dataplane.

# Load Balancer Application



# Load Balancer Application


Recorded Event Trace:



# Load Balancer Application

Recorded Event Trace:

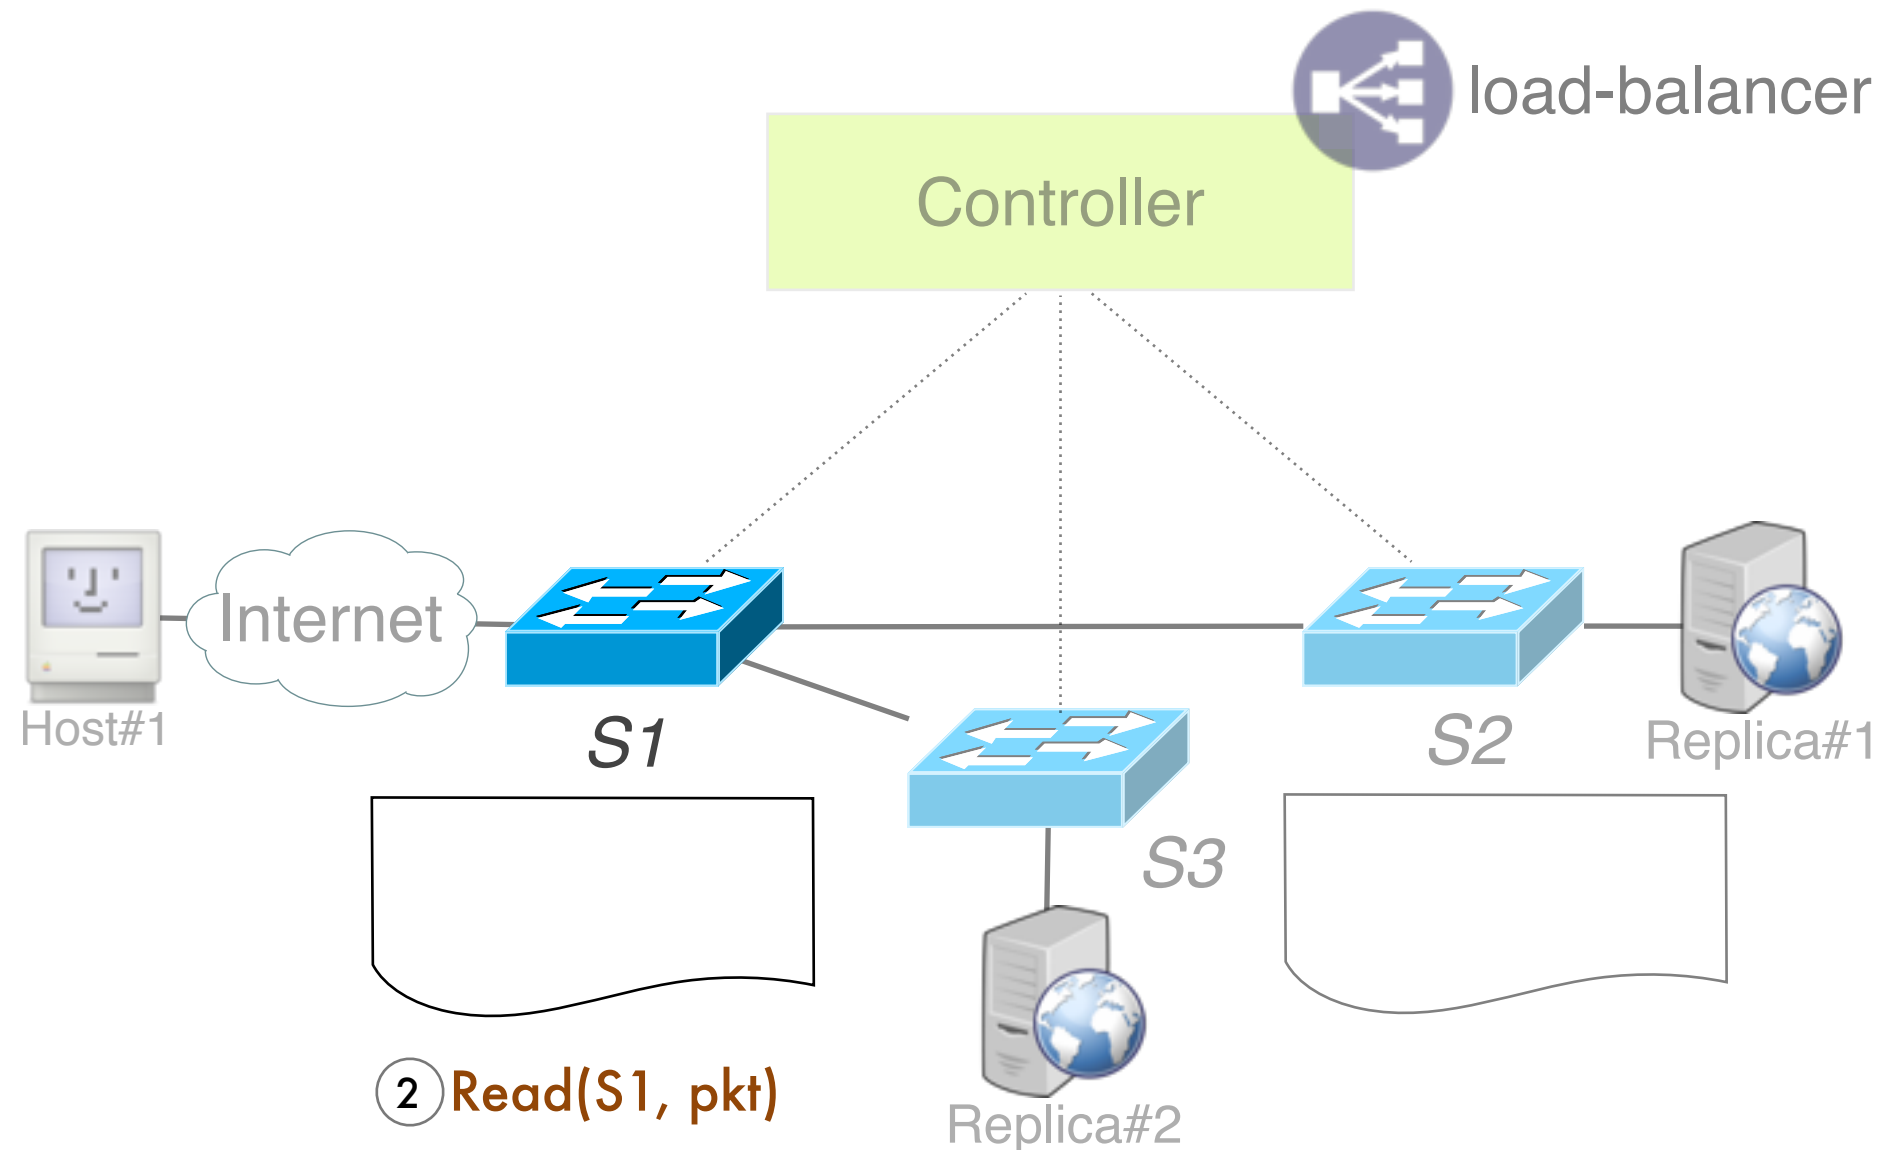
① Host Send (pkt, dst=srv)



# Load Balancer Application

Recorded Event Trace:

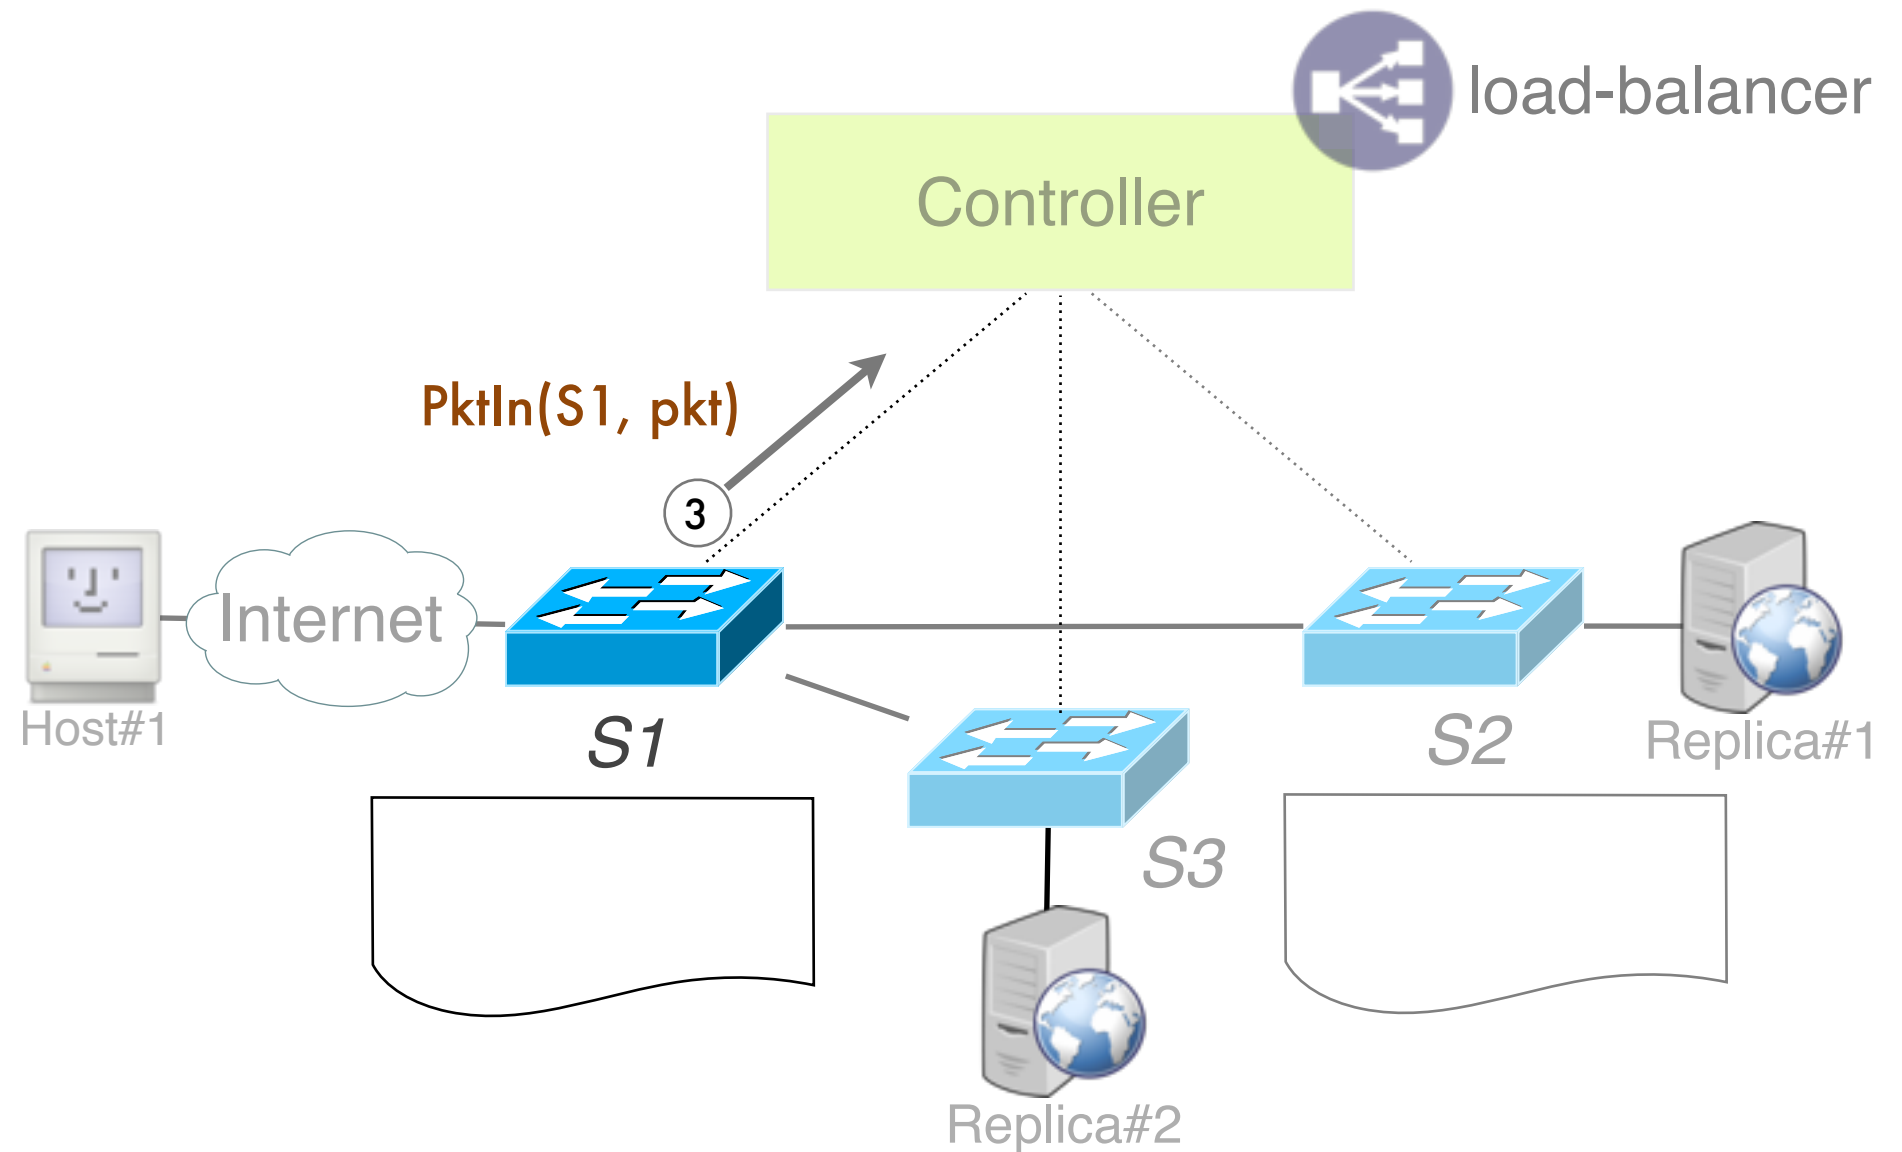
- ① Host Send (pkt, dst=svr)
- ② Read(S1, pkt)



# Load Balancer Application

Recorded Event Trace:

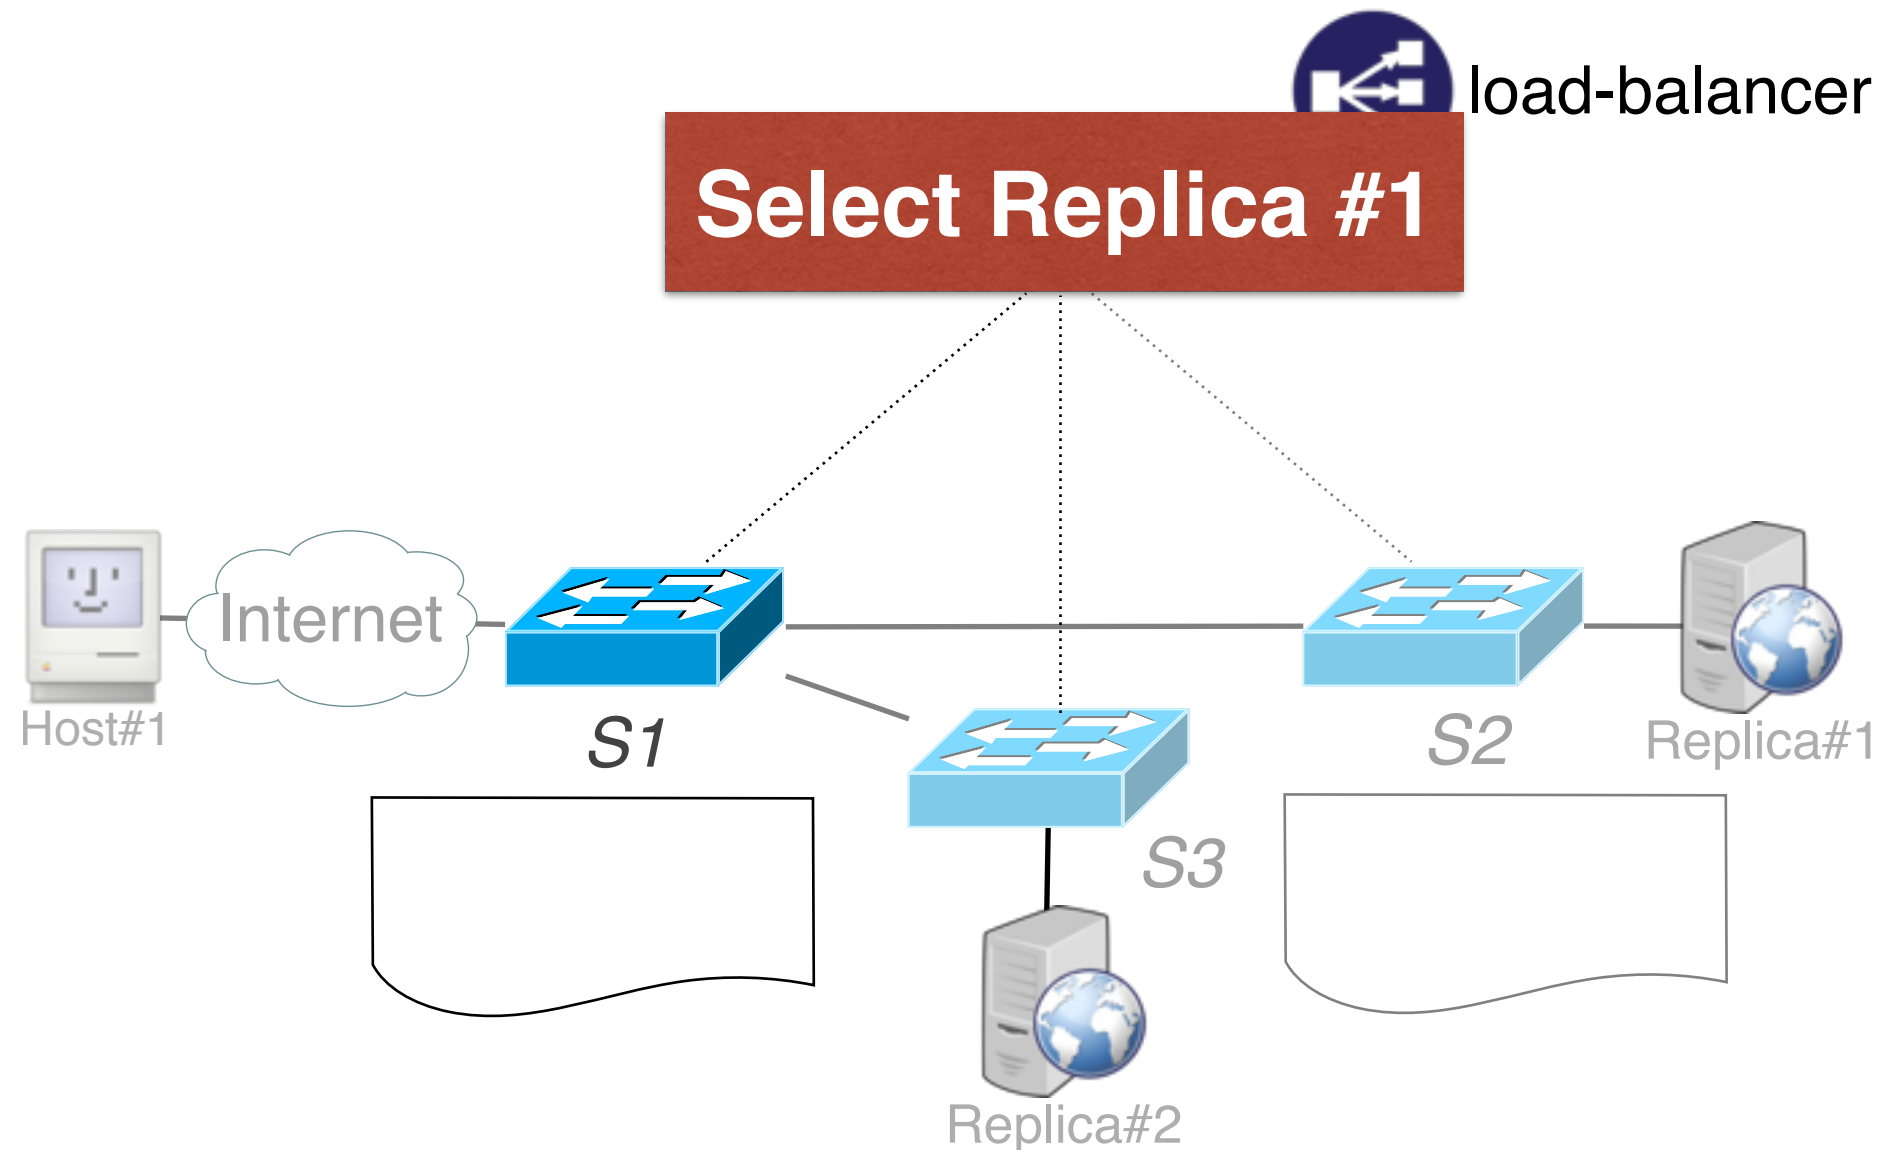
- 1 Host Send (pkt, dst=srv)
- 2 Read(S1, pkt)
- 3 PktIn(S1, pkt)



# Load Balancer Application

Recorded Event Trace:

- 1 Host Send (pkt, dst=srv)
- 2 Read(S1, pkt)
- 3 PktIn(S1, pkt)

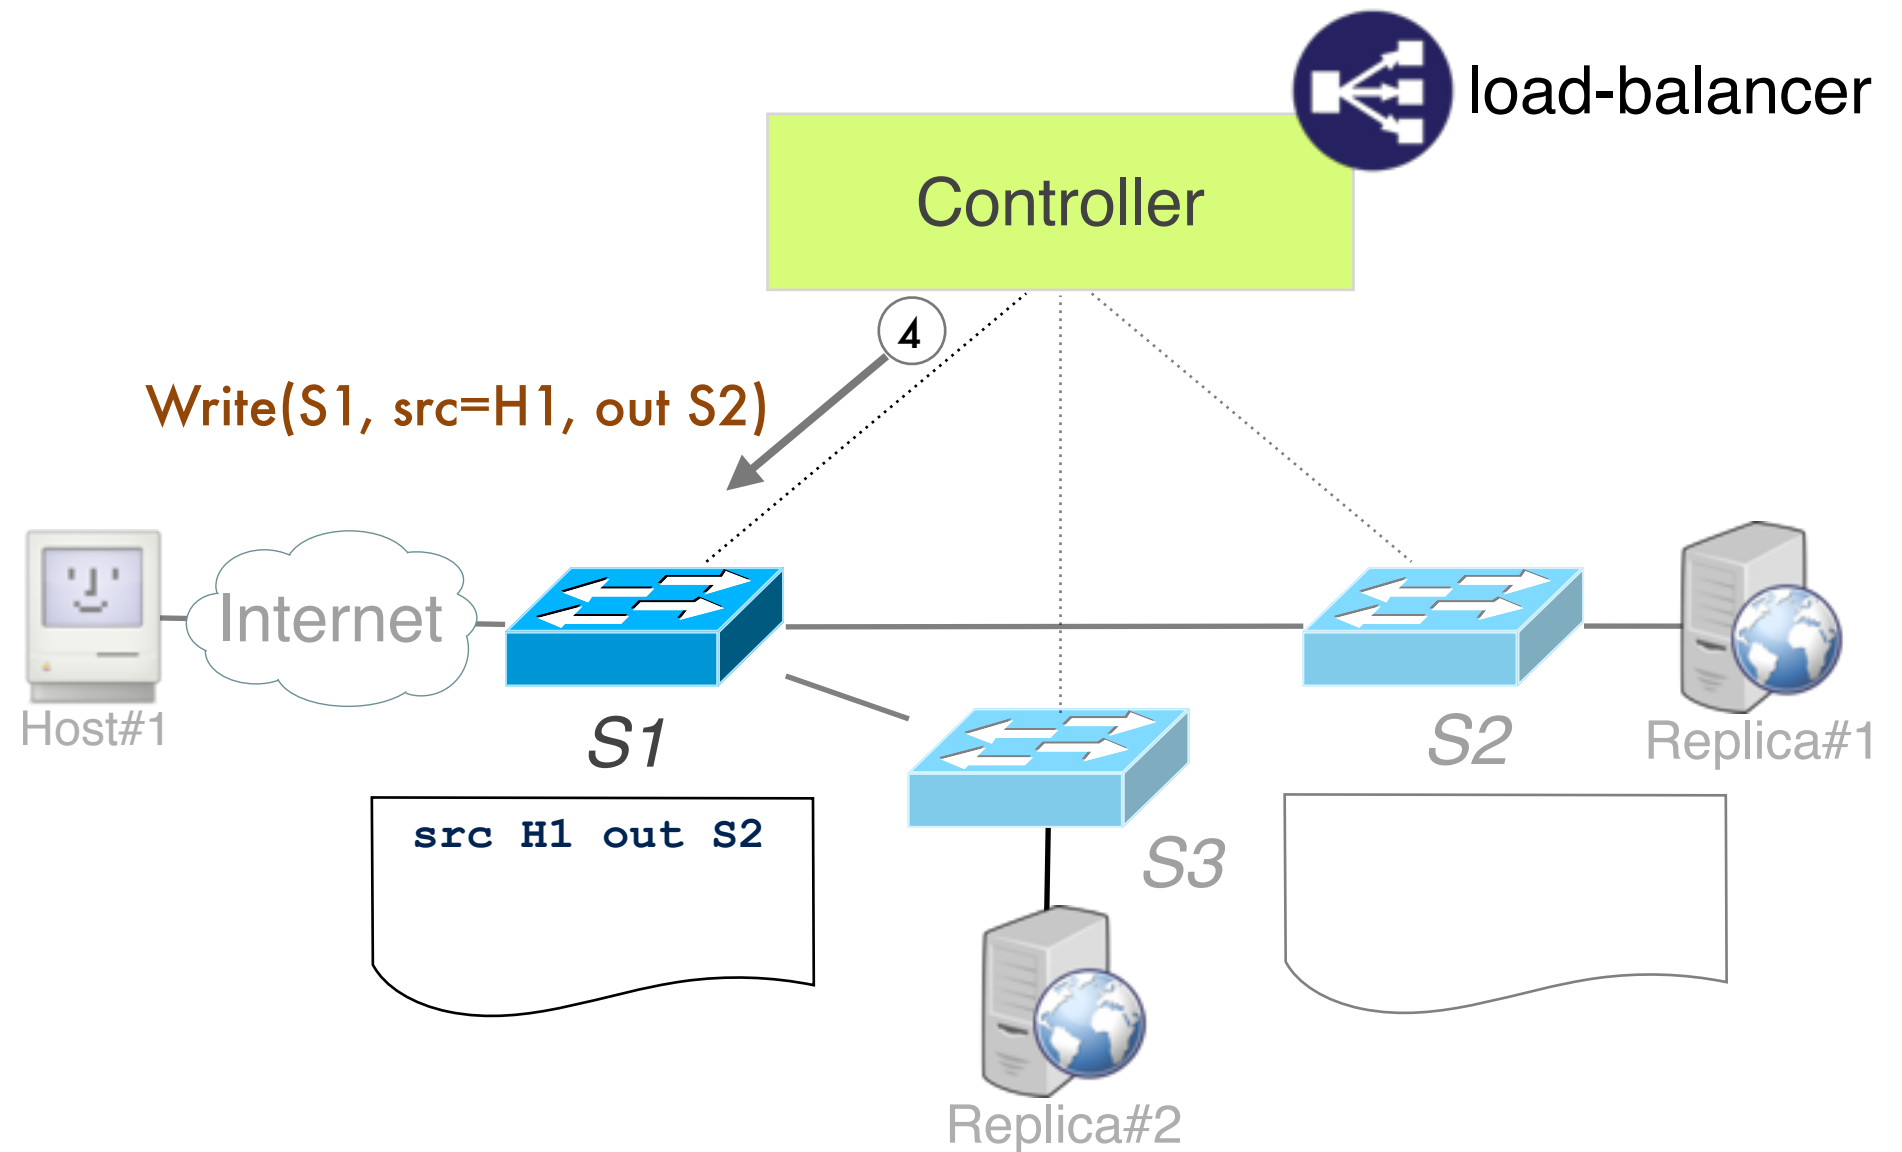




# Load Balancer Application

Recorded Event Trace:

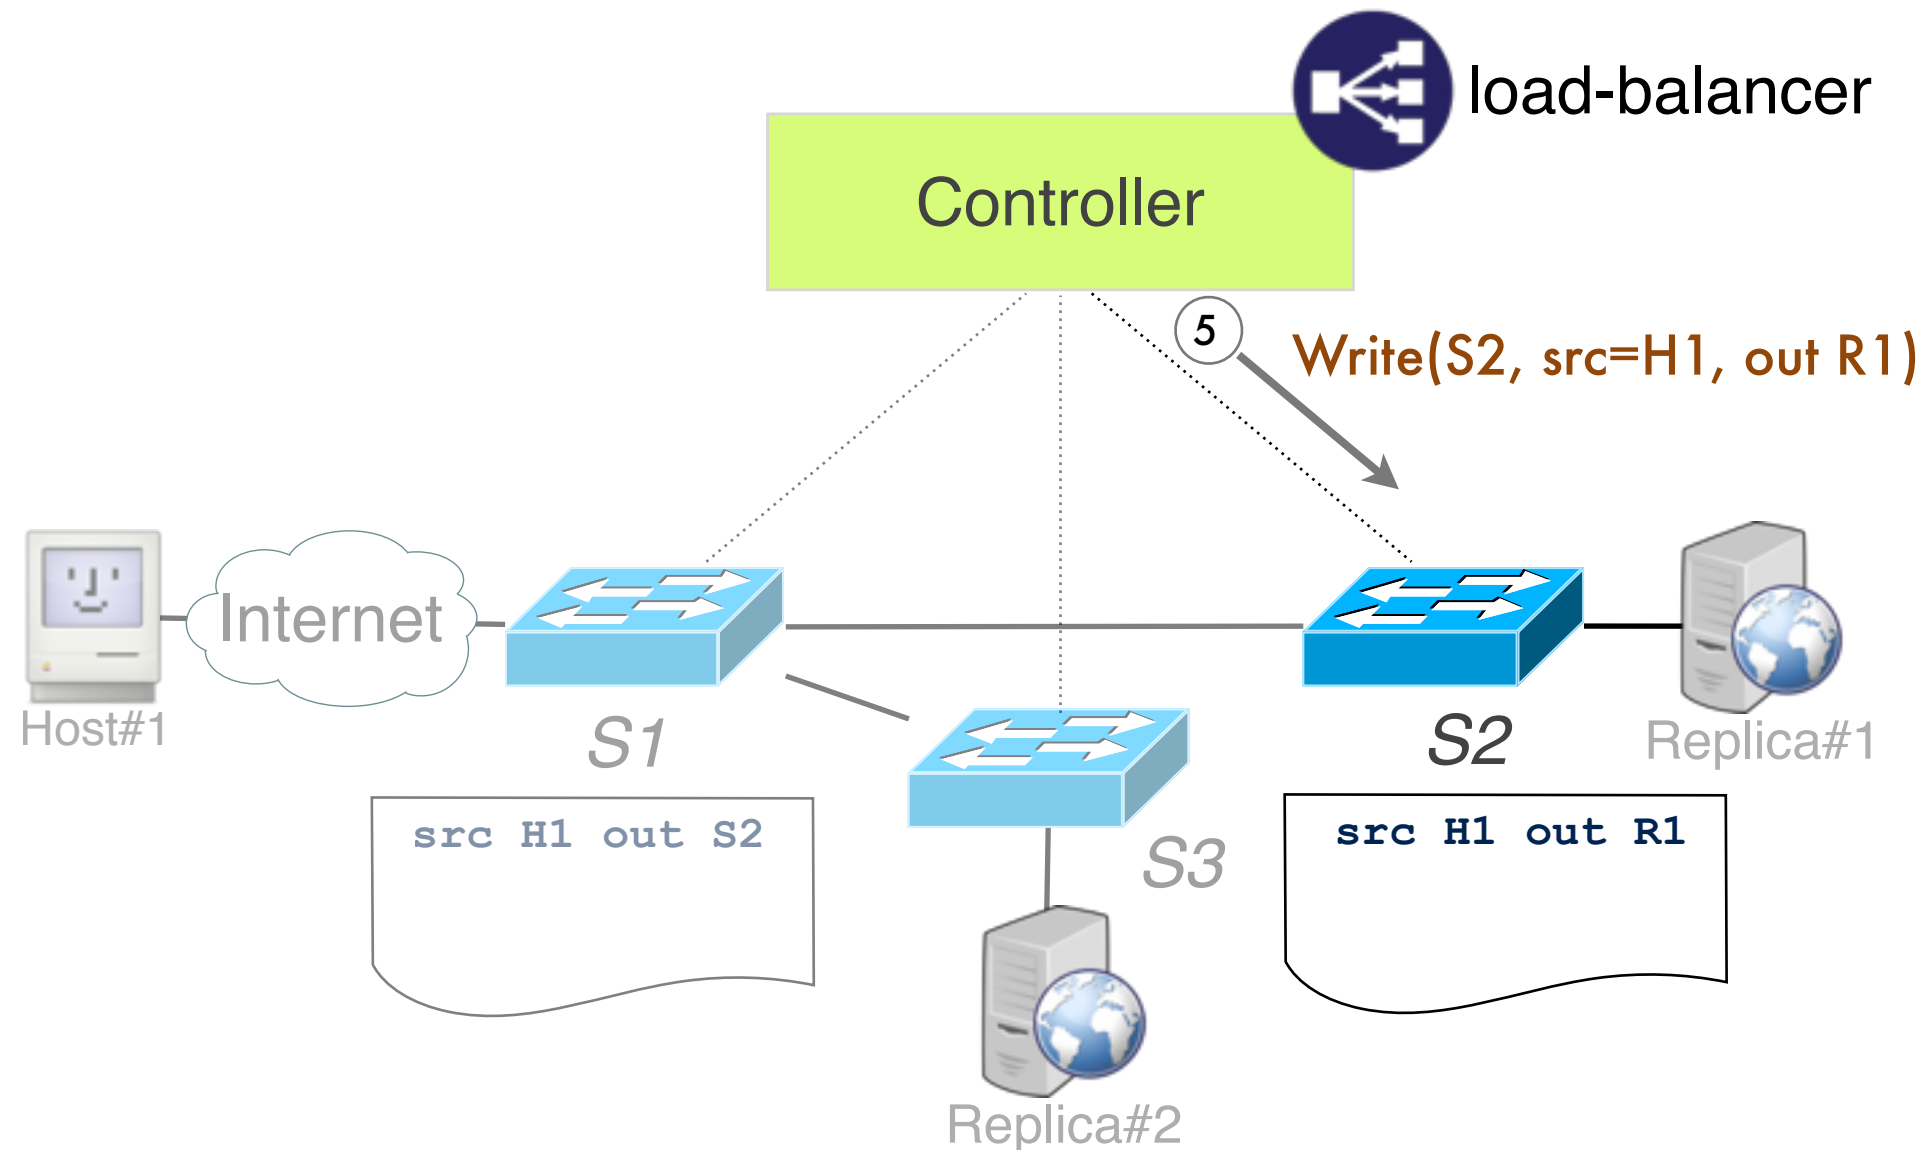
- 1 Host Send (pkt, dst=srv)
- 2 Read(S1, pkt)
- 3 PktIn(S1, pkt)
- 4 Write(S1, src=H1, out S2)



# Load Balancer Application

Recorded Event Trace:

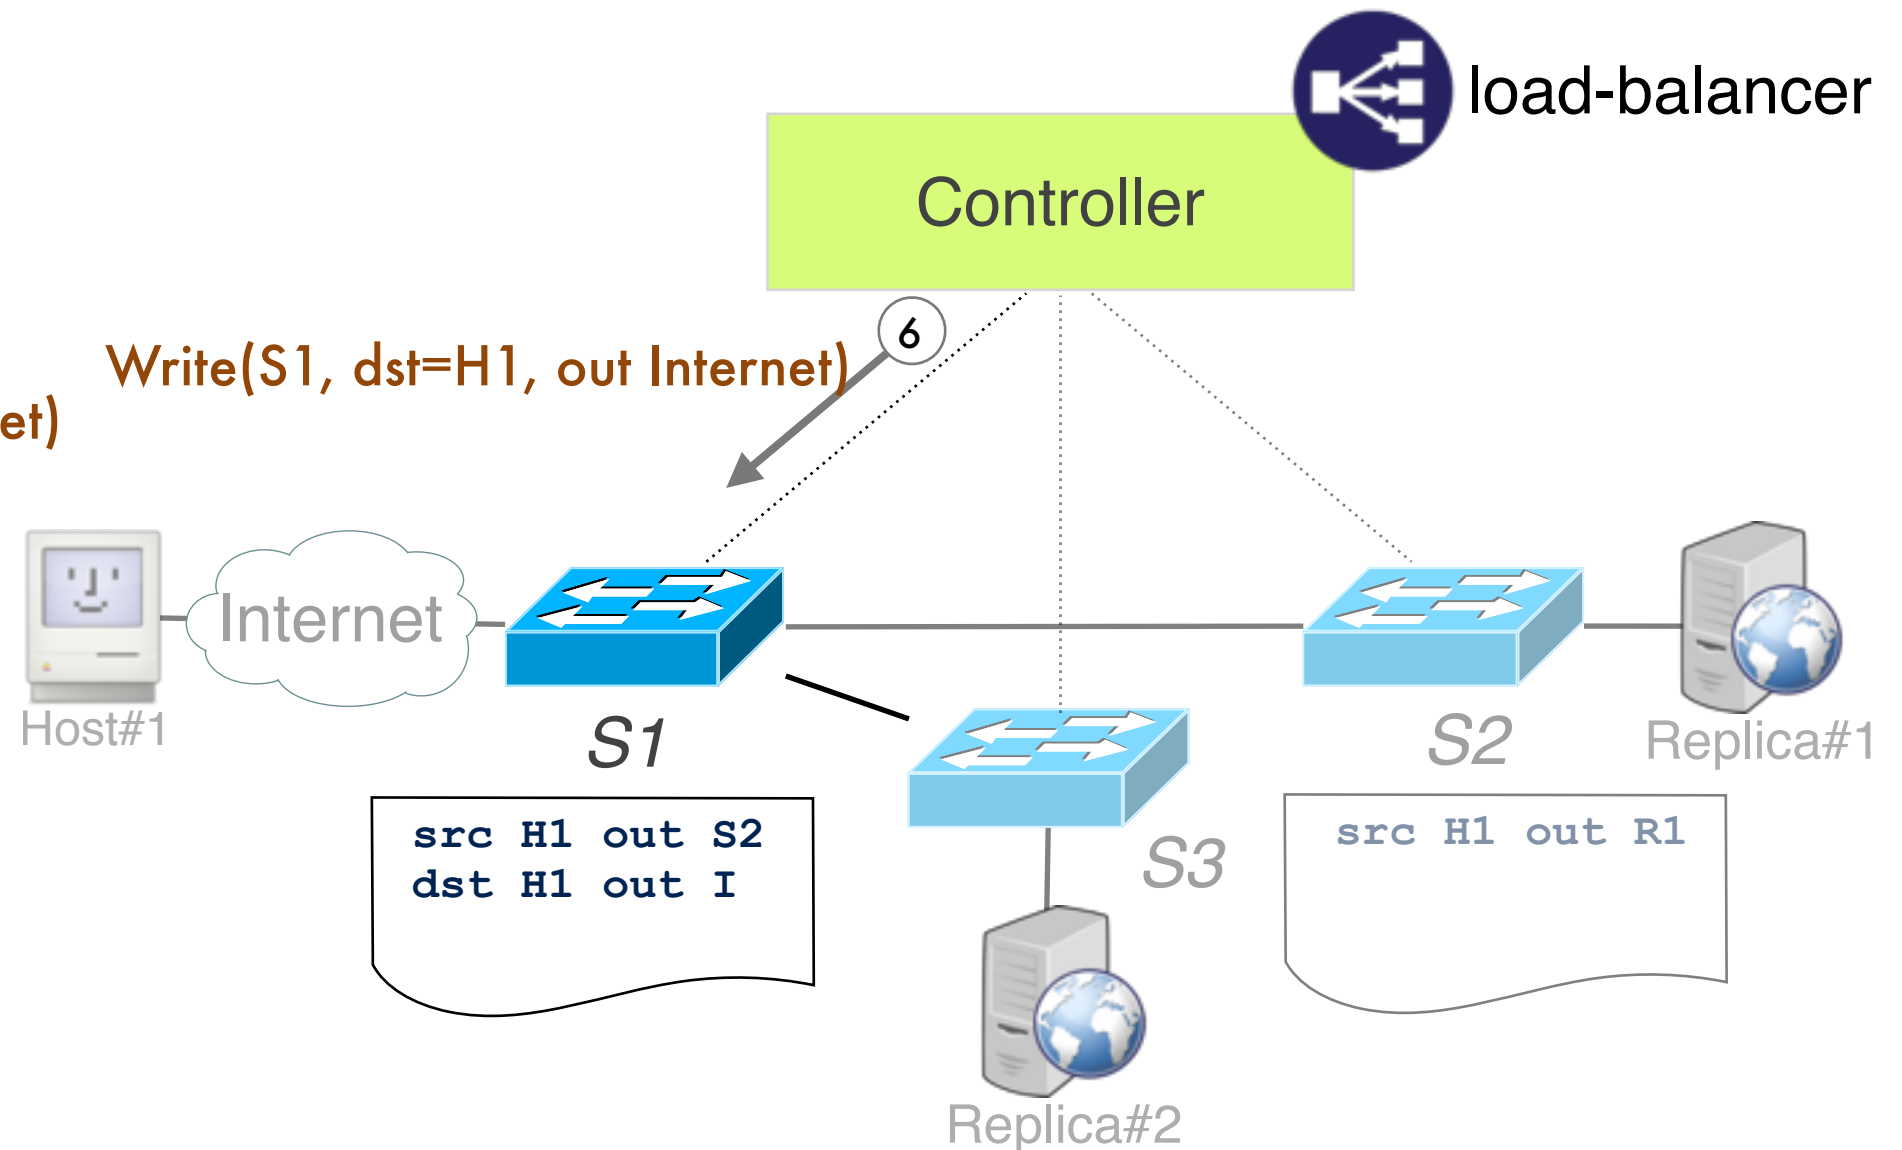
- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑤ Write(S2, src=H1, out R1)



# Load Balancer Application

## Recorded Event Trace:

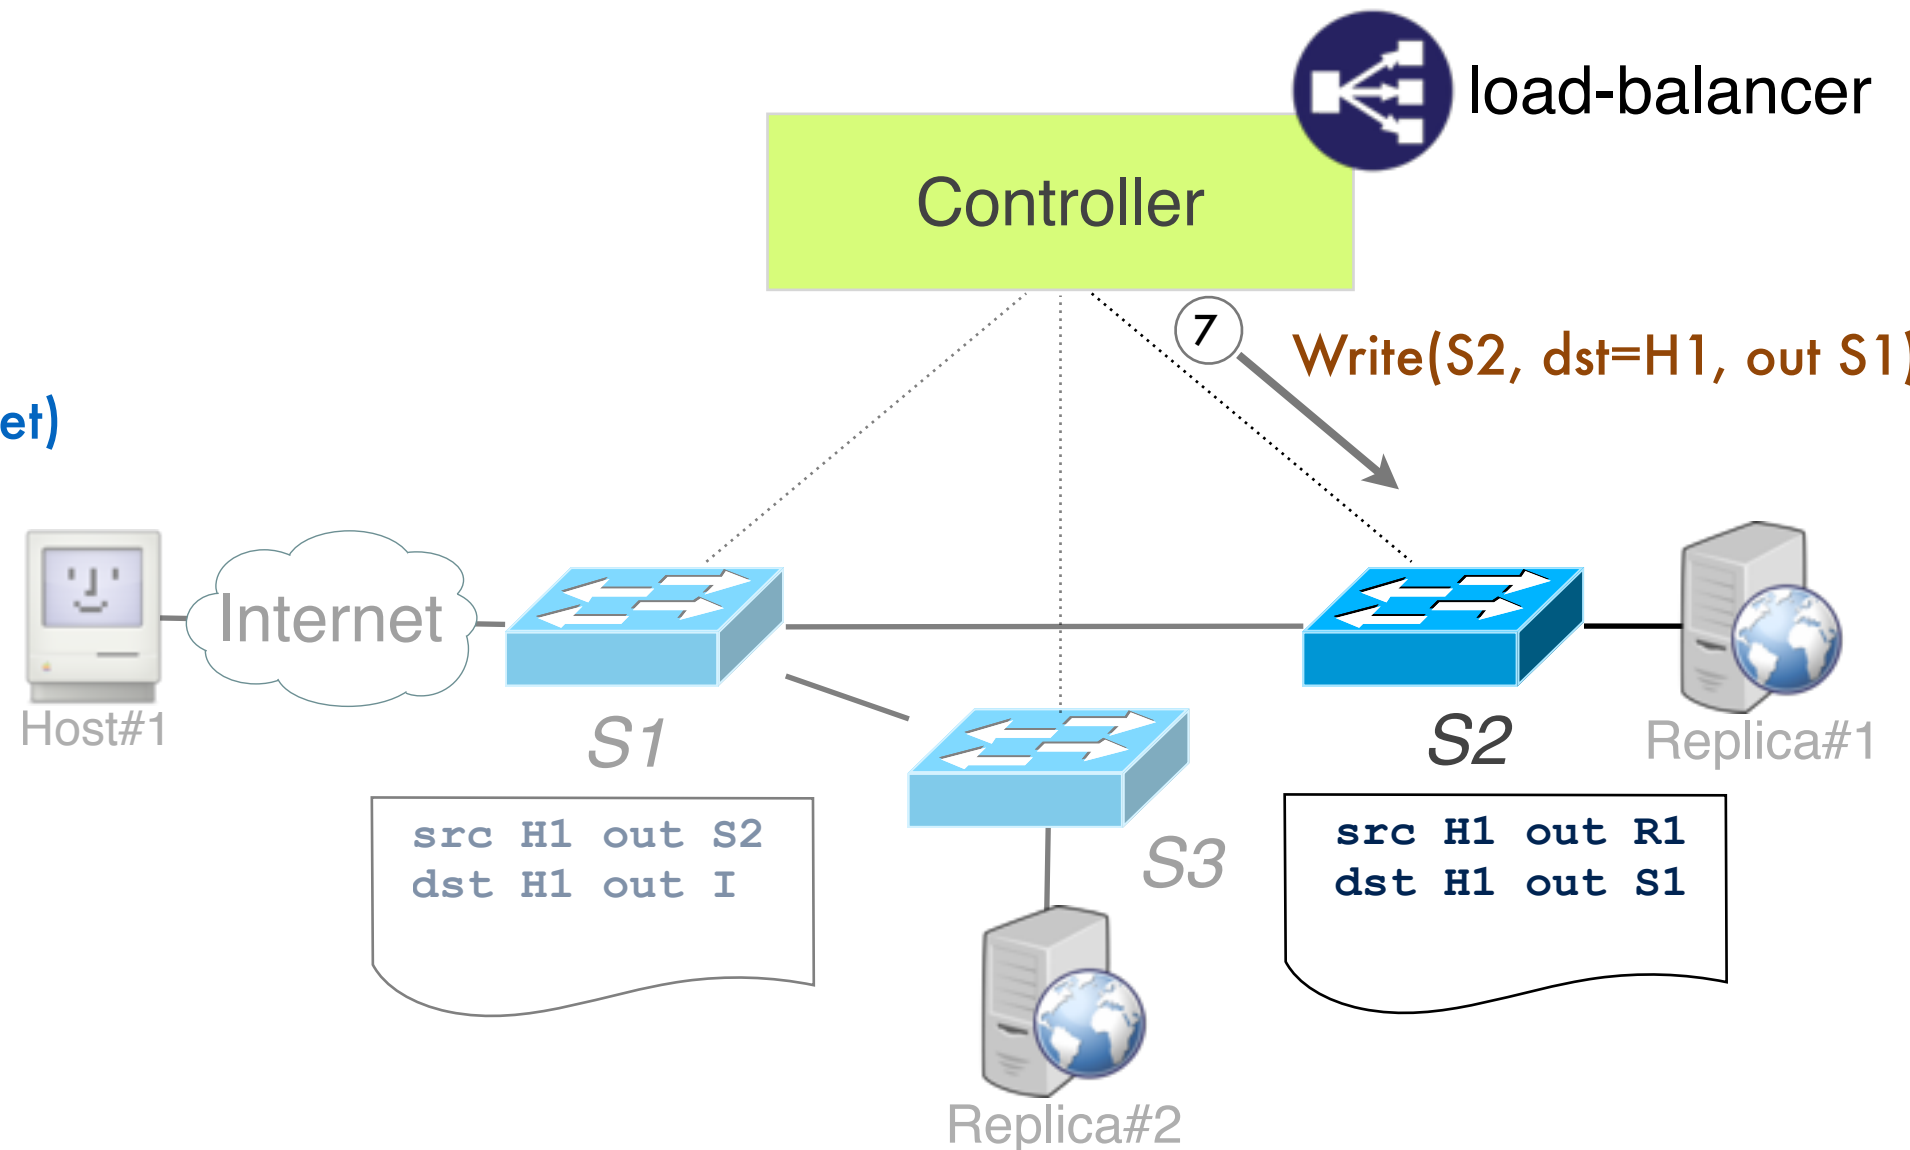
- ① Host Send (pkt, dst=srcv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑥ Write(S1, dst=H1, out Internet)



# Load Balancer Application

## Recorded Event Trace:

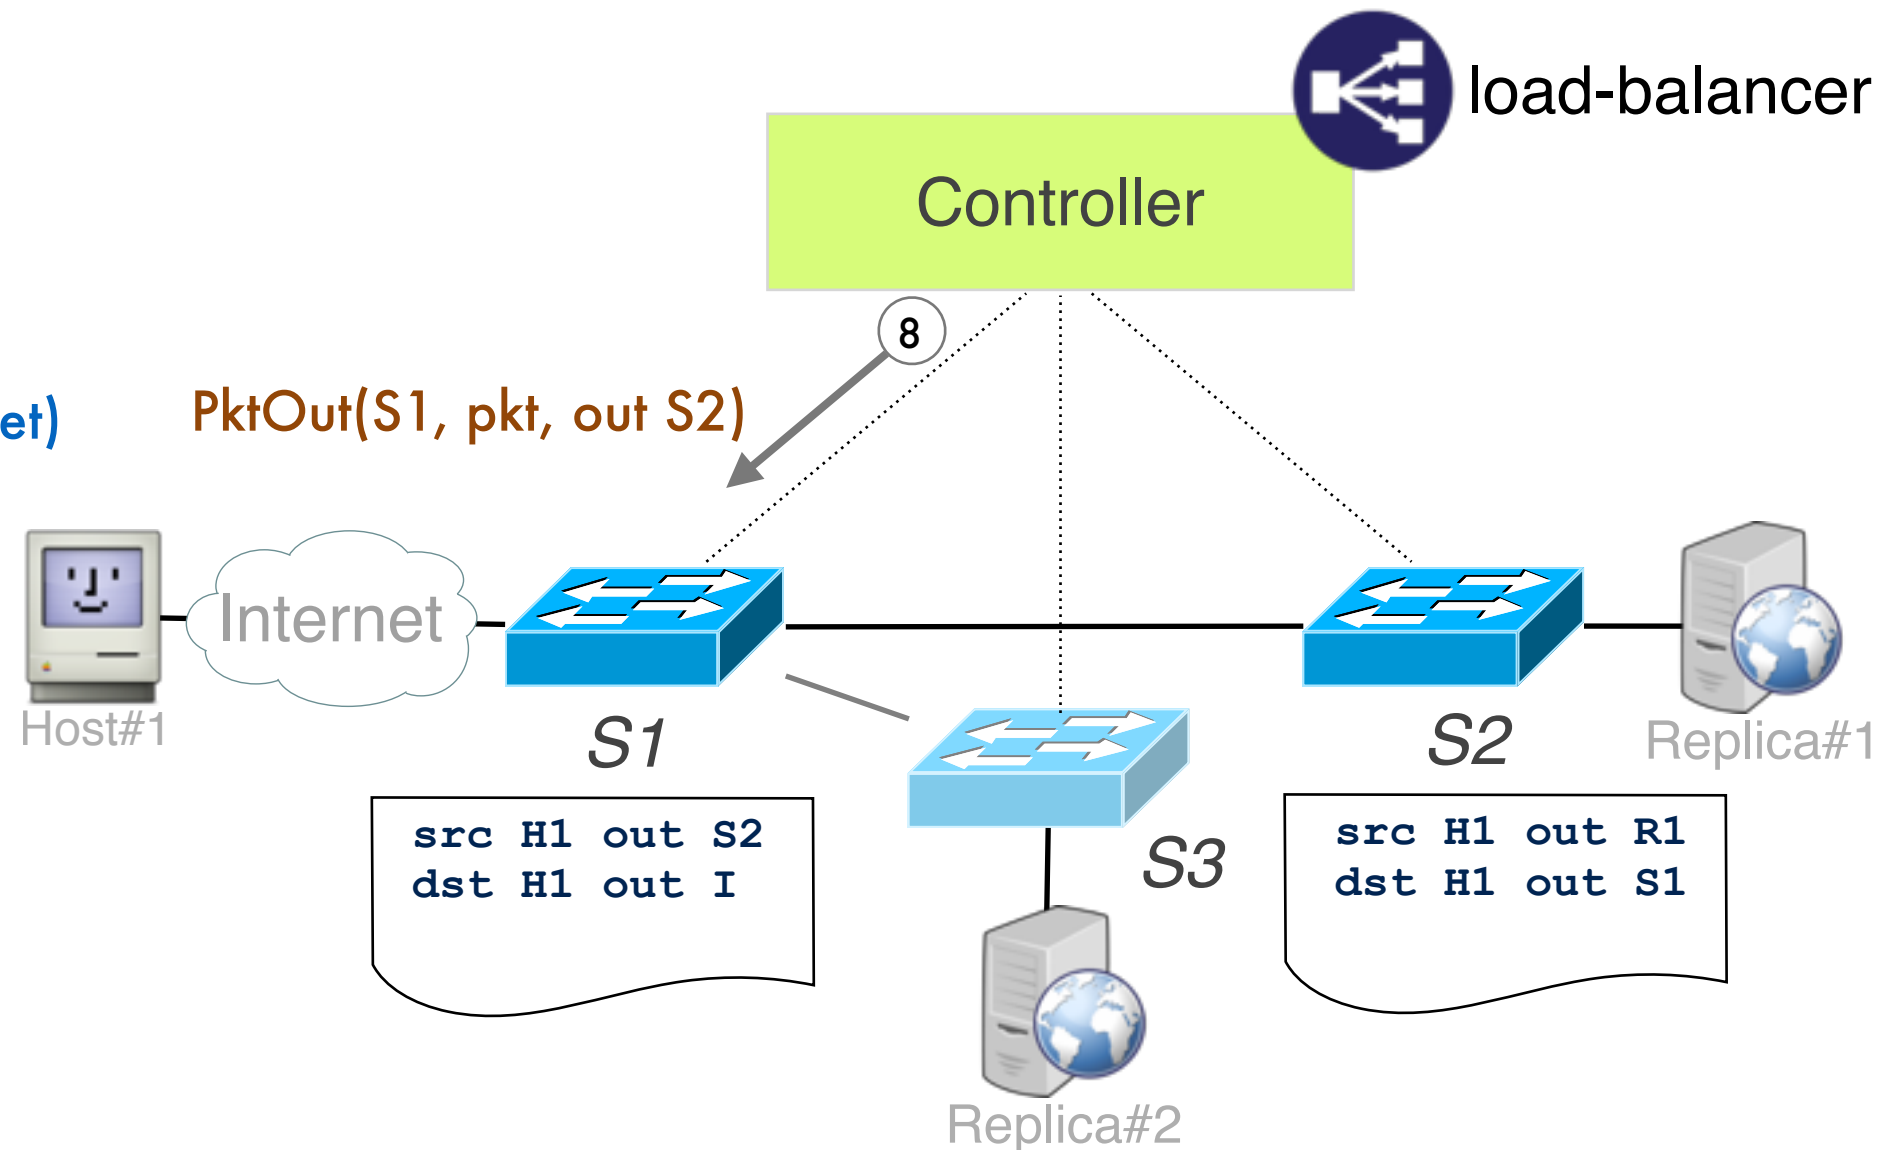
- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)



# Load Balancer Application

## Recorded Event Trace:

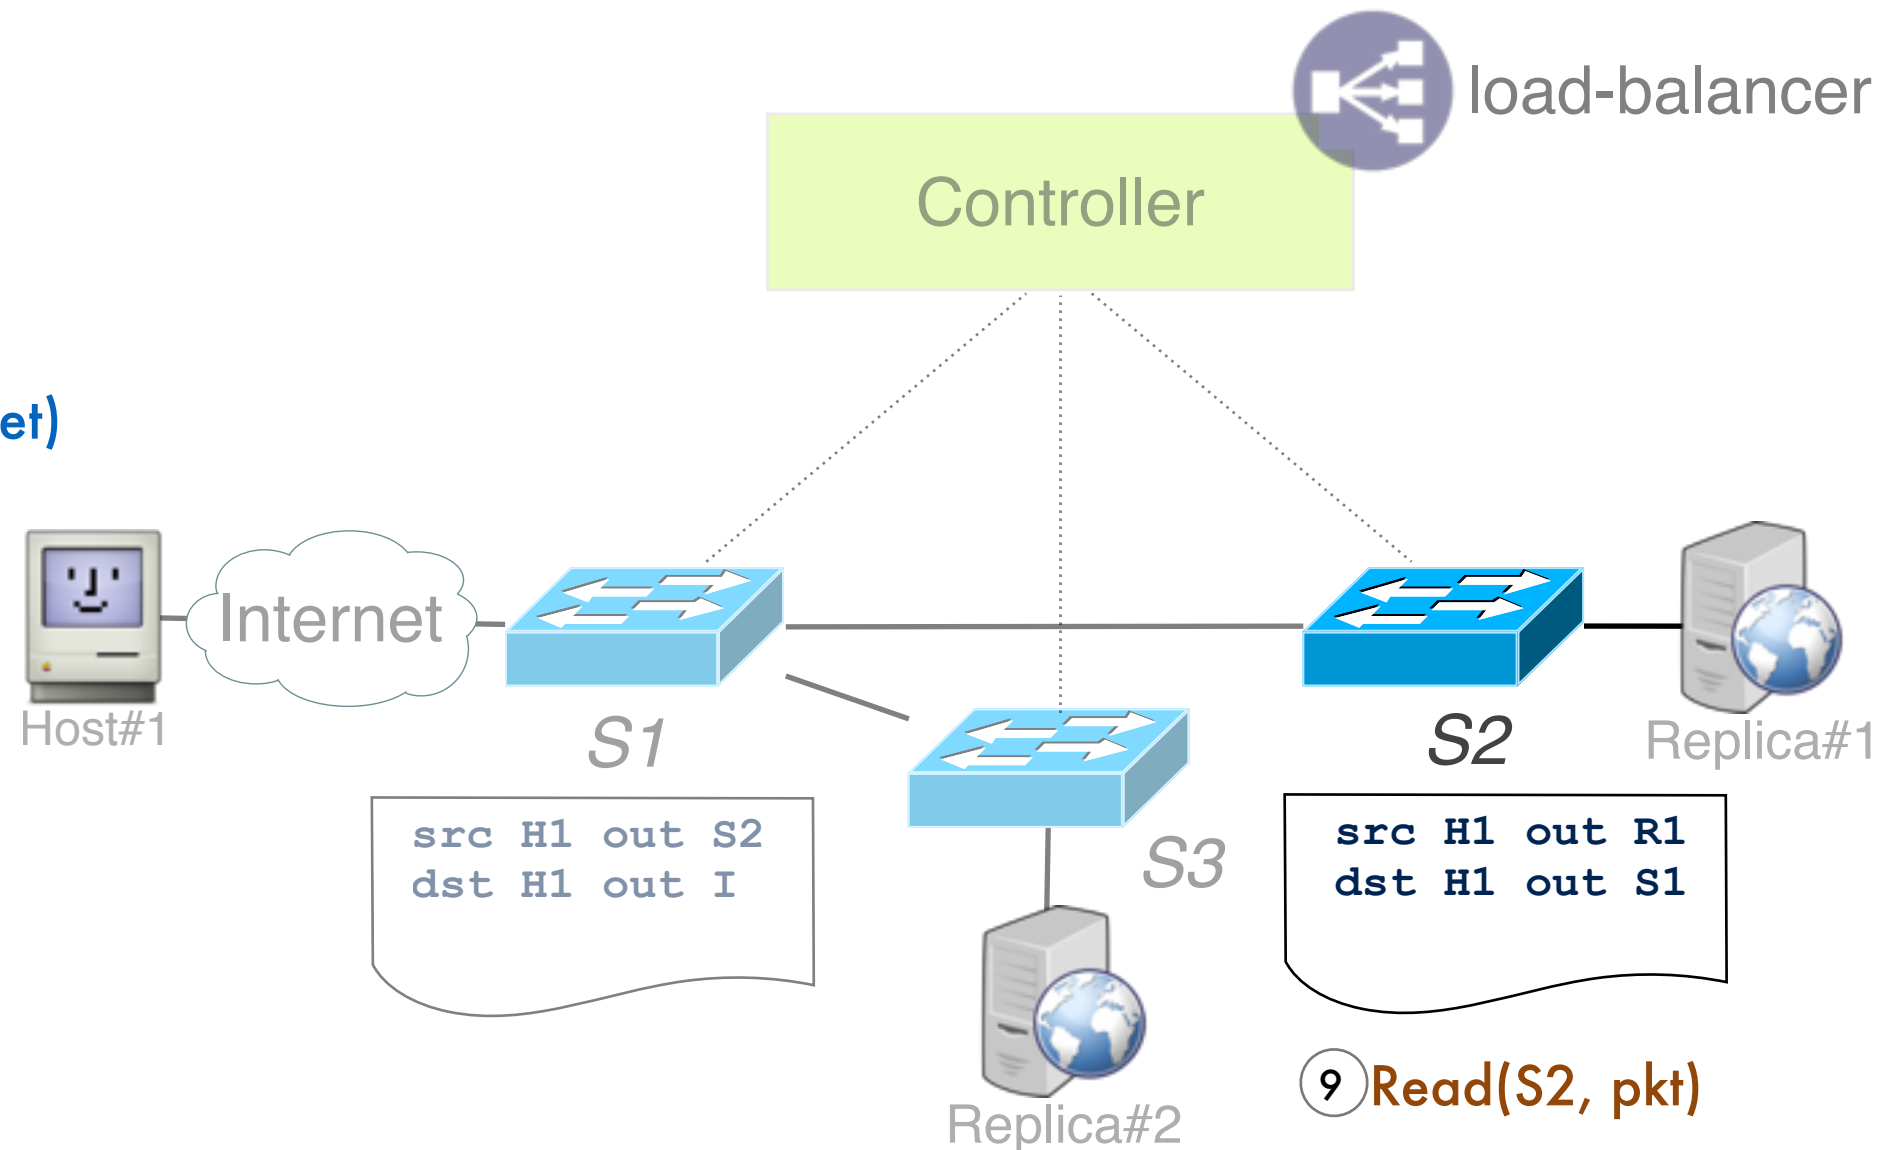
- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)



# Load Balancer Application

## Recorded Event Trace:

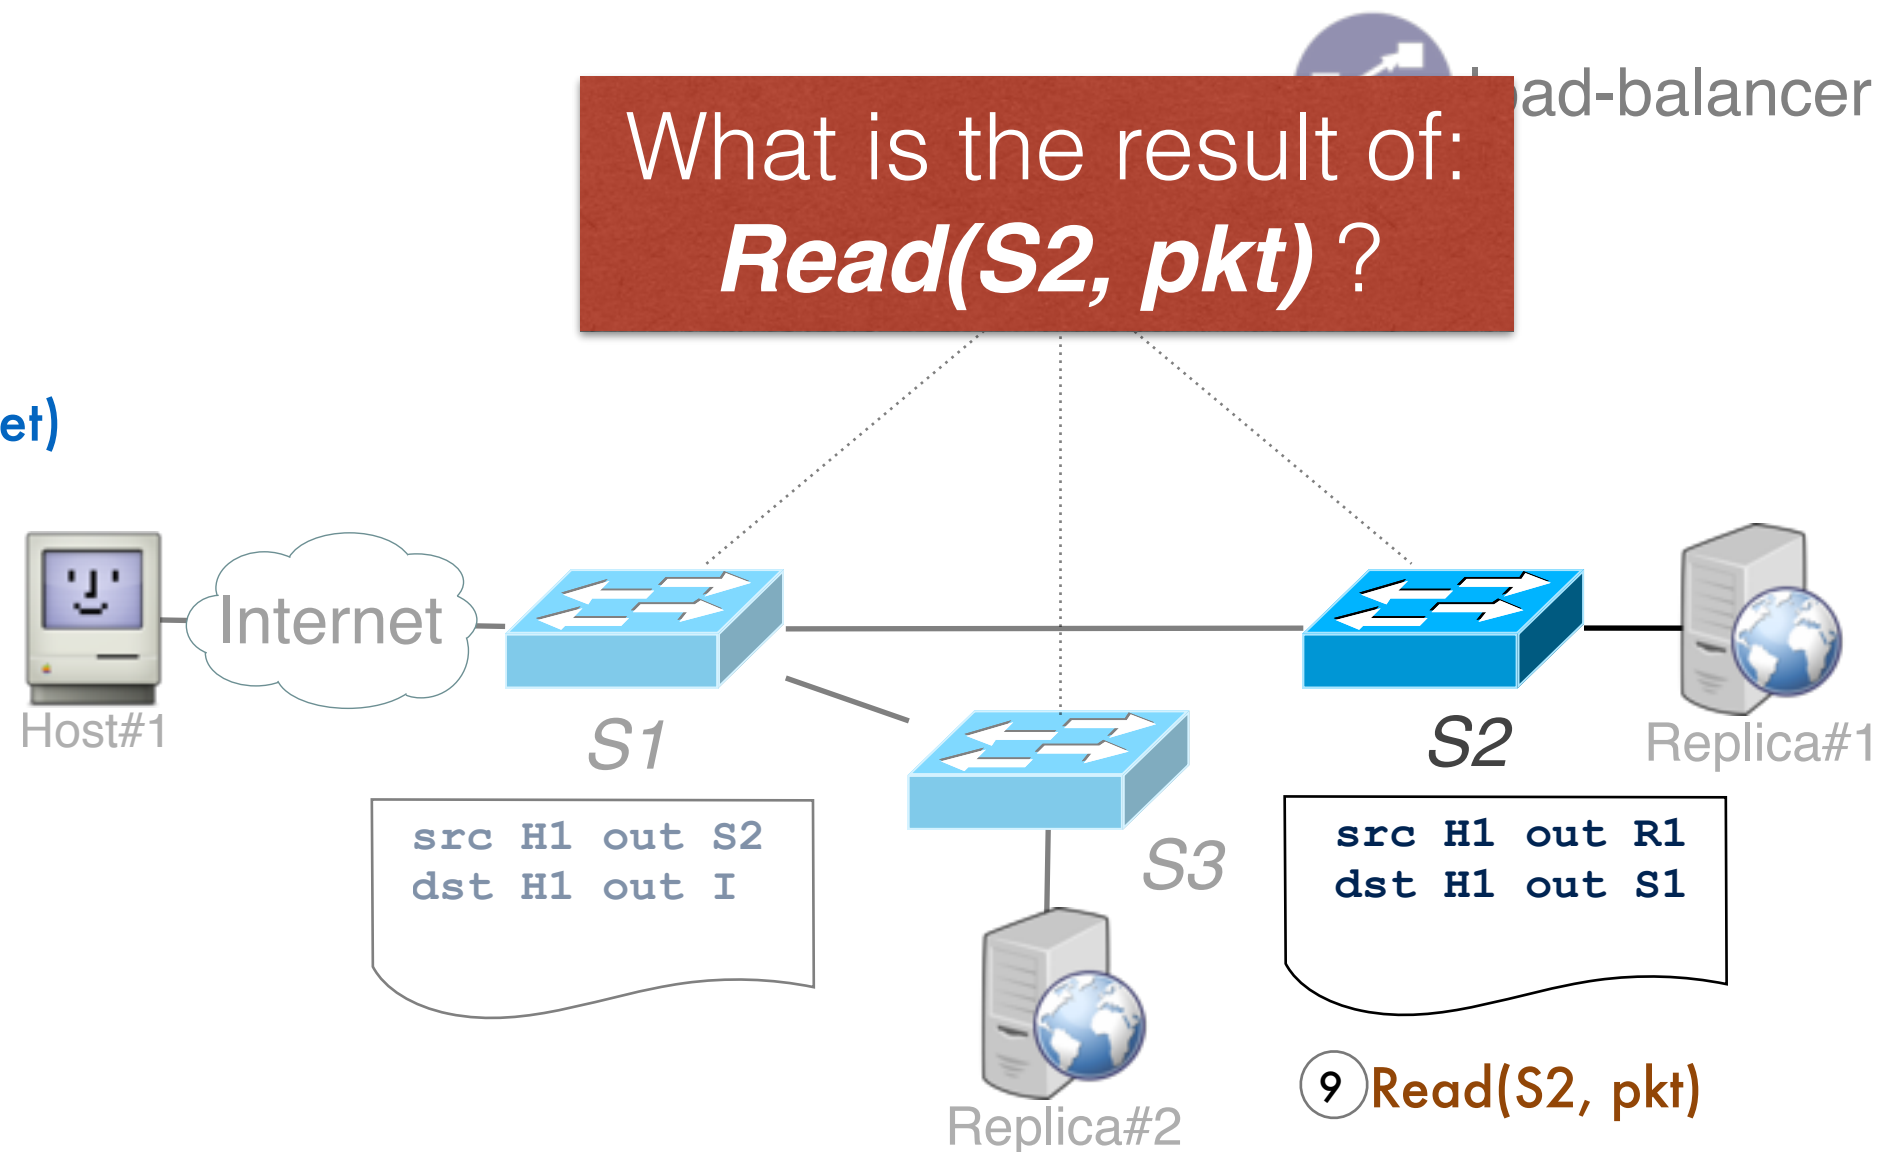
- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑨ Read(S2, pkt)



# Load Balancer Application

Recorded Event Trace:

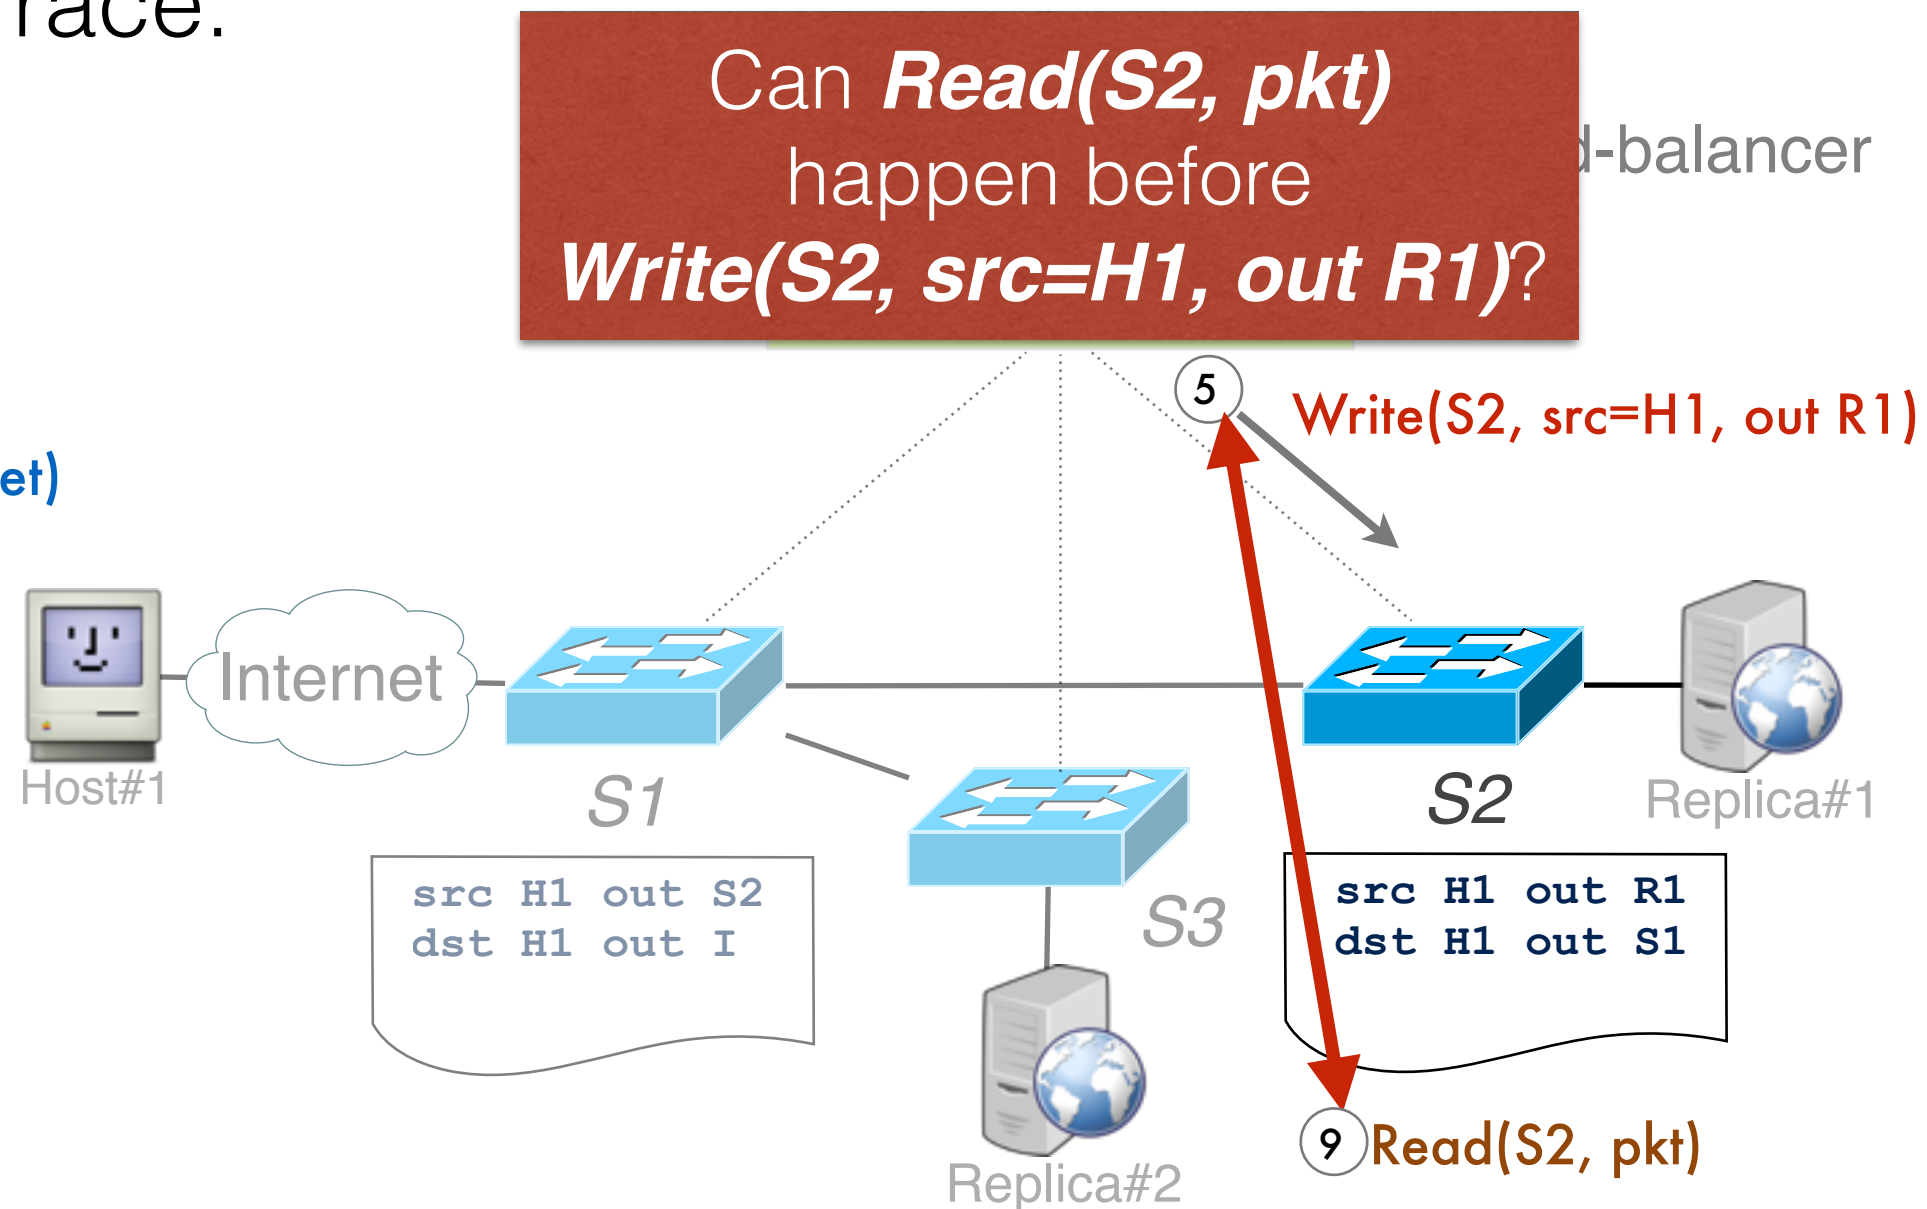
- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑨ Read(S2, pkt)



# Load Balancer Application

## Recorded Event Trace:

- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑤ **Write(S2, src=H1, out R1)**
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑨ **Read(S2, pkt)**






# Load Balancer Application

Recorded Event Trace:

- ① Host Send (pkt, dst=svr)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑤ **Write(S2, src=H1, out R1)**
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑨ **Read(S2, pkt)**

What if *Read(S2, pkt)*  
happen before  
*Write(S2, src=H1, out R1)*?



# Load Balancer Application

Reordered Event Trace:

- ① Host Send (pkt, dst=src)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑨ **Read(S2, pkt)**
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑤ **Write(S2, src=H1, out R1)**

What if *Read(S2, pkt)*  
happen before  
*Write(S2, src=H1, out R1)*?

load-balancer



Internet



S1

src H1 out S2  
dst H1 out I



S3



Replica#2



S2



Replica#1

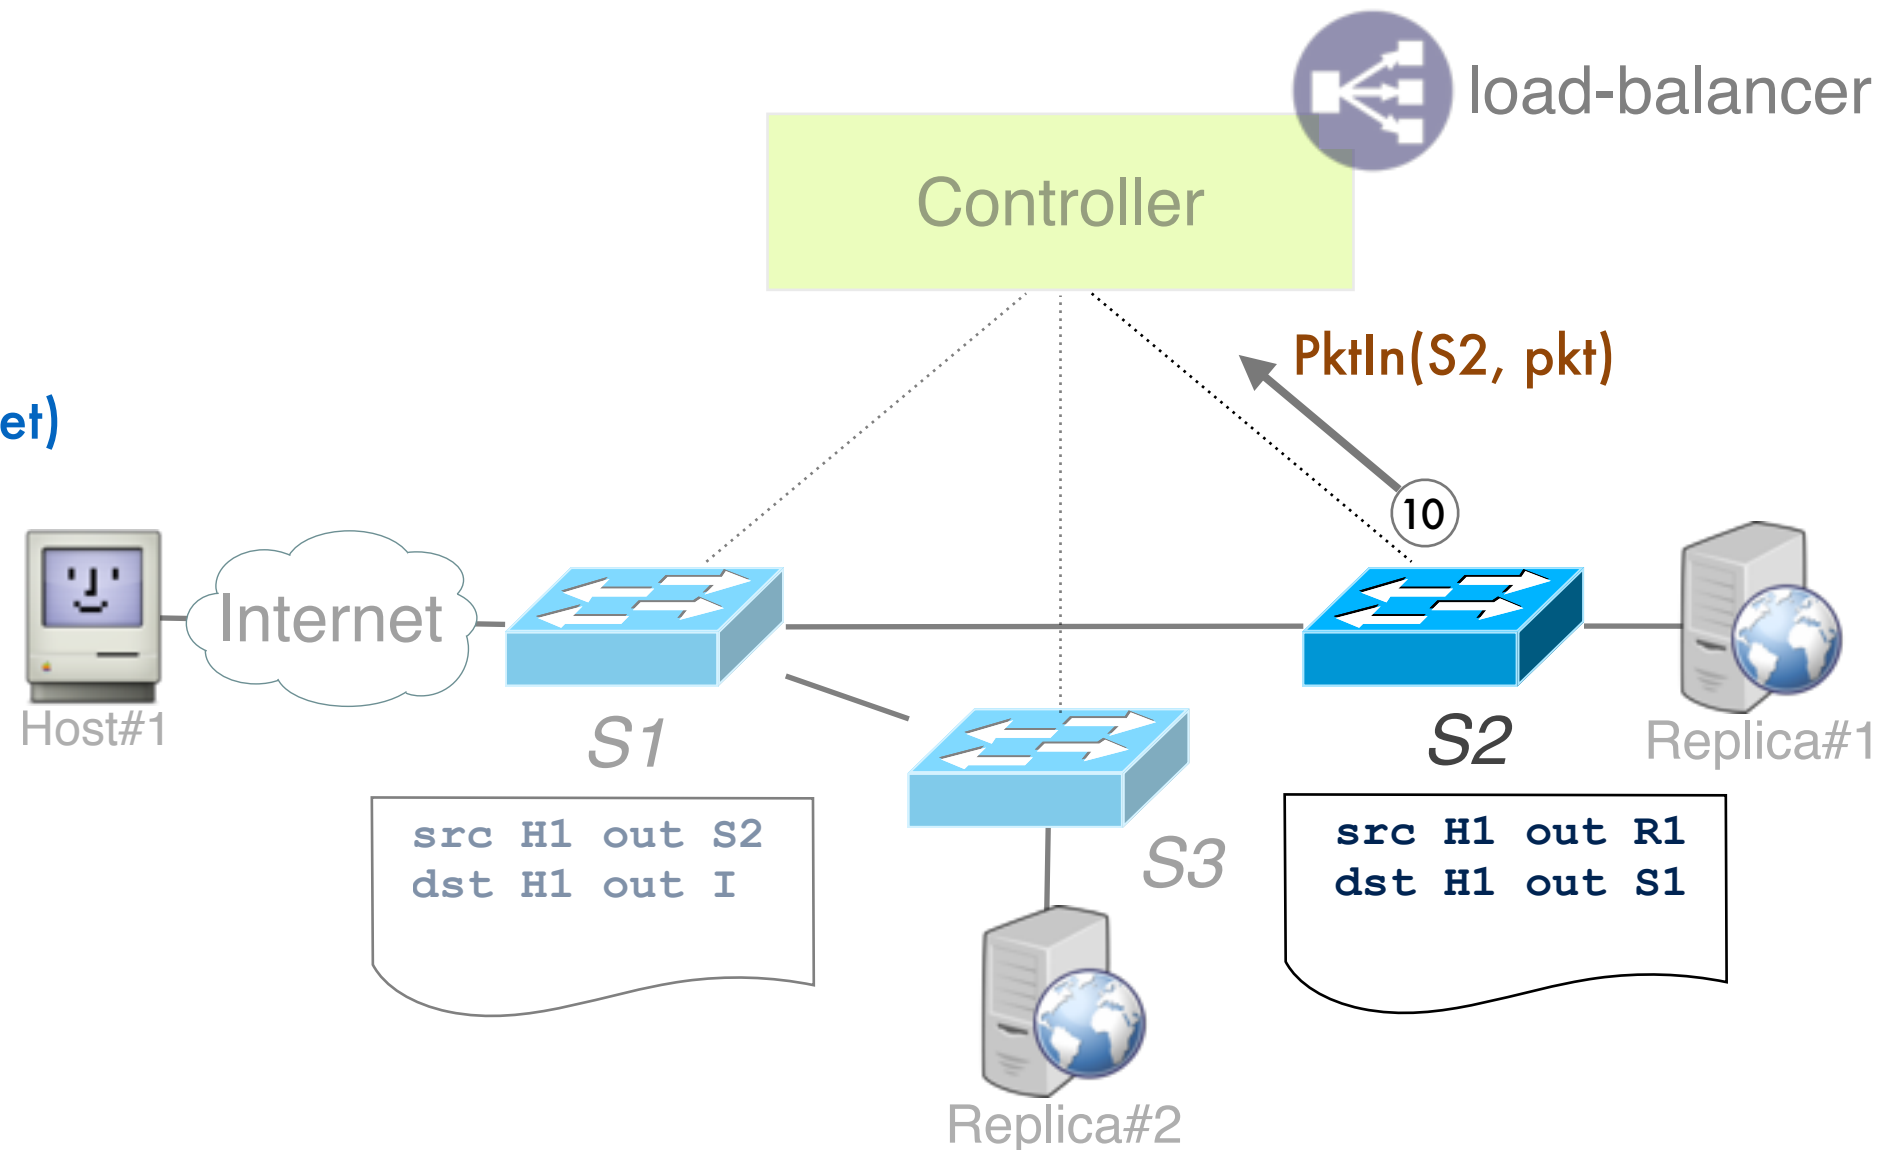
Write(S2, src=H1, out R1)

⑨ Read(S2, pkt)

# Load Balancer Application

## Reordered Event Trace:

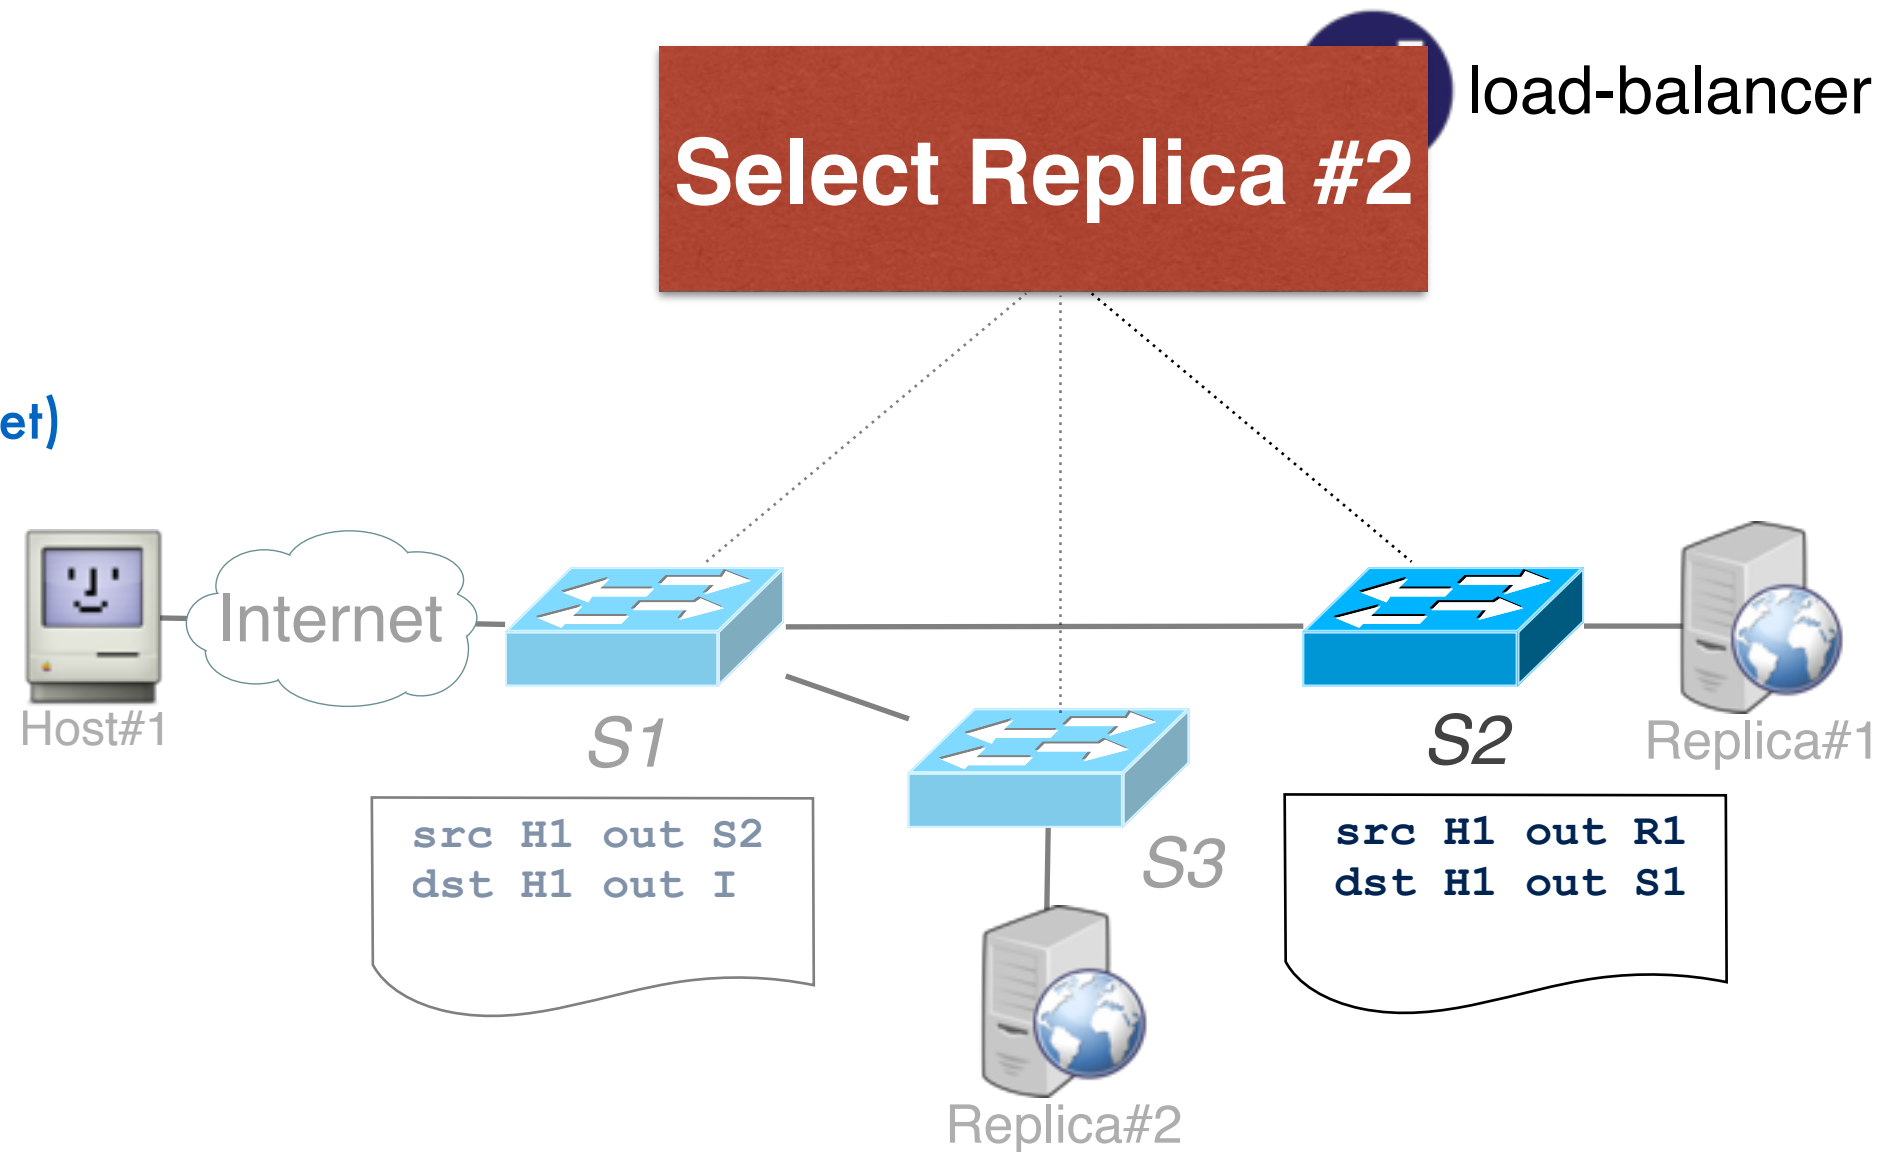
- 1 Host Send (pkt, dst=srv)
- 2 Read(S1, pkt)
- 3 PktIn(S1, pkt)
- 4 Write(S1, src=H1, out S2)
- 9 Read(S2, pkt)
- 6 Write(S1, dst=H1, out Internet)
- 7 Write(S2, dst=H1, out S1)
- 8 PktOut(S1, pkt, S2)
- 5 Write(S2, src=H1, out R1)
- 10 PktIn(S2, pkt)



# Load Balancer Application

Reordered Event Trace:

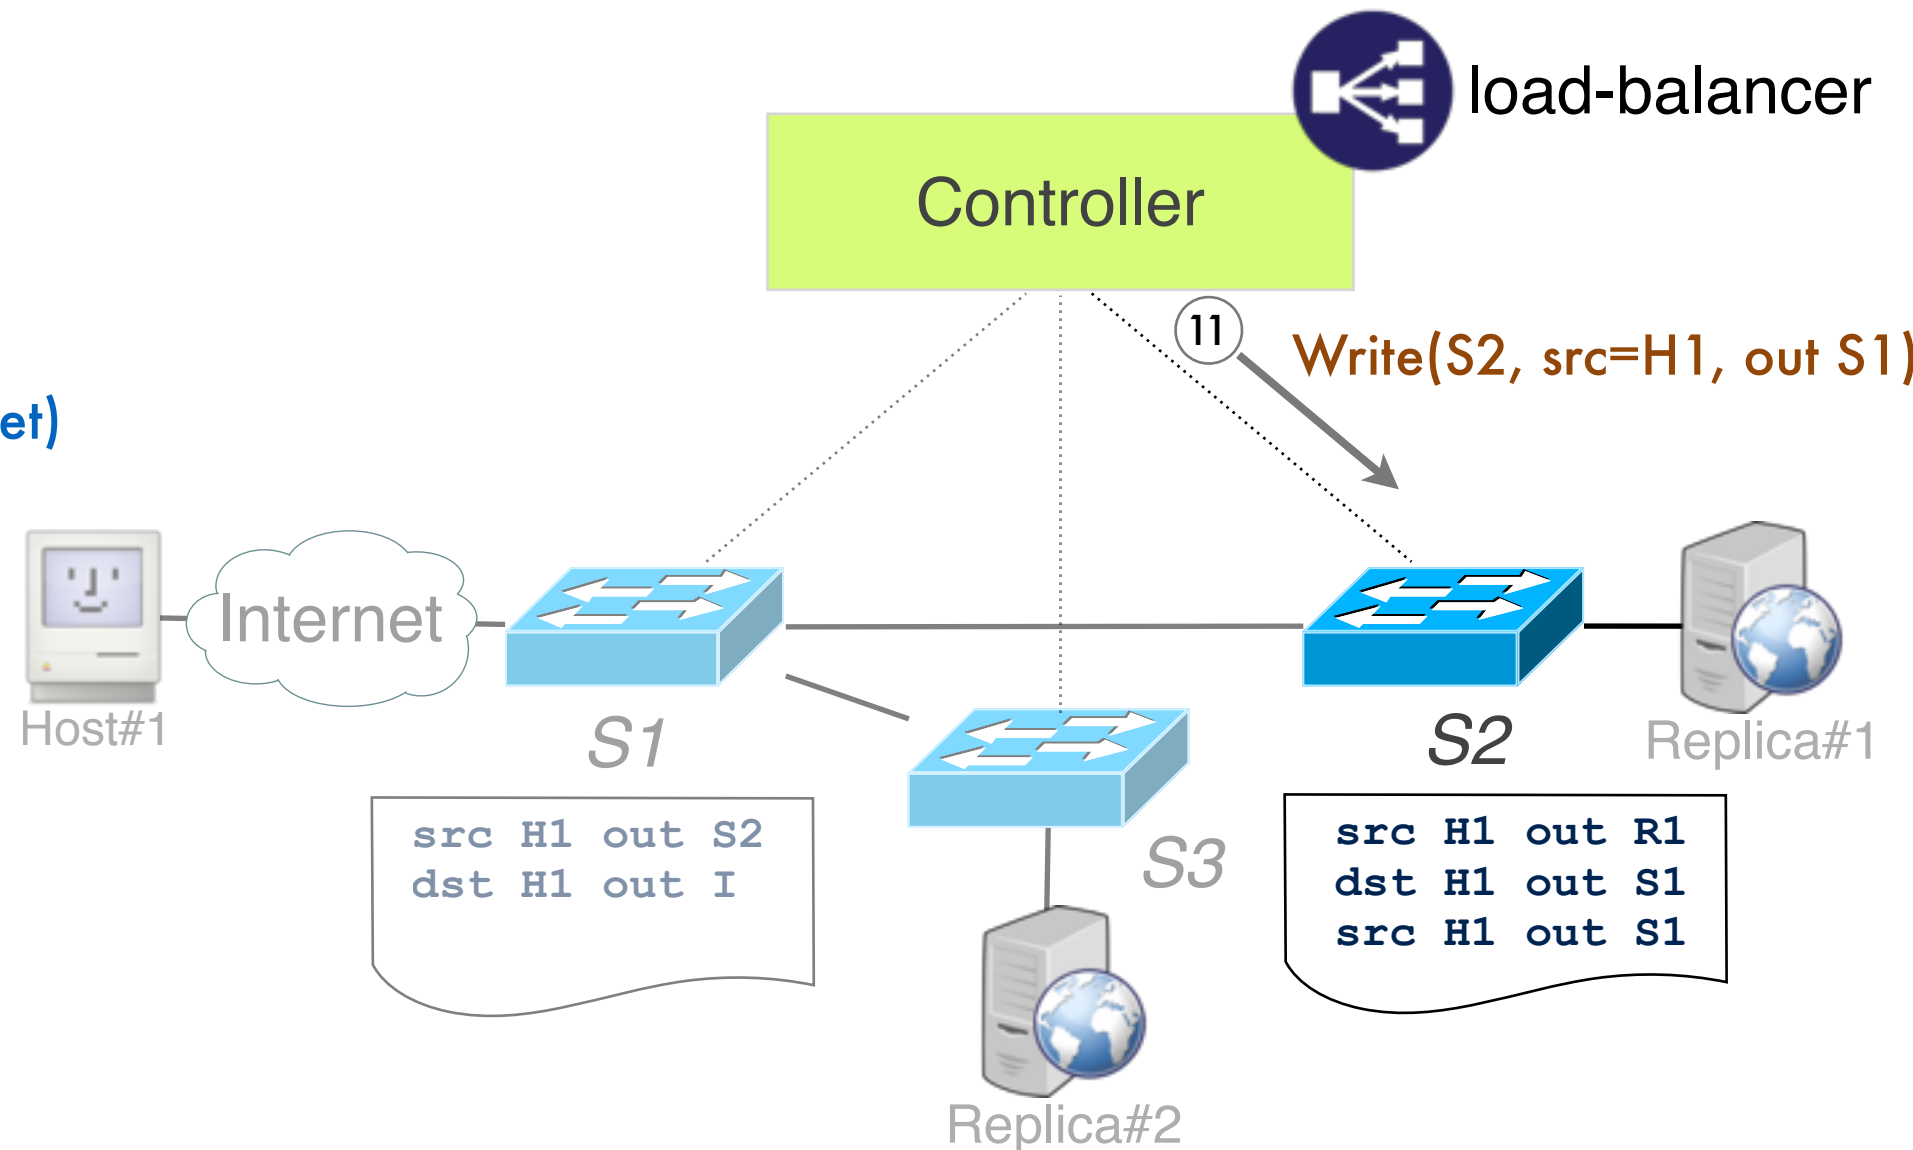
- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑨ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑩ PktIn(S2, pkt)



# Load Balancer Application

## Reordered Event Trace:

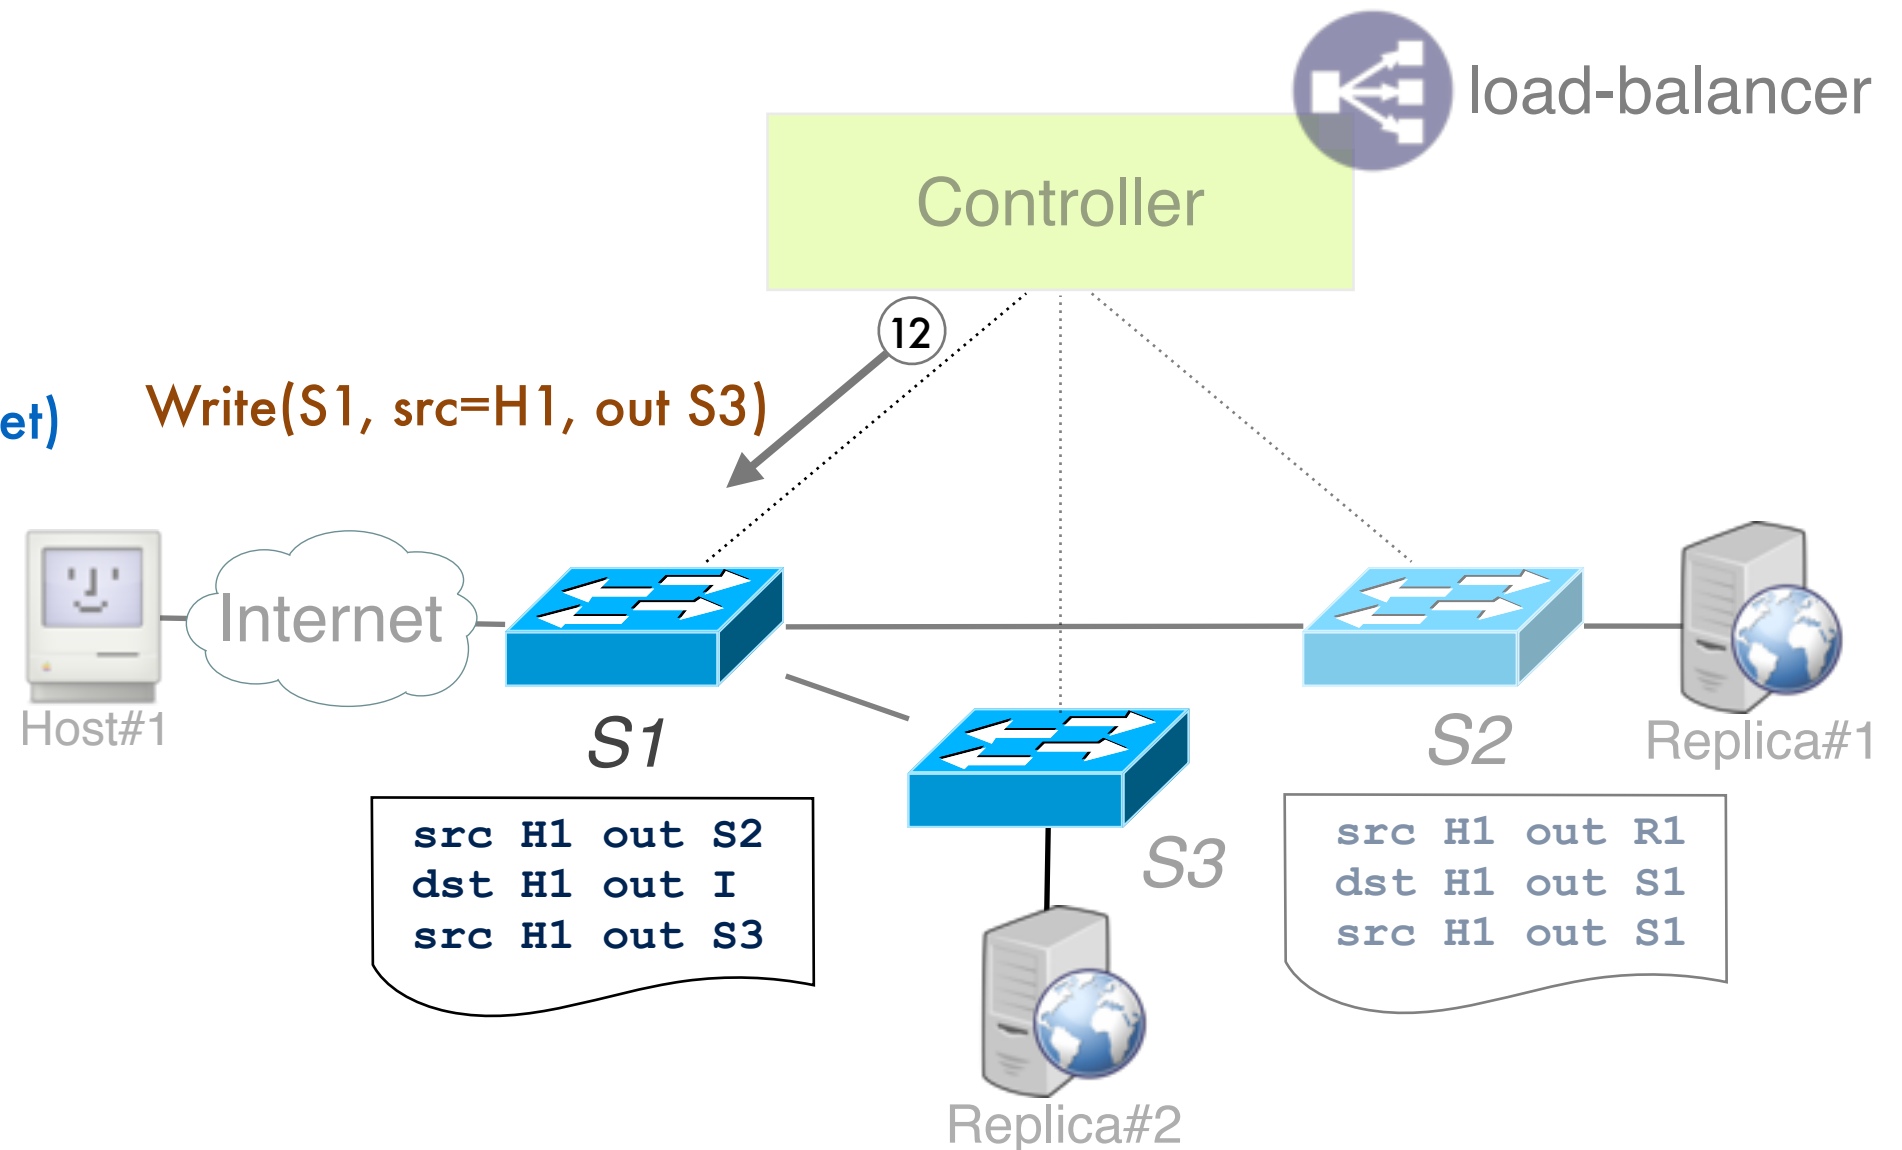
- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑨ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑩ PktIn(S2, pkt)
- ⑪ Write(S2, src=H1, out S1)



# Load Balancer Application

## Reordered Event Trace:

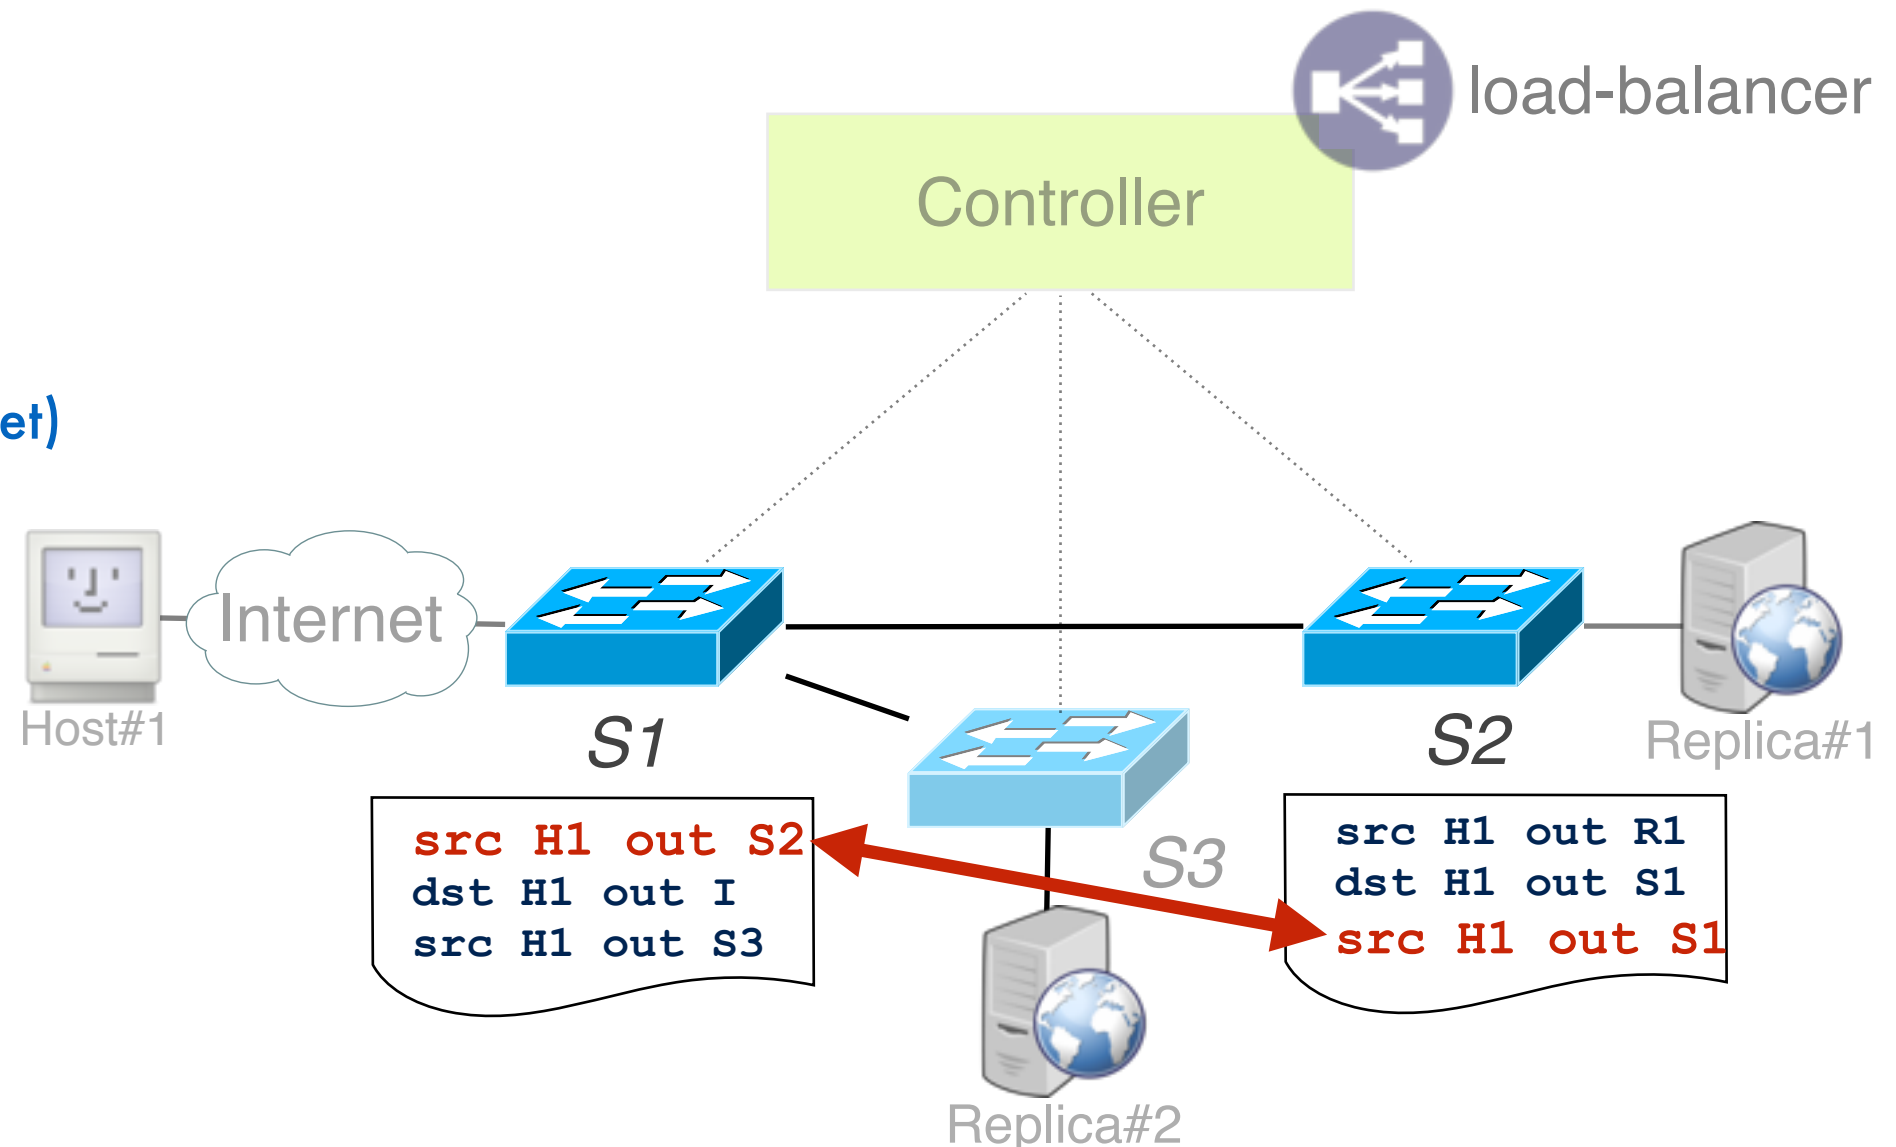
- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑨ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑩ PktIn(S2, pkt)
- ⑪ Write(S2, src=H1, out S1)
- ⑫ Write(S1, src=H1, out S3)



# Load Balancer Application

## Reordered Event Trace:

- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑨ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑩ PktIn(S2, pkt)
- ⑪ Write(S2, src=H1, out S1)
- ⑫ Write(S1, src=H1, out S3)



**Forwarding Loop!!!**

**SDNRacer** detected this **real bug** in  
**Floodlight's** Load Balancer.





# Key Observation

The cause of this bug is **interference** on the **Flow Table** caused by **concurrent writes** by the controller and **reads** triggered by packets in the network.



# SDNRacer



**Detecting concurrency violations**



Precise notion of interference



Checks for high-level properties



Implementation and evaluation


# Data Races in SDN

**Data Race:** two *unordered* events accessing the Flow Table where one is a write.

- Flow tables are memory locations.
- The controller generate writes events.
- Packets trigger read events.

# Formalizing Asynchrony in SDN

Need to identify *causality* between events.



# Happens-Before for SDN

- A switch may send a *PktIn* message after reading it

**$\text{Read}(s1, \text{pkt}) \rightarrow \text{PktIn}(s1, \text{pkt})$**

- A controller may issue *Write* after *PktIn*

**$\text{PktIn}(s1, \text{pkt}) \rightarrow \text{write}(s1, \text{match predicate, out S2})$**

**$\text{PktIn}(s1, \text{pkt}) \rightarrow \text{write}(s2, \text{match predicate, out R1})$**

# HB-relation example


## Events Trace:

- ① Host Send (pkt, dst=src)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑨ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑩ PktIn(S2, pkt)

# HB-relation example

## Events Trace:

- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑤ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑨ Write(S2, src=H1, out R1)
- ⑩ PktIn(S2, pkt)



# Analysis Results

For Floodlight Load Balancer

**703,864 Races**



# Analysis Results

For Floodlight Load Balancer

**703,864 Races**

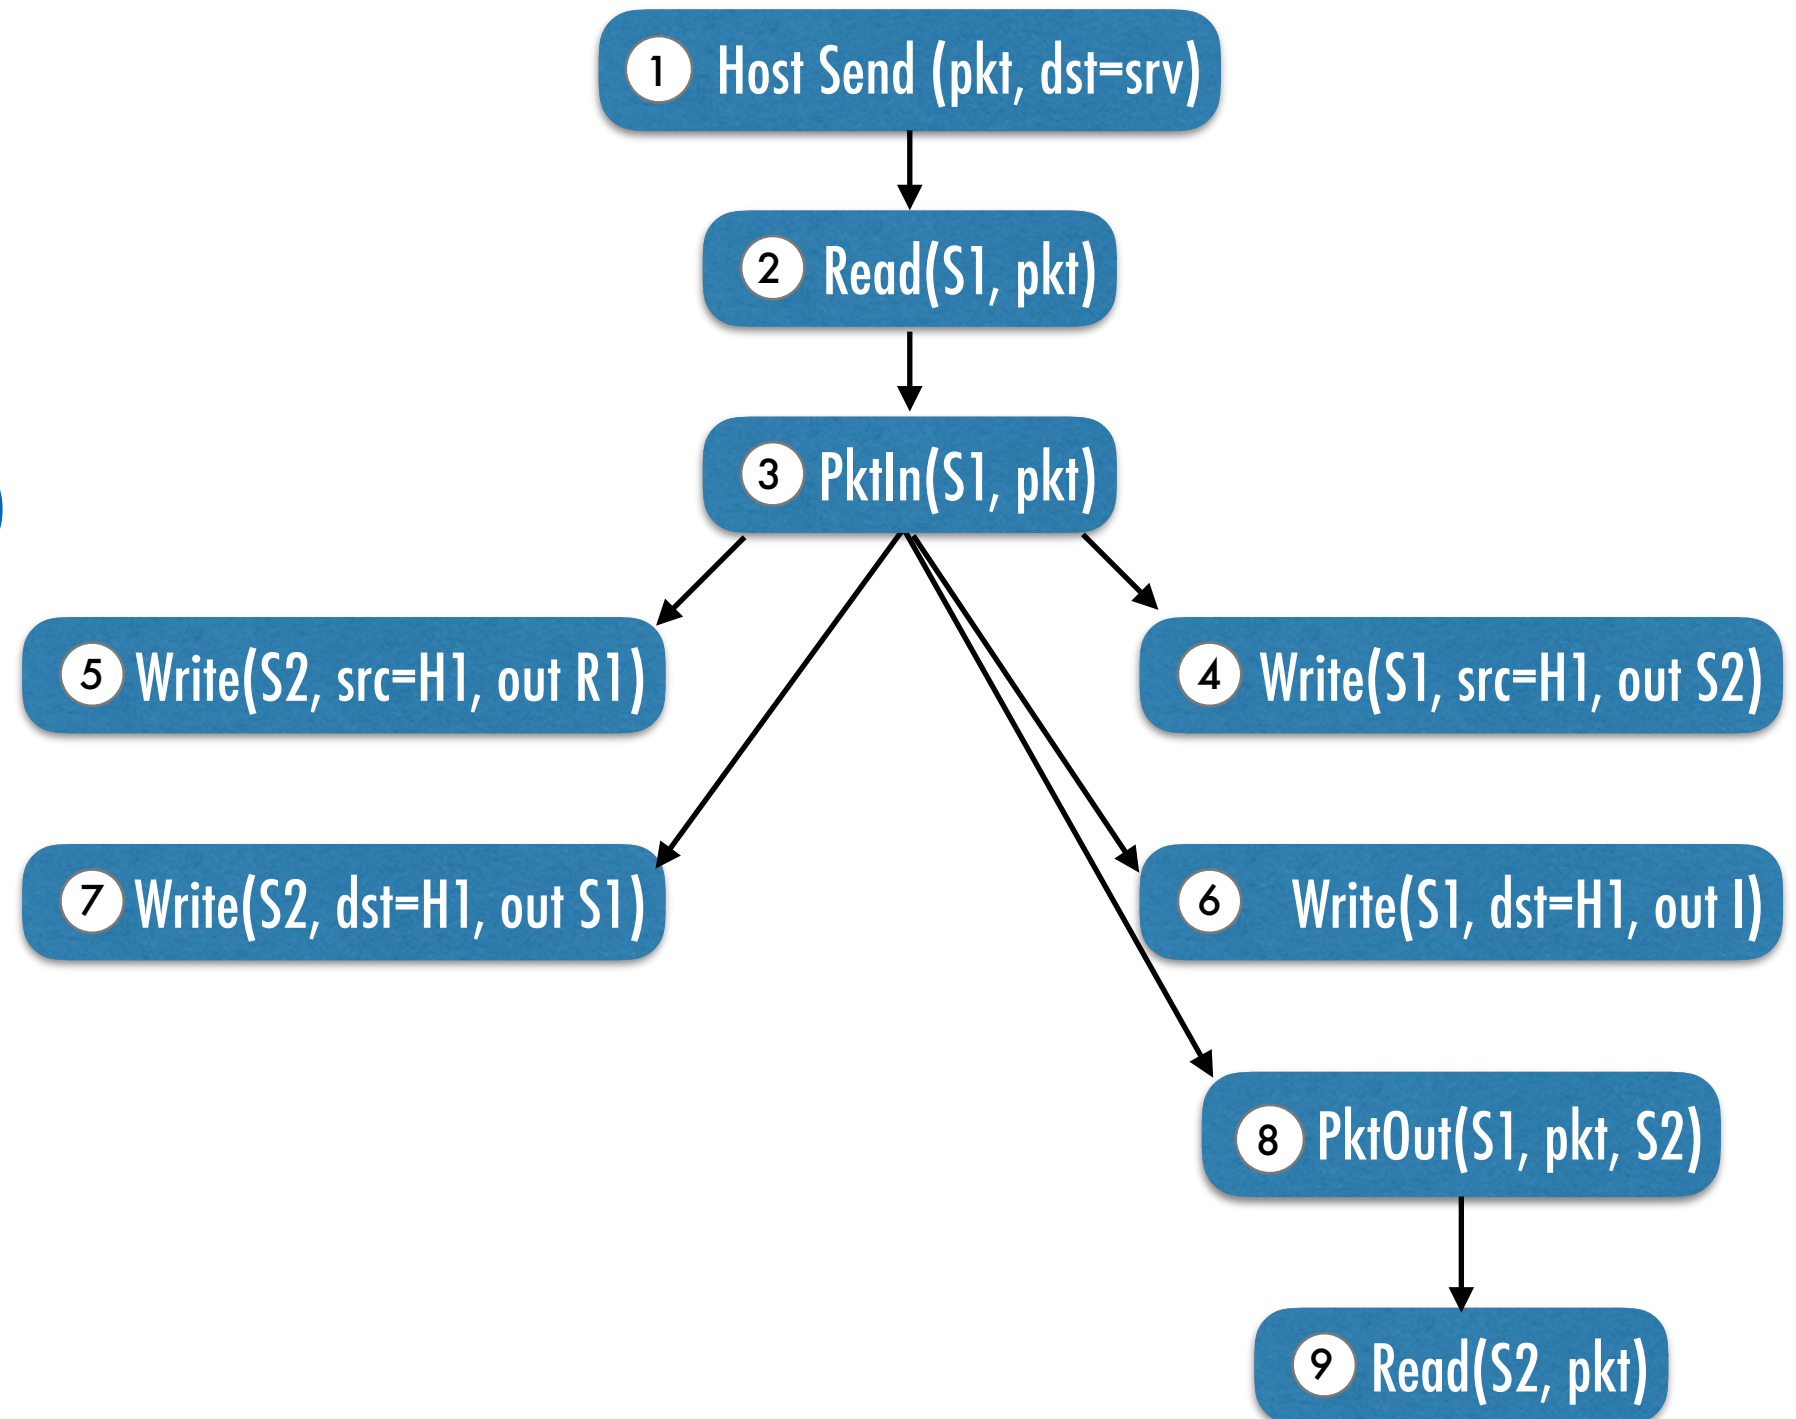
**Too imprecise!!!**

How can we **reduce**  
the number of **reports**?

# HB-relation example

## Events Trace:

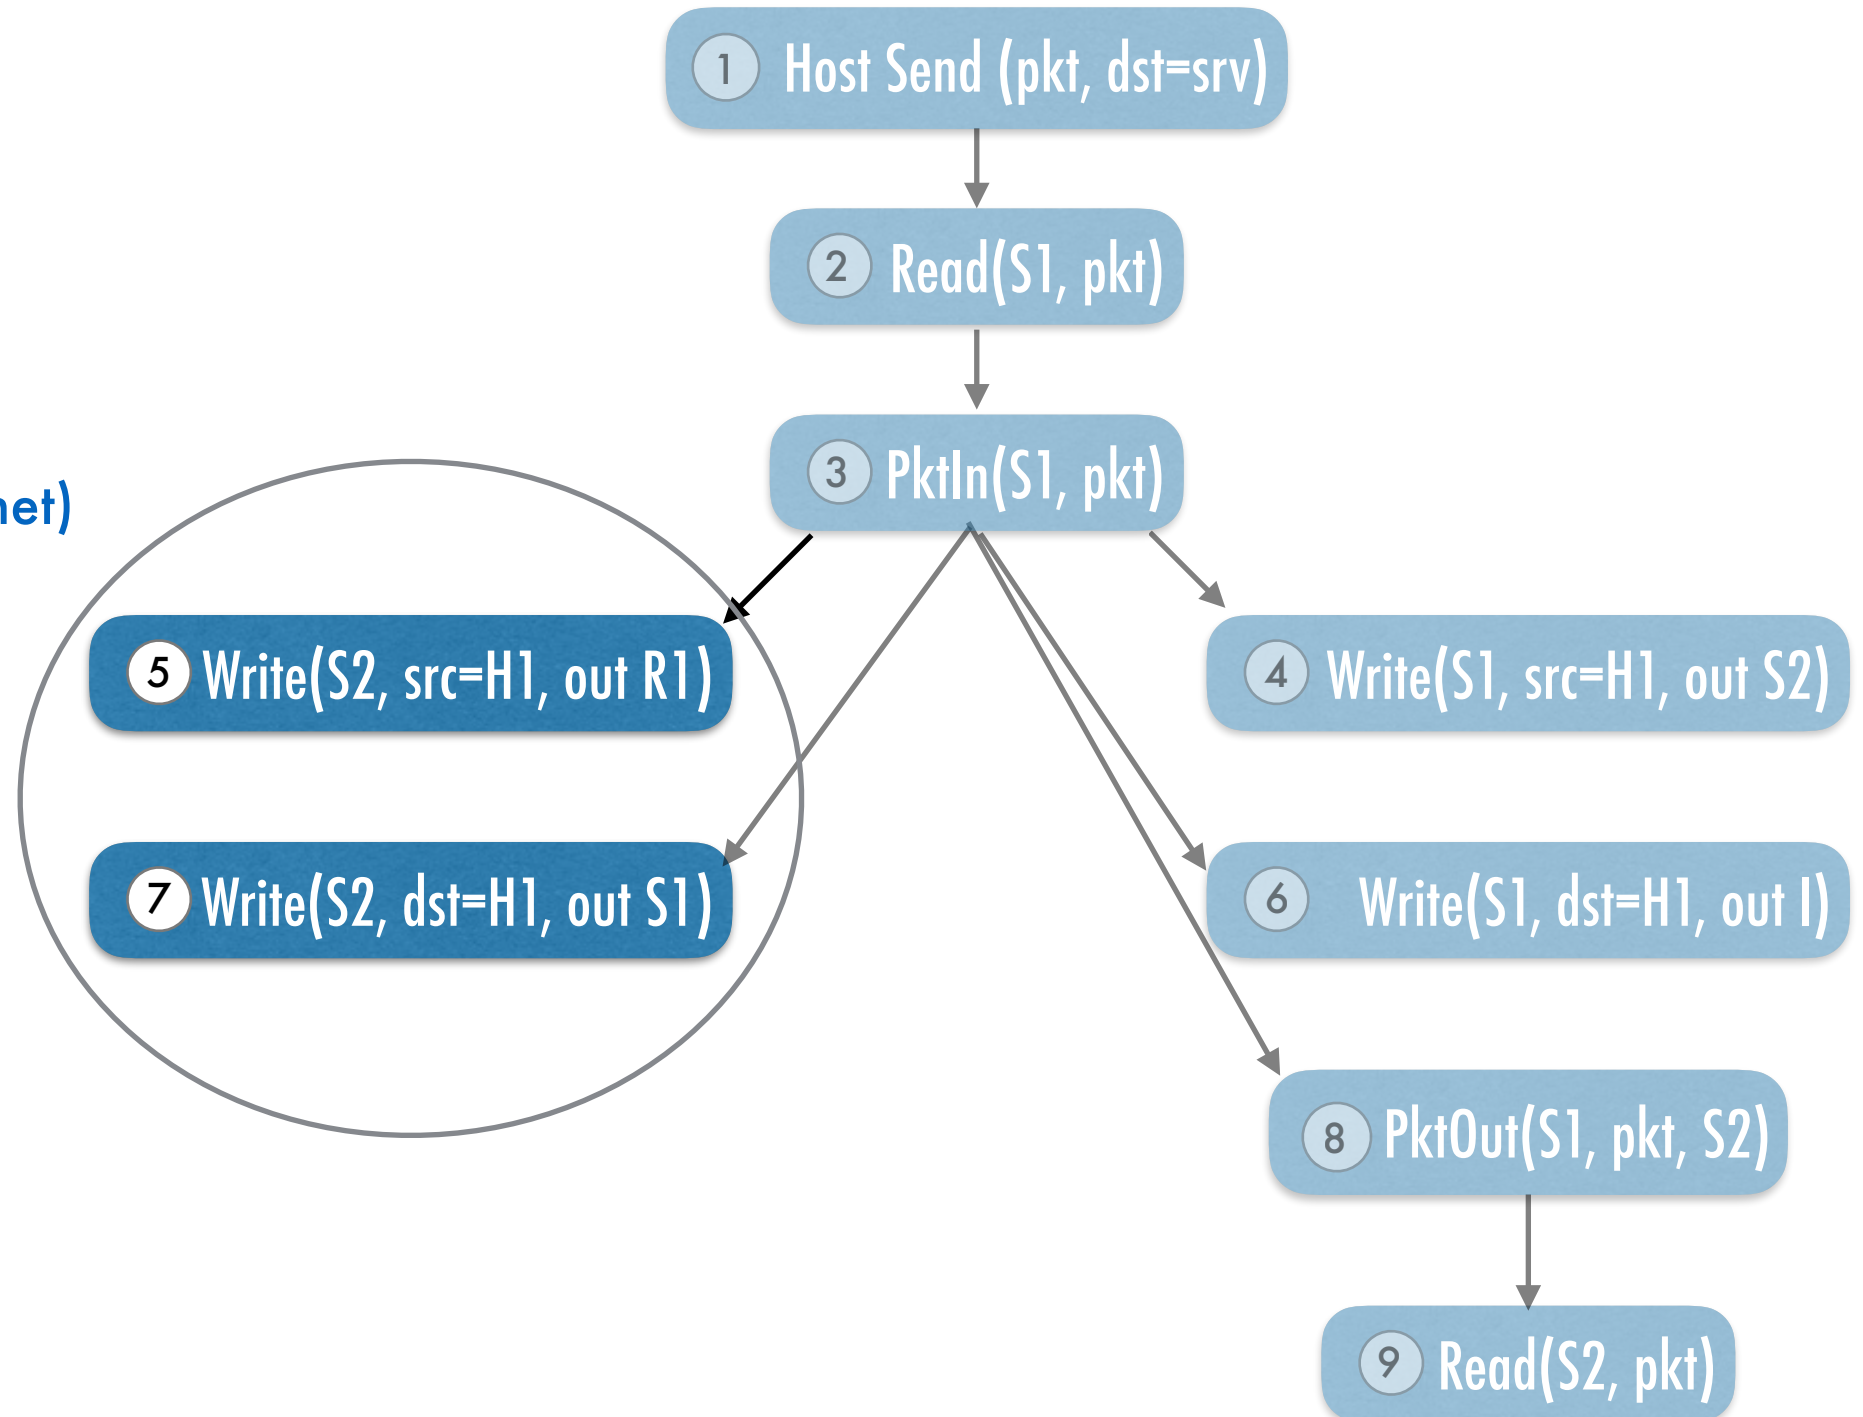
- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑤ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑨ Write(S2, src=H1, out R1)
- ⑩ PktIn(S2, pkt)



# HB-relation example

## Events Trace:

- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑨ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ **Write(S2, dst=H1, out S1)**
- ⑧ PktOut(S1, pkt, S2)
- ⑤ **Write(S2, src=H1, out R1)**
- ⑩ PktIn(S2, pkt)



# Non-Interfering Events

## Events Trace:

- ① Host Send (pkt, dst=src)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑨ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ **Write(S2, dst=H1, out S1)**
- ⑧ PktOut(S1, pkt, S2)
- ⑤ **Write(S2, src=H1, out R1)**
- ⑩ PktIn(S2, pkt)



⑤ Write(S2, src=H1, out R1)

Event ⑤ match predicate true for packets with source H1.

⑦ Write(S2, dst=H1, out S1)

Event ⑦ match predicate true for packets with destination H1.

# Non-Interfering Events

Read (s1, **dst=X**)

Write(s1, **dst=Y**, output port 1)

Write(s1, **dst=Z**, output port 2)



# Events with the same net effect

Write(s1, dst=X, **output port 1**)

Write(s1, dst=X, **output port 1**)



# SDNRacer



Detecting concurrency violations



**Precise notion of interference**



Checks for high-level properties

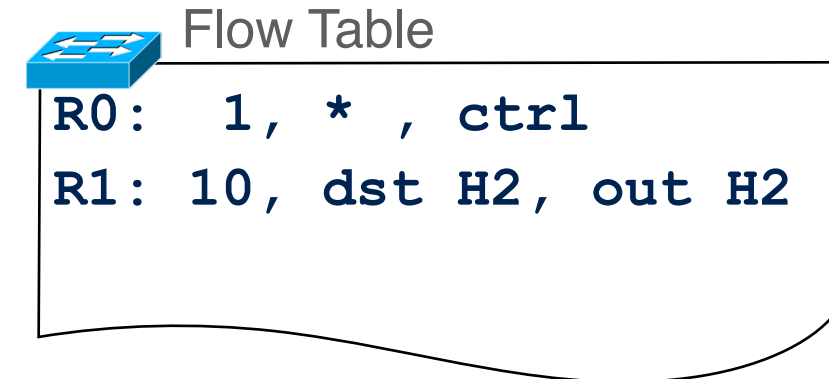


Implementation and evaluation



# Capture interference at the Flow Table

**Observation:** Flow Table can be seen as a high-level ADT with operations:



`read(pkt)/eread`

A packet is matched against entry in the flow table.

`add(eadd, no overlap)`

Add a new entry to the flow table.

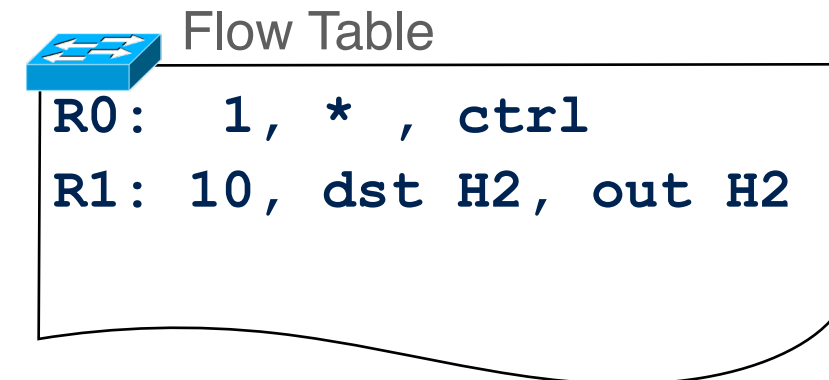
`mod(emod, strict)`

A mod operation modifies existing entries in the flow table.

`del(edel, strict)`

A del operation deletes all entries that match the entry in the flow table.

# Capture interference at the Flow Table

A diagram of a flow table. It features a small blue icon of a network switch with a grid pattern on its top surface, positioned to the left of the title "Flow Table". The table itself is a rounded rectangle with a black border, containing two entries labeled R0 and R1.

R0:	1, *	ctrl
R1:	10, dst H2,	out H2

Capture high-level interference via commutativity:  
Two operations **commute** if **reordering** them have  
the **same effect** on the the **network state**.

# Commutativity Spec of Flow Table ADT

$\varphi_{add(e_{add}, no\_overlap)}^{read(pkt)/e_{read}}$	$:= \neg(e_{read} \neq none \wedge e_{read} = e_{add})$ $:= \neg(pkt.h \subseteq e_{add}.m \wedge (e_{read} = none$ $\vee (e_{read}.p \leq e_{add}.p \wedge e_{read}.a \neq e_{add}.a)))$	if $add <_{\pi} read$ if $read <_{\pi} add$
$\varphi_{mod(e_{mod}, strict)}^{read(pkt)/e_{read}}$	$:= \neg(e_{read} \neq none \wedge e_{read} \overset{strict}{\subseteq} e_{mod} \wedge e_{read}.a = e_{mod}.a)$ $\neg(e_{read} \neq none \wedge pkt.h \subseteq e_{mod}.m \wedge e_{read}.a \neq e_{mod}.a)$	if $mod <_{\pi} read$ if $read <_{\pi} mod$
$\varphi_{del(e_{del}, strict)}^{read(pkt)/e_{read}}$	$:= \neg(pkt.h \subseteq e_{del}.m)$ $\neg(e_{read} \neq none \wedge deletes(e_{del}, e_{read}, strict))$	if $del <_{\pi} read$ if $read <_{\pi} del$
$\varphi_{mod(e_{mod}, strict_{mod})}^{del(e_{del}, strict_{del})}$	$:= \neg(deletes(e_{del}, e_{mod}, true))$ $\neg(e_{del}.m \cap e_{mod}.m \neq \emptyset)$	if $strict_{mod}$ otherwise
$\varphi_{del(e_{del}, strict)}^{add(e_{add}, no\_overlap)}$	$:= \neg(deletes(e_{del}, e_{add}, strict) \vee (no\_overlap \wedge e_{add} \cap e_{del} \neq \emptyset))$	
$\varphi_{mod(e_2, strict_2)}^{mod(e_1, strict_1)}$	$\neg(e_1.m \cap e_2.m \neq \emptyset \wedge e_1.a \neq e_2.a)$ $:= \neg(e_1.m = e_2.m \wedge e_1.p = e_2.p \wedge e_1.a \neq e_2.a)$ $\neg((e_1 \overset{strict_2}{\subseteq} e_2 \vee e_2 \overset{strict_1}{\subseteq} e_1) \wedge e_1.a \neq e_2.a)$	if $\neg strict_1 \wedge \neg strict_2$ if $strict_1 \wedge strict_2$ otherwise
$\varphi_{mod(e_{mod}, strict)}^{add(e_{add}, no\_overlap)}$	$:= \neg(e_{add} \overset{strict}{\subseteq} e_{mod} \wedge e_{add}.a \neq e_{mod}.a)$ $\neg(e_{add} \cap e_{mod} \neq \emptyset)$	if $\neg no\_overlap$ otherwise
$\varphi_{add(e_2, no\_overlap_2)}^{add(e_1, no\_overlap_1)}$	$:= \neg(e_1.m \cap e_2.m \neq \emptyset \wedge e_1.p = e_2.p)$ $\neg(e_1.m = e_2.m \wedge e_1.p = e_2.p \wedge e_1.a \neq e_2.a)$	if $no\_overlap_1 \vee no\_overlap_2$ otherwise

**Figure 3:** Commutativity specification of an OpenFlow switch. Two *read* or two *del* operations always commute.

# Commutativity Specification of SDN

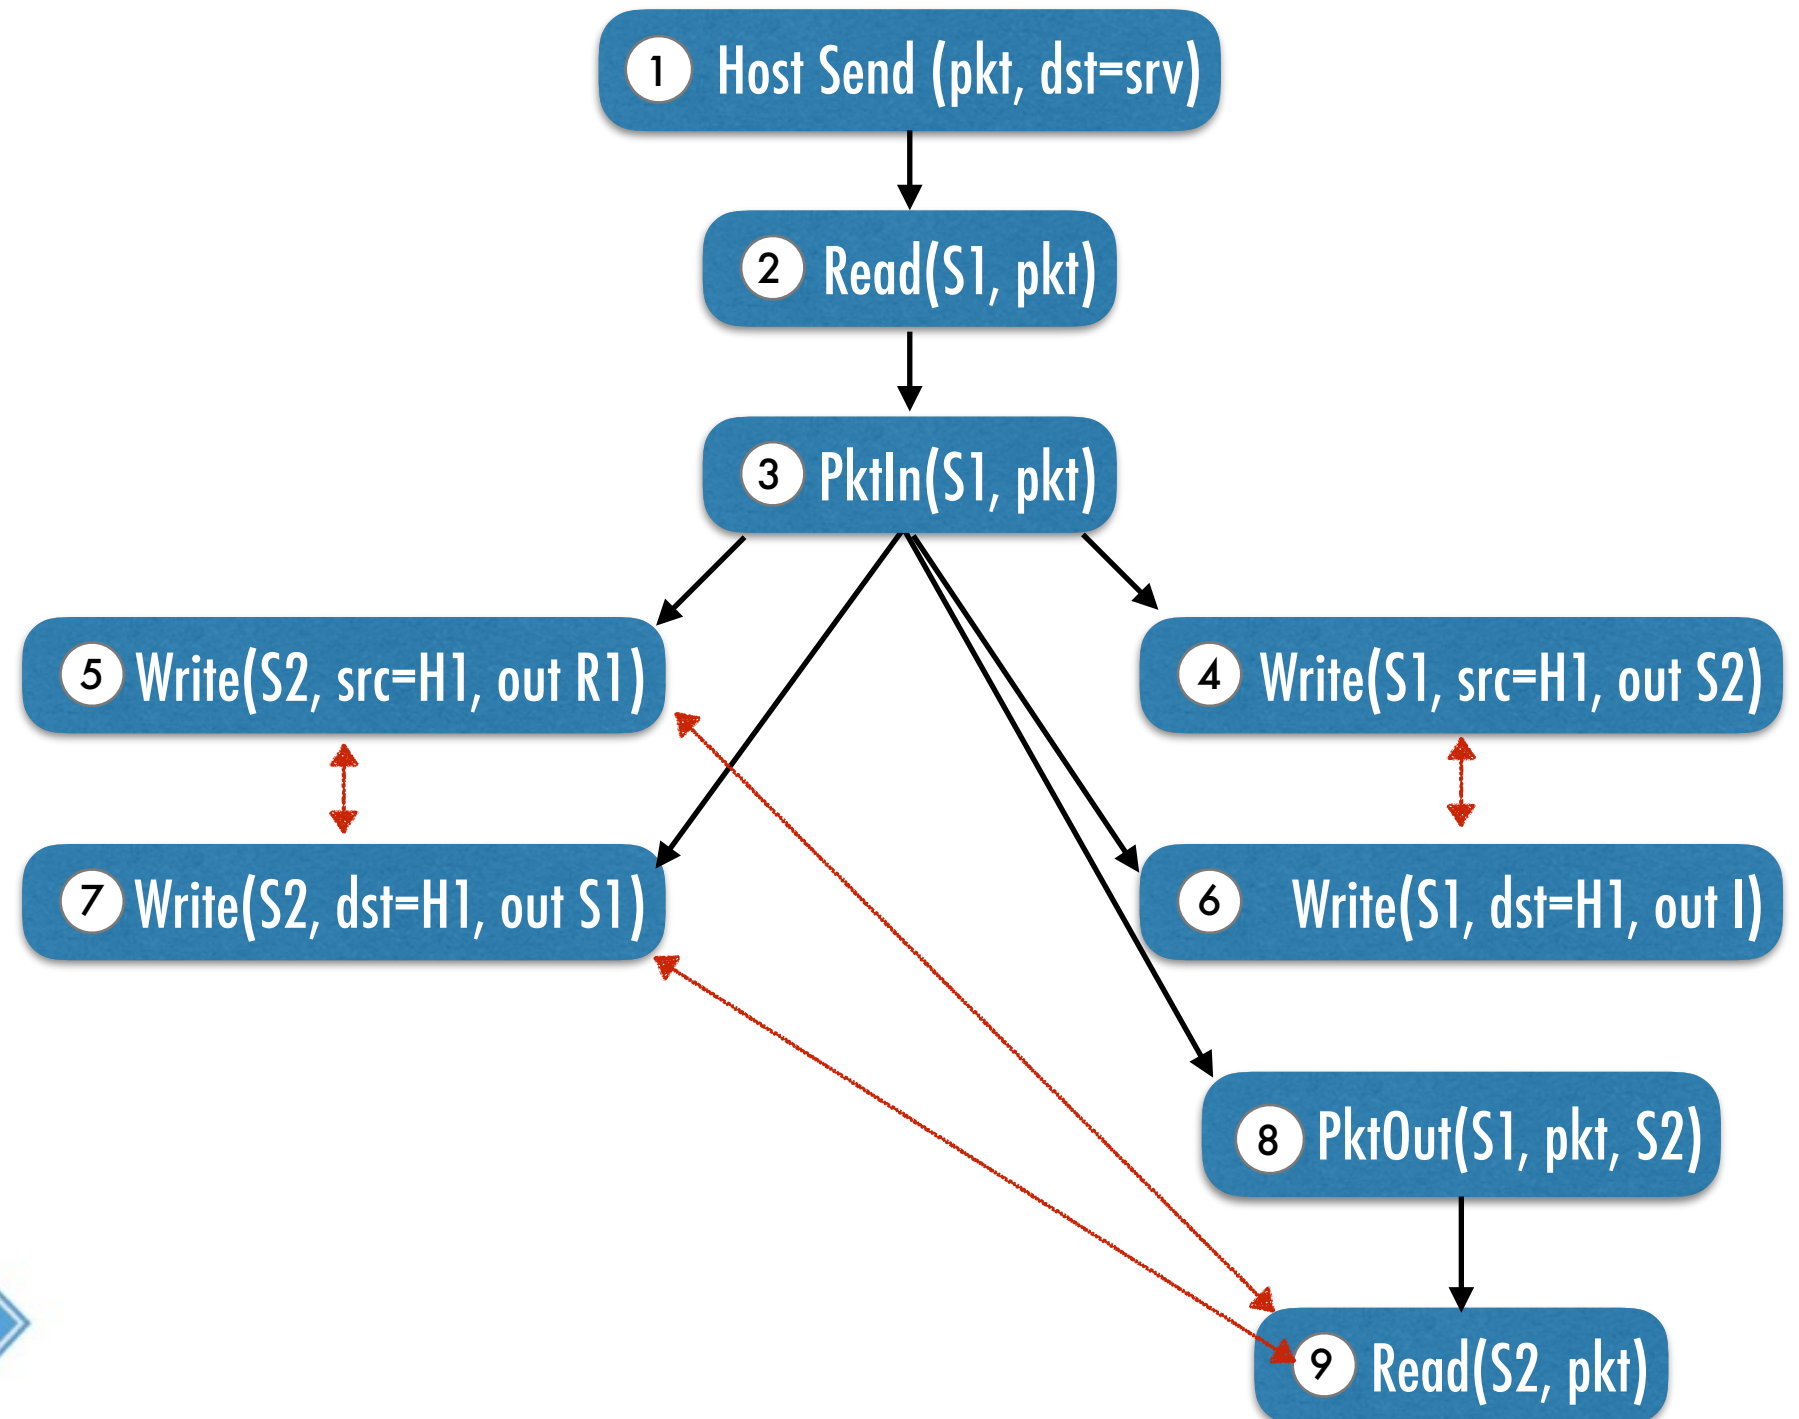
$$\begin{array}{lll}
 \varphi_{\text{add}(e_{\text{add}}, \text{no\_overlap})}^{\text{read}(pkt)/e_{\text{read}}} & := \neg(e_{\text{read}} \neq \text{none} \wedge e_{\text{read}} = e_{\text{add}}) & \text{if } \text{add} <_{\pi} \text{read} \\
 & := \neg(\text{pkt}.h \subseteq e_{\text{add}}.m \wedge (e_{\text{read}} = \text{none} & \text{if } \text{read} <_{\pi} \text{add} \\
 & \quad \vee (e_{\text{read}}.p \leq e_{\text{add}}.p \wedge e_{\text{read}}.a \neq e_{\text{add}}.a))) \\
 \\
 \varphi_{\text{mod}(e_{\text{mod}}, \text{strict})}^{\text{read}(pkt)/e_{\text{read}}} & := \neg(e_{\text{read}} \neq \text{none} \wedge e_{\text{read}} \stackrel{\text{strict}}{\subseteq} e_{\text{mod}} \wedge e_{\text{read}}.a = e_{\text{mod}}.a) & \text{if } \text{mod} <_{\pi} \text{read} \\
 & \quad \neg(e_{\text{read}} \neq \text{none} \wedge \text{pkt}.h \subseteq e_{\text{mod}}.m \wedge e_{\text{read}}.a \neq e_{\text{mod}}.a) & \text{if } \text{read} <_{\pi} \text{mod} \\
 \\
 \varphi_{\text{del}(e_{\text{del}}, \text{strict})}^{\text{read}(pkt)/e_{\text{read}}} & := \neg(\text{pkt}.h \subseteq e_{\text{del}}.m) & \text{if } \text{del} <_{\pi} \text{read}
 \end{array}$$

$$\varphi_{\text{add}(e_1, \text{no\_overlap}_1)}^{\text{add}(e_2, \text{no\_overlap}_2)} := \neg(e_1.m \cap e_2.m \neq \emptyset \wedge e_1.p = e_2.p)$$

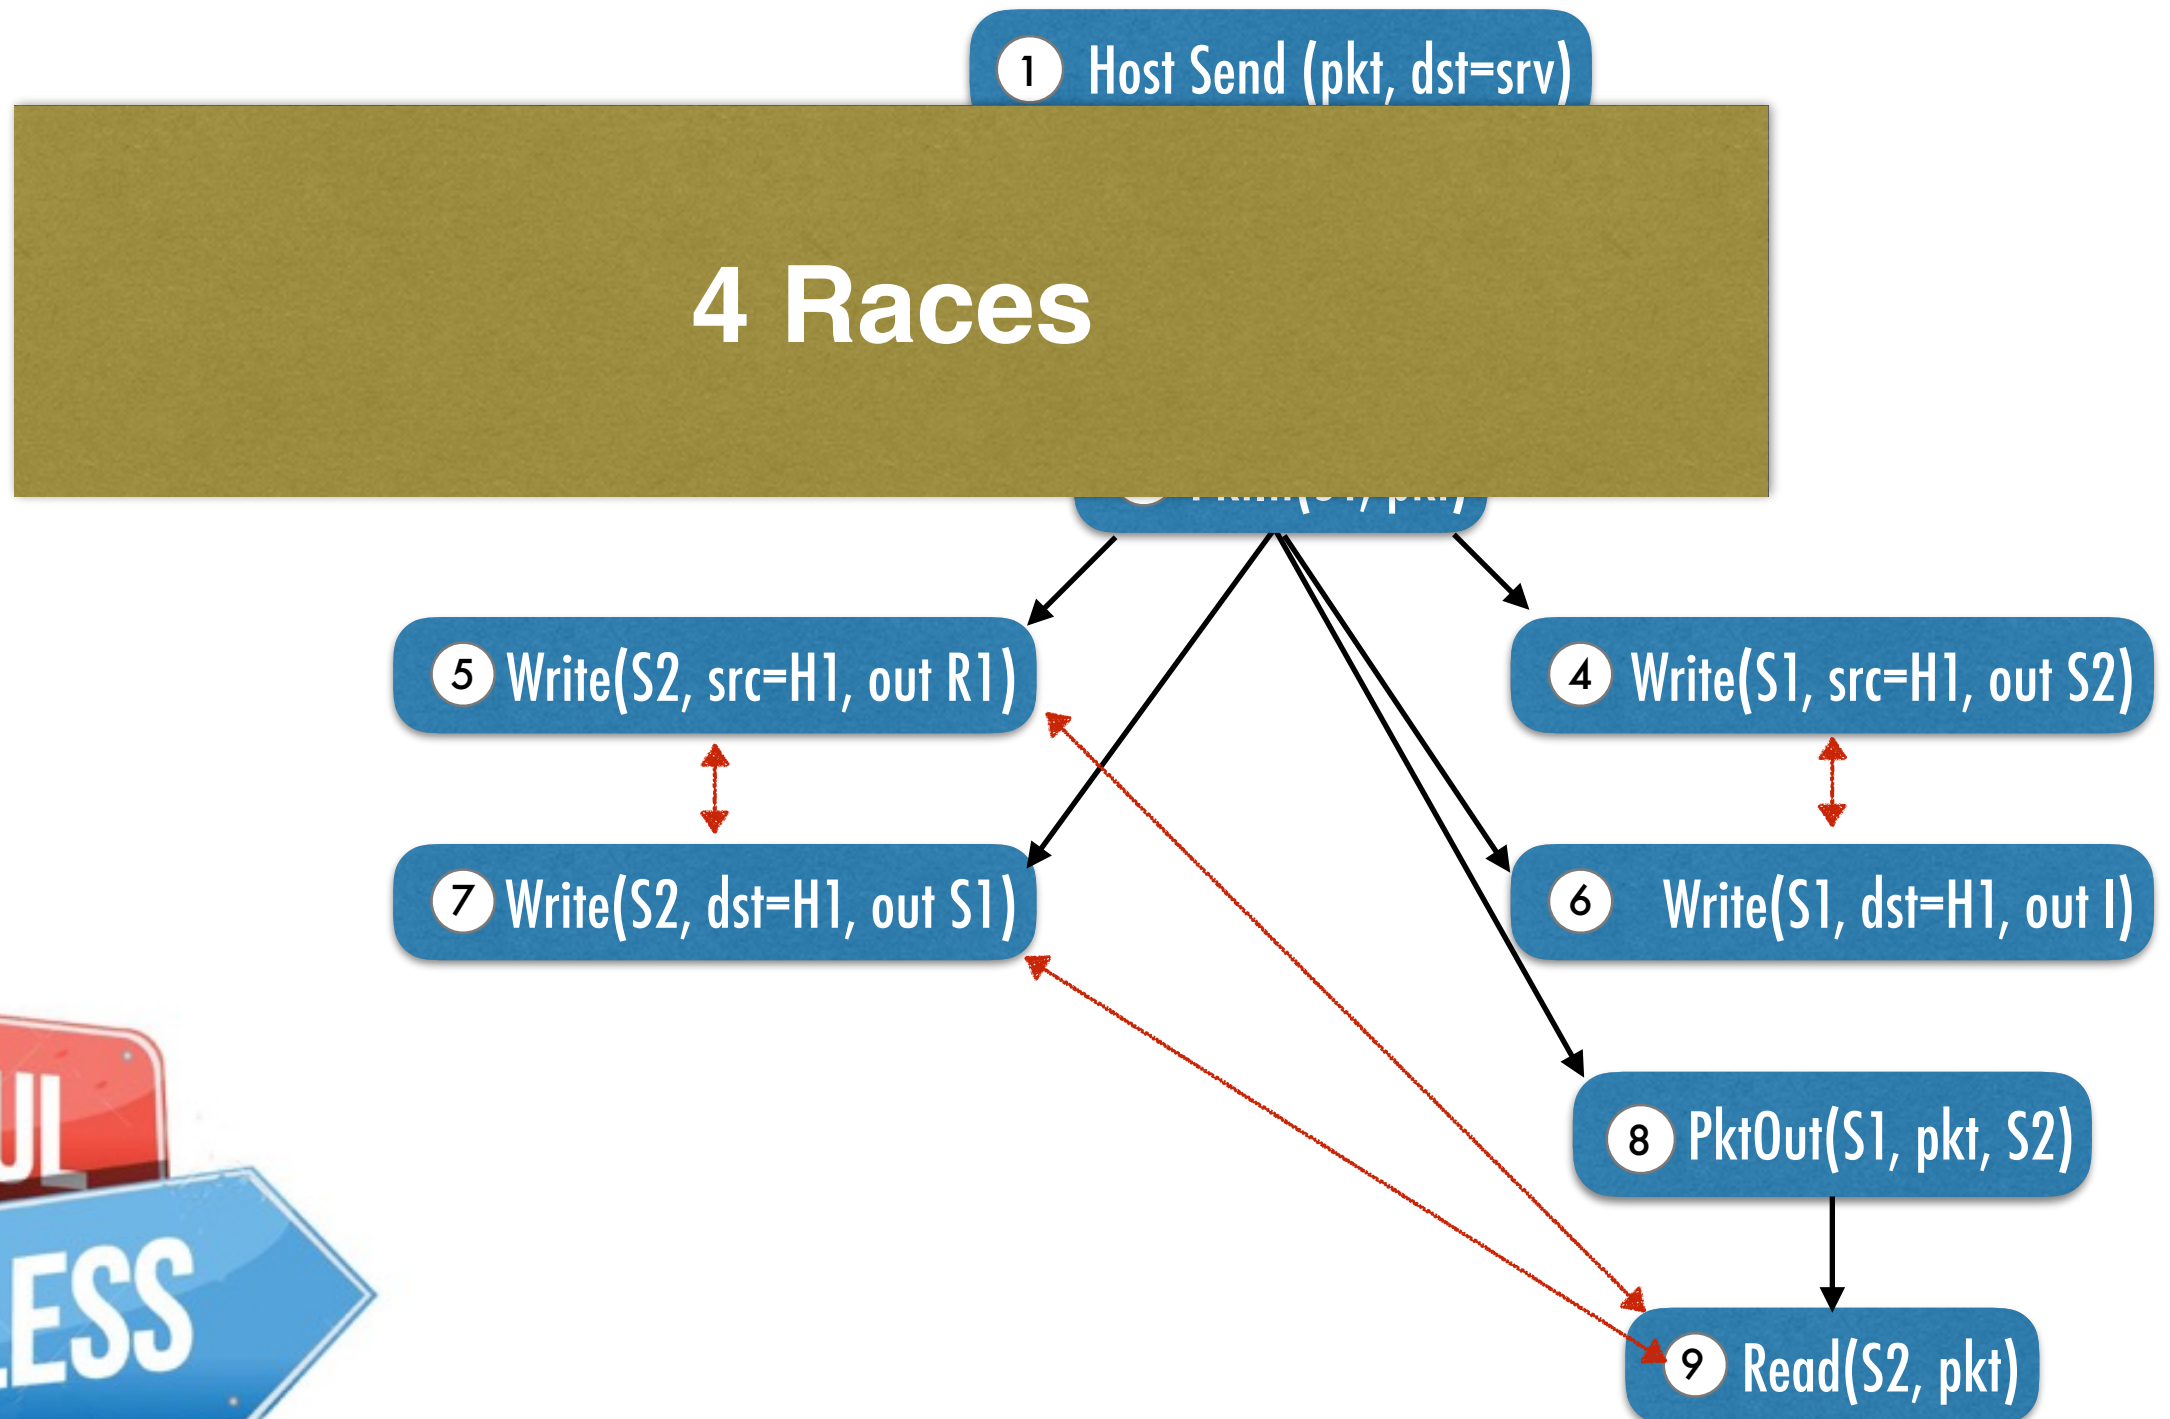
$$\begin{array}{lll}
 \varphi_{\text{mod}(e_1, \text{strict}_1)}^{\text{mod}(e_2, \text{strict}_2)} & := \neg(e_1.m \cap e_2.m \neq \emptyset \wedge e_1.a \neq e_2.a) & \text{if } \neg \text{strict}_1 \wedge \neg \text{strict}_2 \\
 & := \neg(e_1.m = e_2.m \wedge e_1.p = e_2.p \wedge e_1.a \neq e_2.a) & \text{if } \text{strict}_1 \wedge \text{strict}_2 \\
 & \quad \neg((e_1 \stackrel{\text{strict}_2}{\subseteq} e_2 \vee e_2 \stackrel{\text{strict}_1}{\subseteq} e_1) \wedge e_1.a \neq e_2.a) & \text{otherwise} \\
 \\
 \varphi_{\text{add}(e_{\text{add}}, \text{no\_overlap})}^{\text{add}(e_{\text{add}}, \text{no\_overlap})} & := \neg(e_{\text{add}} \stackrel{\text{strict}}{\subseteq} e_{\text{mod}} \wedge e_{\text{add}}.a \neq e_{\text{mod}}.a) & \text{if } \neg \text{no\_overlap} \\
 & \quad \neg(e_{\text{add}} \cap e_{\text{mod}} \neq \emptyset) & \text{otherwise} \\
 \\
 \varphi_{\text{add}(e_1, \text{no\_overlap}_1)}^{\text{add}(e_2, \text{no\_overlap}_2)} & := \neg(e_1.m \cap e_2.m \neq \emptyset \wedge e_1.p = e_2.p) & \text{if } \text{no\_overlap}_1 \vee \text{no\_overlap}_2 \\
 & \quad \neg(e_1.m = e_2.m \wedge e_1.p = e_2.p \wedge e_1.a \neq e_2.a) & \text{otherwise}
 \end{array}$$

**Figure 3:** Commutativity specification of an OpenFlow switch. Two *read* or two *del* operations always commute.

# Without Commutativity

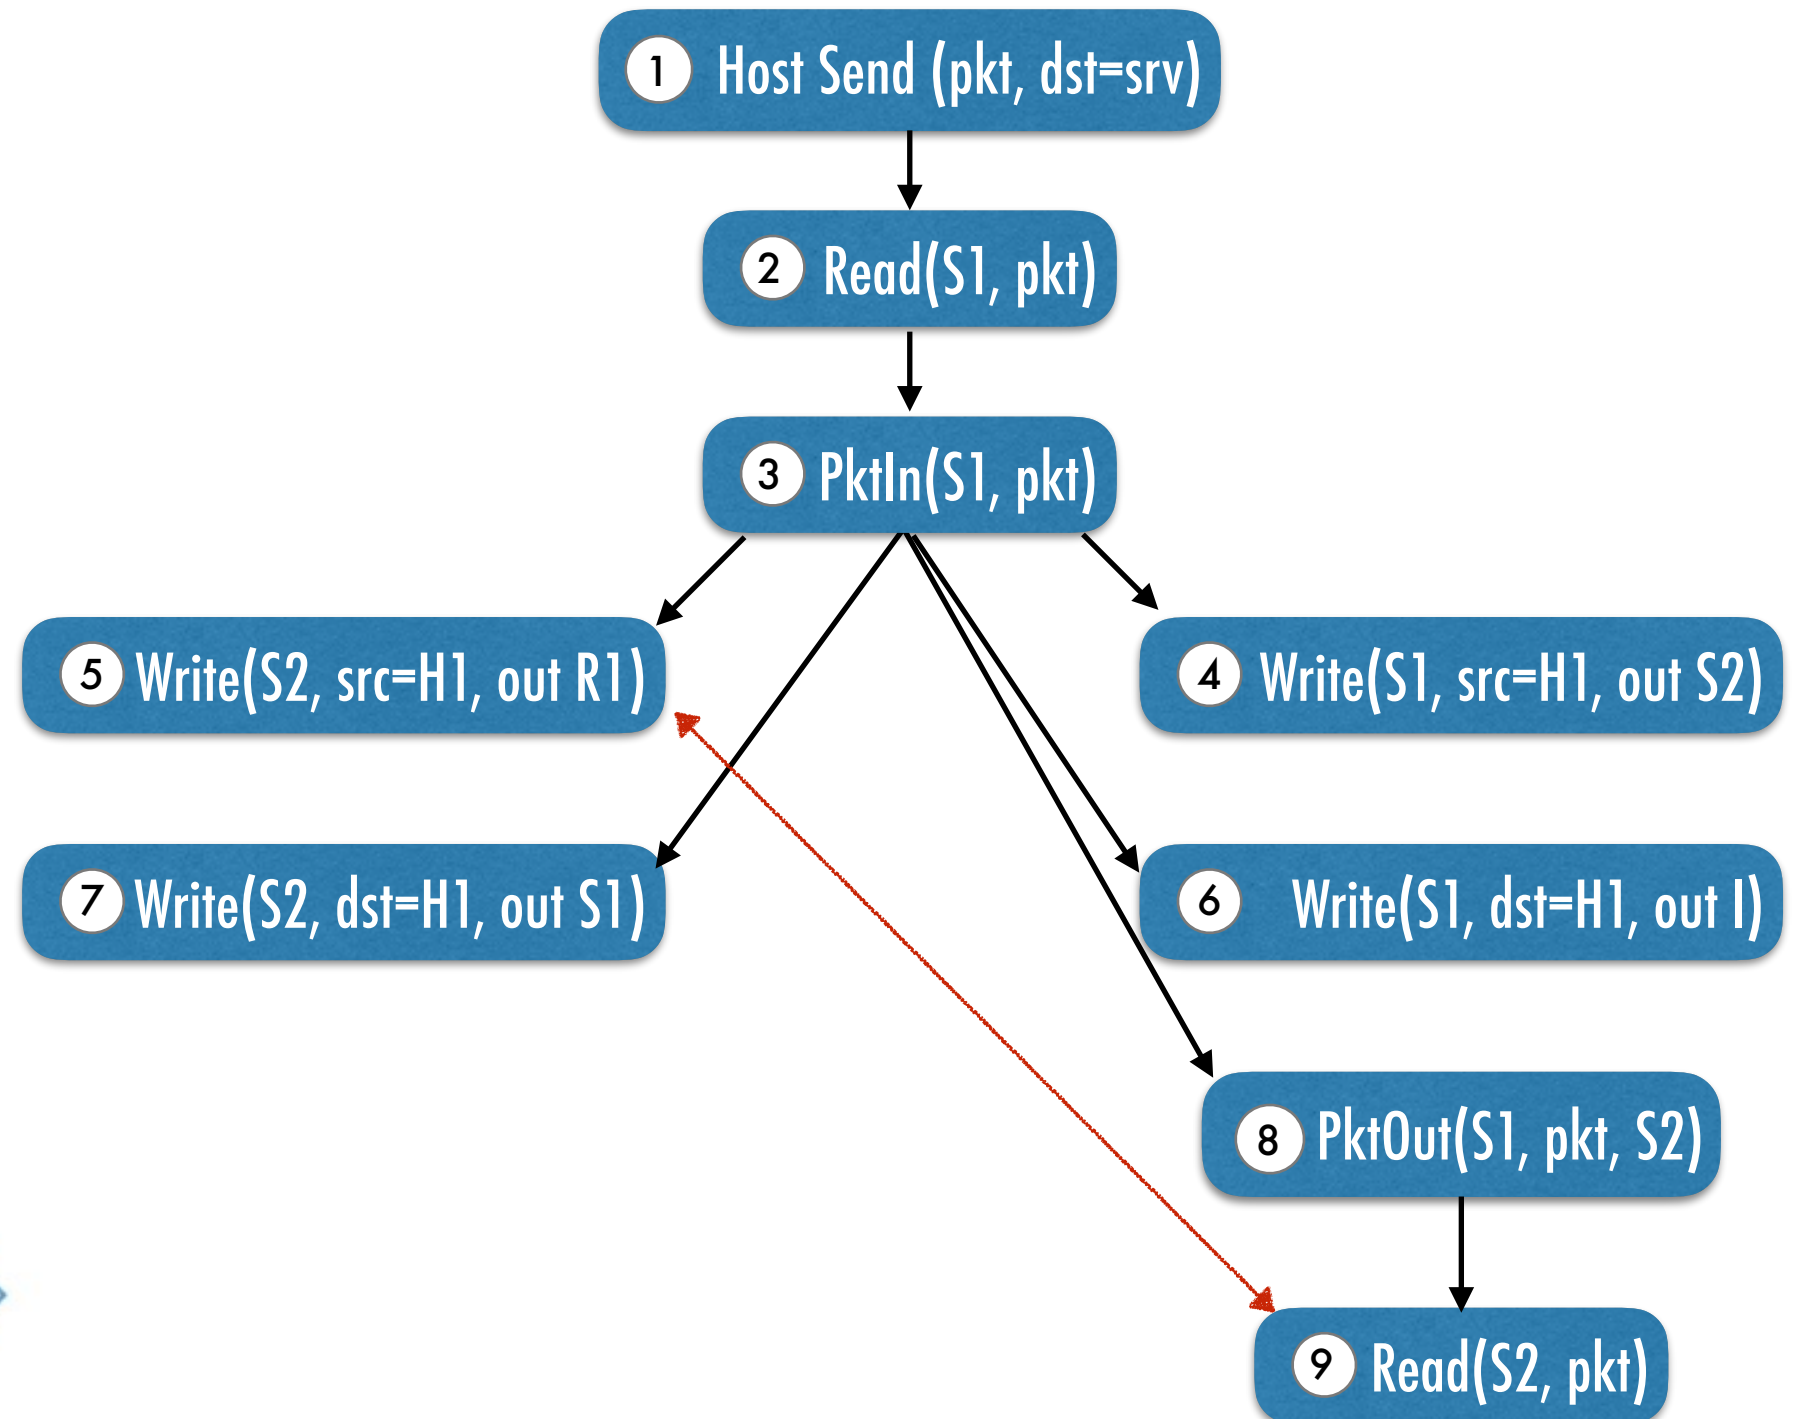


# Without Commutativity

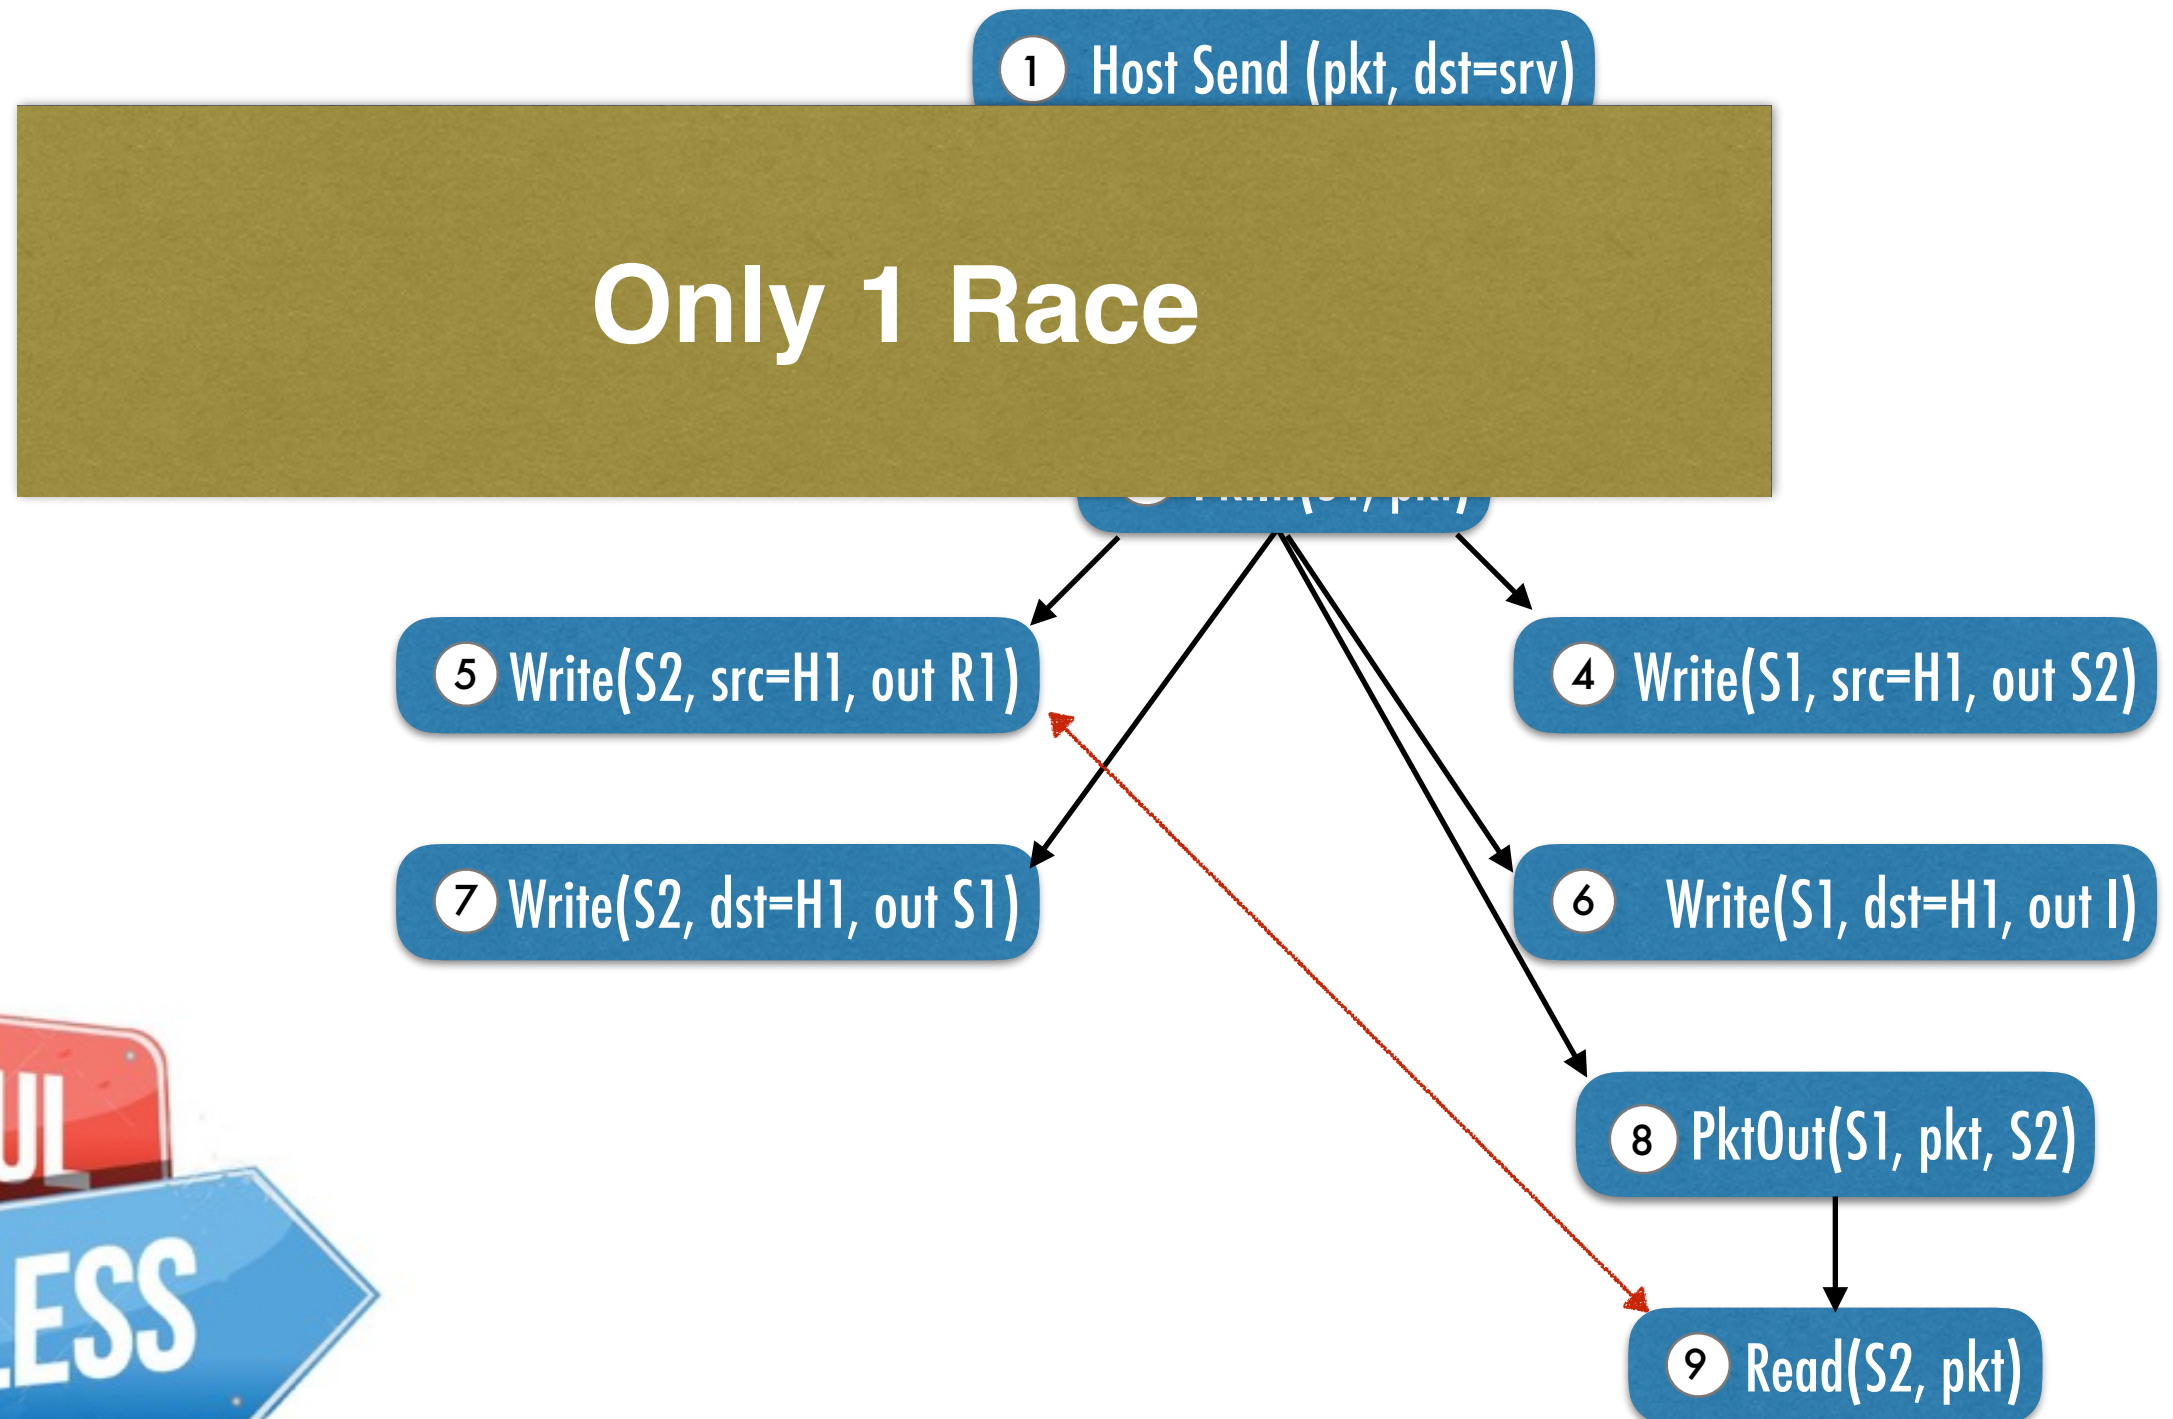




# Using Commutativity



# Using Commutativity





# Effect of Commutativity

For **Floodlight** Load Balancer

**703,864 Races**



**18,706 Races**

# Effect of Commutativity

For **Floodlight** Load Balancer

**703,864 Races**



**18,706 Races**

**What about the remaining 18,706 races???**

# Infeasible races

**Write(S1, pkt, out Replica 1)**



# Infeasible races

**Write(S1, pkt, out Replica 1)**



# Infeasible races

**Write(S1, pkt, out Replica 1)**



**Read(S1, pkt)**



# Time-Based Filter

Add **HB edges** between events that are **cannot be** reordered **due physical limits** of the network.



# Overall reduction

For Floodlight Load Balancer

**703,864 Races**



Commutativity filter

**18,706 Races**



Time-based filter

**2214 (0.31%) Remaining**

Can we leverage the  
race detector to check  
for high-level properties?





# SDNRacer



Detecting concurrency violations



Precise notion of interference



**Checks for high-level properties**



Implementation and evaluation

# Network Update

A set of write events that together reflects high-level network-wide policy change.



# Network Update Example

## Events Trace:

- ① Host Send (pkt, dst=src)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑨ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑩ PktIn(S2, pkt)
- ⑪ Write(S2, src=H1, out S1)
- ⑫ Write(S1, src=H1, out S3)

# Network Update Example

## Events Trace:

- ① Host Send (pkt, dst=src)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ **Write(S1, src=H1, out S2)**
- ⑨ Read(S2, pkt)
- ⑥ **Write(S1, dst=H1, out Internet)**
- ⑦ **Write(S2, dst=H1, out S1)**
- ⑧ PktOut(S1, pkt, S2)
- ⑤ **Write(S2, src=H1, out R1)**
- ⑩ PktIn(S2, pkt)
- ⑪ Write(S2, src=H1, out S1)
- ⑫ Write(S1, src=H1, out S3)

## Update #1

Send Traffic from H1  
to Replica#1

④ Write(S1, src=H1, out S2)

⑤ Write(S2, src=H1, out R1)

⑥ Write(S1, dst=H1, out I)

⑦ Write(S2, dst=H1, out S1)

# Network Update Example

## Events Trace:

- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ **Write(S1, src=H1, out S2)**
- ⑨ Read(S2, pkt)
- ⑥ **Write(S1, dst=H1, out Internet)**
- ⑦ **Write(S2, dst=H1, out S1)**
- ⑧ PktOut(S1, pkt, S2)
- ⑤ **Write(S2, src=H1, out R1)**
- ⑩ PktIn(S2, pkt)
- ⑪ **Write(S2, src=H1, out S1)**
- ⑫ **Write(S1, src=H1, out S3)**

## Update #1

Send Traffic from H1  
to Replica#1

④ Write(S1, src=H1, out S2)

⑤ Write(S2, src=H1, out R1)

⑥ Write(S1, dst=H1, out I)

⑦ Write(S2, dst=H1, out S1)

## Update #2

Send Traffic from H1  
to Replica#2

⑪ Write(S2, src=H1, out S1)

⑫ Write(S1, src=H1, out S3)

# Network Update Isolation Property

Network updates are isolated if they are **serializable**.



# Network Update Isolation

**SDNRacer** checks if there are **no** data races between write events of different update sets.





# Network Update Example

## Events Trace:

- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ **Write(S1, src=H1, out S2)**
- ⑨ Read(S2, pkt)
- ⑥ **Write(S1, dst=H1, out Internet)**
- ⑦ **Write(S2, dst=H1, out S1)**
- ⑧ PktOut(S1, pkt, S2)
- ⑤ **Write(S2, src=H1, out R1)**
- ⑩ PktIn(S2, pkt)
- ⑪ **Write(S2, src=H1, out S1)**
- ⑫ **Write(S1, src=H1, out S3)**

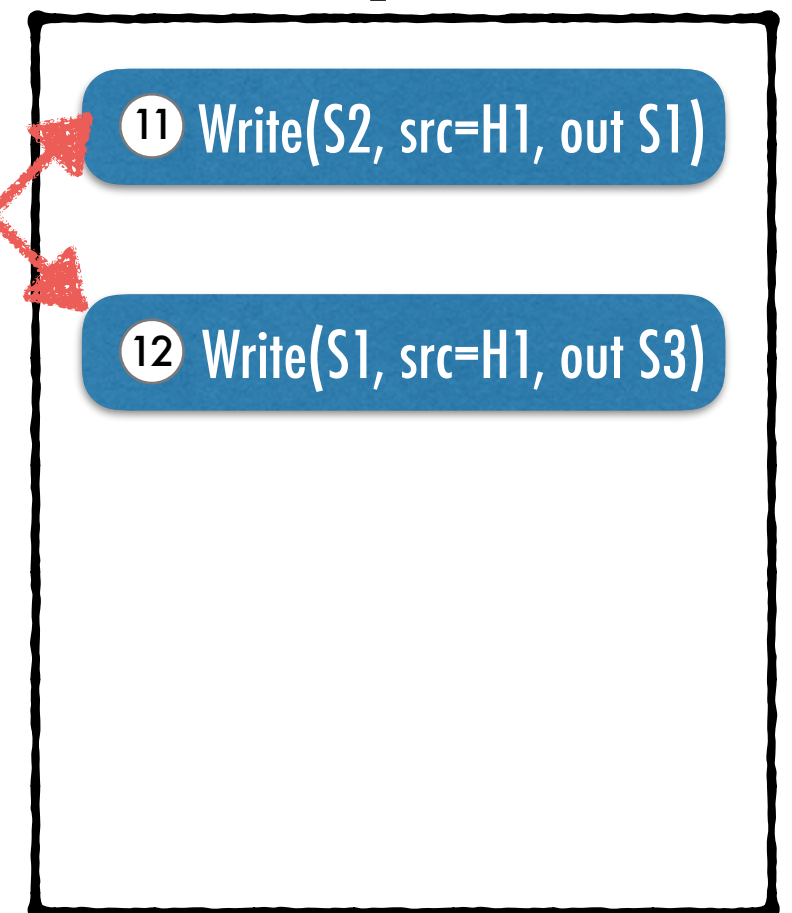
## Update #1

Send Traffic from H1  
to Replica#1



## Update #2

Send Traffic from H1  
to Replica#2





# Network Update Example

## Events Trace:

- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S1)
- ⑤ Read(S2, pkt)
- ⑥ Write(S1, src=H1, out S1)
- ⑦ Write(S2, src=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑨ Write(S2, src=H1, out R1)
- ⑩ PktIn(S2, pkt)
- ⑪ Write(S2, src=H1, out S1)
- ⑫ Write(S1, src=H1, out S3)

### Update #1

Send Traffic from H1  
to Replica#1

### Update #2

Send Traffic from H1  
to Replica#2

Update #1 and Update #2 are not serializable!!!

⑥ Write(S1, dst=H1, out I)

⑦ Write(S2, dst=H1, out S1)

S1)

S3)

# Packet Coherence Property

A **packet** is **coherent** if it is processed entirely using **one consistent global network configuration**.

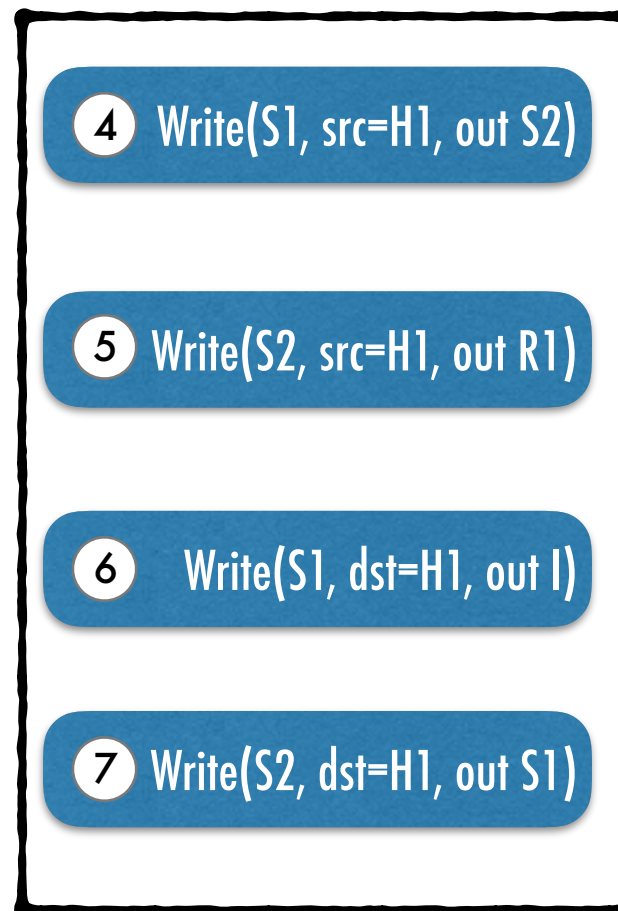


# Packet Coherence Example


## Events Trace:

- ① Host Send (pkt, dst=svr)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑨ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑩ PktIn(S2, pkt)
- ⑪ Write(S2, src=H1, out S1)
- ⑫ Write(S1, src=H1, out S3)

### Update #1 Send Traffic from H1 to Replica#1



### Update #2 Send Traffic from H1 to Replica#2



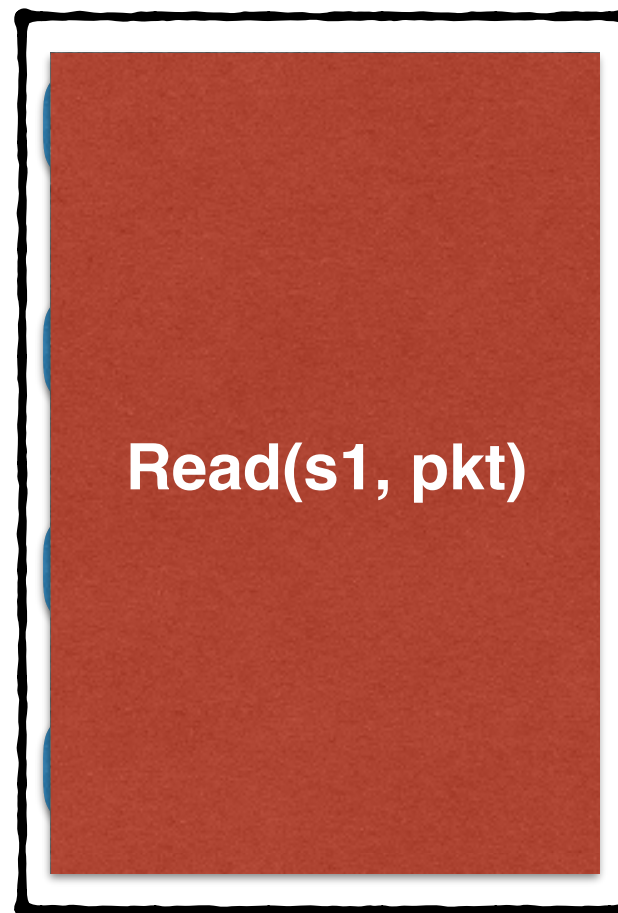
# Packet Coherence Example

## Events Trace:

- ① Host Send (pkt, dst=srv)
- ② Read(S1, pkt)
- ③ PktIn(S1, pkt)
- ④ Write(S1, src=H1, out S2)
- ⑨ Read(S2, pkt)
- ⑥ Write(S1, dst=H1, out Internet)
- ⑦ Write(S2, dst=H1, out S1)
- ⑧ PktOut(S1, pkt, S2)
- ⑤ Write(S2, src=H1, out R1)
- ⑩ PktIn(S2, pkt)
- ⑪ Write(S2, src=H1, out S1)
- ⑫ Write(S1, src=H1, out S3)

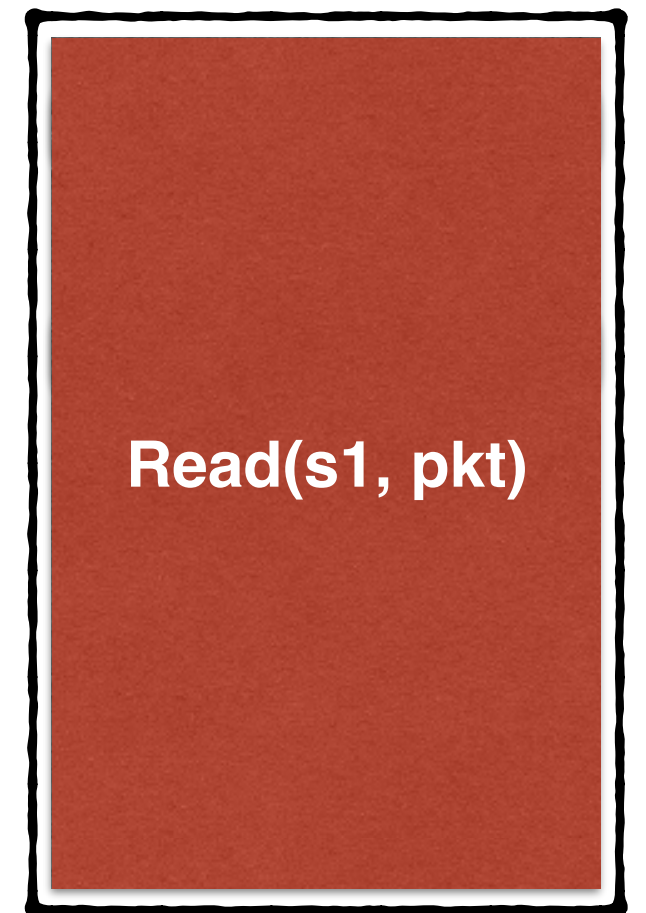
### Update #1

Send Traffic from H1  
to Replica#1



### Update #2

Send Traffic from H1  
to Replica#2



# Packet Coherence Check

1. Take note of the first version the packet sees.
2. Check at every switch if the read operations triggered by the packet races with write operations from different version than originally observed.

# SDNRacer: Guarantees

- **SDNRacer** checks are *more general* than other tools which take a *snapshot* of the network.
- **SDNRacer** guarantees that the properties will hold for all possible reordering of the given trace.

**Benefit:** Fewer traces to explore

# SDNRacer



Detecting concurrency violations



Precise notion of interference



Checks for high-level properties



**Implementation and evaluation**

# Implementation

- [Instrumentation](#) of SDN troubleshooting system ([STS](#)) [SIGCOMM'14].
- [Concurrency analyzer](#) that implements happens-before rules, commutativity specification and speculative time-based filter.
- [Property Checker](#) for network update isolation and packet coherence violations.
- Around 3,000 of Python code

<http://sdnracer.ethz.ch>



# Evaluation: Controllers

Tested with [off-the-shelf](#) controllers:



# Evaluation: Applications

We tested **SDNRacer** with five different applications shipped with each controller:

- MAC-learning
- Forwarding
- Circuit Pusher
- Admission Control
- Load Balancer

# Experimental Setup

- Three different network topologies: Single Switch, Two switches, and Binary Tree.
- Run every controller and every application, if possible, with randomly generated input.
- We collected 29 traces.

**SDNRacer** filters more than **90%** of all  
races in **89%** of examined traces.


# Detected Bugs

**SDNRacer** detected **two** update isolation violations


Violation #1: Floodlight Load Balancer distributes flows inconsistently

Violation #2: POX forwarding module deletes rules installed by other modules

In **90%** of the studied traces **SDNRacer** can analyze the traces in **less** than **30 seconds**.



In **90%** of the cases **SDNRacer** can analyze the traces in **less** than **30 seconds**.



# Conclusion

## **Happens-Before Model for SDN**

Captures asynchrony of SDN

## **Flow Table Commutativity Spec**

Captures Interference

## **Concurrency Analysis**

- Race Freedom
- Network Update Isolation
- Packet Coherence

## **Implementation and Evaluation**

- Found bugs in existing apps: ONOS, POX, Floodlight

<http://sdnracer.ethz.ch>